

mapreduce-on-yarn任务提交源码分析

mapreduce-on-yarn任务提交源码分析

一 官方 wordcount 任务执行案例

二 提交任务入口类

2.1 基于 hadoop/yarn jar 命令提交 (RunJar)

2.2 基于 mapred jar 命令提交 (JobClient)

2.3 使用 RunJar 作为入口类

三 任务启动类入口

3.1 连接 ResourceManager 获取 ClientRMService 代理 RPC 客户端

3.1.1 创建 Cluster

3.1.1.1 创建 YARNRunner

3.1.1.1.1 创建 ResourceMgrDelegate(真正 RPC 客户端 YarnClientImpl)

3.1.1.2 初始化 YANRRunner

3.2 提交任务 Job

3.2.1 向 ResourceManager 的 ClientRMService RPC 申请任务 ApplicationId

3.2.2 获取任务临时目录

3.2.3 提交 job.jar 信息到临时目录

3.2.4 切片相关信息(job.split, job.splitmetainfo)

3.2.4.1 计算切片

3.2.4.2 上传切片信息到临时目录

3.2.5 上传任务配置信息文件 job.xml

3.2.6 RPC 客户端提交任务到 Yarn (调用 YARNRunner.submitJob())

3.2.6.1 提交任务信息封装成 ApplicationSubmissionContext (启动 MRAppMaster)

3.2.6.2 提交任务(调用 YarnClientImpl.submitApplication())

四 ResourceManager 的 ClientRMService 接收到客户端发送的 RPC 请求 (启动 AM)

4.1 备注说明

4.2 启动 AM (MRAppMaster) 调用 ApplicationMasterLauncher.handle()

4.2.1 执行启动 AM 调用 AMLauncher.run()

4.2.1.1 获取 ContainerManagerImpl 的 RPC 客户端代理

4.2.1.2 发送 RPC 请求到 NodeManager 的 ContainerManagerImpl RPC 服务启动容器

4.2.1.2.1 备注说明

4.2.1.2.2 调度执行 AM 容器(ContainerScheduler.handle(SCHEDULE_CONTAINER))

4.2.1.2.3 真正启动 AM 容器(调用ContainersLauncher.handle(LAUNCH_CONTAINER))

五 启动 MRAppMaster 进程 (MR 任务的ApplicationMaster)

5.1 创建 MRAppMaster 组合服务

5.2 初始化并启动 AM

5.2.1 初始化 MRAppMaster (调用 MRAppMaster.serviceInit())

5.2.1.1 异步事件分发器 AsyncDispatcher.serviceInit()

5.2.1.2 任务完成监控服务 TaskAttemptFinishingMonitor.serviceInit()

5.2.1.3 提交者事件处理 CommitterEventHandler.serviceInit()

5.2.1.4 Task监听服务 TaskAttemptListenerImpl.serviceInit()

5.2.1.5 任务推测执行服务 DefaultSpeculator.serviceInit()

5.2.1.6 任务清除目录服务 StagingDirCleaningService.serviceInit()

5.2.1.7 容器申请服务 ContainerAllocatorRouter.serviceInit()

5.2.1.8 启动容器服务 ContainerLauncherRouter.serviceInit()

5.2.2 启动 MRAppMaster (调用 MRAppMaster.serviceStart())

5.2.2.1 启动 MRClientService RPC 服务 其调用 serviceStart()

5.2.2.2 异步事件分发器 AsyncDispatcher.serviceStart()

5.2.2.3 任务完成监控服务 TaskAttemptFinishingMonitor.serviceStart()

- 5.2.2.4 提交者事件处理 CommitterEventHandler.serviceStart()
 - 5.2.2.5 Task监听服务 TaskAttemptListenerImpl.serviceStart()
 - 5.2.2.6 任务推测执行服务 DefaultSpeculator.serviceStart()
 - 5.2.2.7 任务清除目录服务 StagingDirCleaningService.serviceStart()
 - 5.2.2.8 容器申请服务 ContainerAllocatorRouter.serviceStart()
 - 5.2.2.8.1 创建从 RM 申请容器服务 RMContainerAllocator
 - 5.2.2.8.2 初始化调用 RMContainerAllocator.serviceInit()
 - 5.2.2.8.3 启动调用 RMContainerAllocator.serviceStart()
 - 5.2.2.9 启动容器服务 ContainerLauncherRouter.serviceStart()
 - 5.2.2.9.1 创建启动容器服务 ContainerLauncherImpl
 - 5.2.2.9.2 初始化调用 ContainerLauncherImpl.serviceInit()
 - 5.2.2.9.3 启动调用 ContainerLauncherImpl.serviceStart()
 - 5.4 AM 的所有组件基本启动完成开始执行任务
 - 5.4.1 启动 MR 任务创建待执行 MapTask、ReduceTask 的源码分析
 - 5.4.2 开始执行任务经过多种状态转换最终调用 SetupCompletedTransition.doTransition() 执行 MapTask 或者 ReduceTask
 - 5.4.3 通过各种状态转换之后向 RM 的 ContainerManagerImpl RPC 服务申请容器启动MapTask或者 ReduceTask源码分析
 - 5.4.3.1 处理 MapTask 容器请求
 - 5.4.3.1.1 AM 申请容器启动 MapTask 容器是通过发送心跳向 RM 的 ContainerManagerImpl RPC 服务申请资源(调用ApplicationMasterService.allocate())
 - 5.4.3.1.2 容器申请好之后开始做状态准备以及分配容器
 - 5.4.3.1.3 分配容器的状态转换之后最终调用 ContainerAssignedTransition.doTransition(TaskAttemptEventType.TA_ASSIGNED) 构建启动容器的进程(YarnChild)上下文环境
 - 5.4.3.1.4 启动容器进程(YarnChild)相关环境准备好之后连接 NM 的 ContainerManagerImpl 服务发送 RPC 请求启动容器 YarnChild 进程(调用 ContainerManagerImpl.startContainers() 类似启动 MRAppMaster 进程)
 - 5.4.3.2 处理 ReduceTask 容器请求(类似启动 MapTask)
- ## 六 启动 MR 真正处理任务 YarnChild 进程
- 6.1 获取 AM(MRAppMaster) 的 Task 监听服务 TaskAttemptListenerImpl RPC 服务客户端代理
 - 6.2 发送 RPC 请求 TaskAttemptListenerImpl RPC 服务获取执行 Task
 - 6.3 执行 MapTask或者ReduceTask 业务逻辑
 - 6.3.1 执行 Mapper 调用 MapTask.run()
 - 6.3.1 执行 Reducer调用 ReduceTask.run()


```
# hadoop.sh/yarn.sh
jar)
    if [[ -n "${YARN_OPTS}" ]] || [[ -n "${YARN_CLIENT_OPTS}" ]]; then
        hadoop_error "WARNING: Use \"yarn jar\" to launch YARN applications."
    fi
    if [[ -z $1 || $1 = "--help" ]]; then
        echo "Usage: hadoop jar <jar> [mainClass] args..."
        exit 0
    fi
    HADOOP_CLASSNAME=org.apache.hadoop.util.RunJar
;;
```

2.2 基于 mapred jar 命令提交 (JobClient)

```
# mapred.sh
job)
    HADOOP_CLASSNAME=org.apache.hadoop.mapred.JobClient
```

2.3 使用 RunJar 作为入口类

```
/** Run a Hadoop job jar. */
@interfaceAudience.Private
@interfaceStability.Unstable
public class RunJar {
    /** Run a Hadoop job jar. If the main class is not in the jar's manifest,
     * then it must be provided on the command line. */
    public static void main(String[] args) throws Throwable {
        new RunJar().run(args);
    }
}
```

```
public void run(String[] args) throws Throwable {
    String usage = "RunJar jarFile [mainClass] args...";

    if (args.length < 1) {
        System.err.println(usage);
        System.exit(-1);
    }

    int firstArg = 0;
    // 判断入口 jar 是否存在
    String fileName = args[firstArg++];
    File file = new File(fileName);
    if (!file.exists() || !file.isFile()) {
        System.err.println("JAR does not exist or is not a normal file: " +
            file.getCanonicalPath());
        System.exit(-1);
    }
    String mainClassName = null;
```

```

JarFile jarFile;
try {
    // 读取 jar 的启动类
    jarFile = new JarFile(fileName);
} catch (IOException io) {
    throw new IOException("Error opening job jar: " + fileName)
        .initCause(io);
}
Manifest manifest = jarFile.getManifest();
if (manifest != null) {
    mainClassName = manifest.getMainAttributes().getValue("Main-Class");
}
jarFile.close();

if (mainClassName == null) {
    if (args.length < 2) {
        System.err.println(usage);
        System.exit(-1);
    }
    mainClassName = args[firstArg++];
}
mainClassName = mainClassName.replaceAll("/", ".");

File tmpDir = new File(System.getProperty("java.io.tmpdir"));
ensureDirectory(tmpDir);

final File workDir;
try {
    workDir = File.createTempFile("hadoop-unjar", "", tmpDir);
} catch (IOException ioe) {
    // If user has insufficient perms to write to tmpDir, default
    // "Permission denied" message doesn't specify a filename.
    System.err.println("Error creating temp dir in java.io.tmpdir "
        + tmpDir + " due to " + ioe.getMessage());
    System.exit(-1);
    return;
}

if (!workDir.delete()) {
    System.err.println("Delete failed for " + workDir);
    System.exit(-1);
}
ensureDirectory(workDir);

ShutdownHookManager.get().addShutdownHook(
    new Runnable() {
        @Override
        public void run() {
            FileUtil.fullyDelete(workDir);
        }
    }, SHUTDOWN_HOOK_PRIORITY);

```

```

unJar(file, workDir);

// 创建类加载器并替换默认当前线程的类加载器
ClassLoader loader = createClassLoader(file, workDir);
Thread.currentThread().setContextClassLoader(loader);

// 反射获取任务启动类的 main()
Class<?> mainClass = Class.forName(mainClassName, true, loader);
Method main = mainClass.getMethod("main", String[].class);
List<String> newArgsSubList = Arrays.asList(args)
    .subList(firstArg, args.length);
String[] newArgs = newArgsSubList
    .toArray(new String[newArgsSubList.size()]);
try {
    // 调用任务启动类的 main() (比如 wordCount.main())
    main.invoke(null, new Object[]{newArgs});
} catch (InvocationTargetException e) {
    throw e.getTargetException();
}
}

```

三 任务启动类入口

```

public class WordCount {

    public static class TokenizerMapper
        extends Mapper<Object, Text, Text, IntWritable> {

        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(Object key, Text value, Context context
        ) throws IOException, InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken());
                context.write(word, one);
            }
        }
    }

    public static class IntSumReducer
        extends Reducer<Text, IntWritable, Text, IntWritable> {
        private IntWritable result = new IntWritable();

        public void reduce(Text key, Iterable<IntWritable> values,
            Context context
        ) throws IOException, InterruptedException {
            int sum = 0;

```

```

        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}

public static void main(String[] args) throws Exception {
    // 创建 Configuration 并解析入口参数
    Configuration conf = new Configuration();
    String[] otherArgs = new GenericOptionsParser(conf, args)
        .getRemainingArgs();
    if (otherArgs.length < 2) {
        System.err.println("Usage: wordcount <in> [<in>...] <out>");
        System.exit(2);
    }
    // MapReduce 程序的标配
    Job job = Job.getInstance(conf, "word count");
    job.setJarByClass(WordCount.class);
    job.setMapperClass(TokenizerMapper.class);
    job.setCombinerClass(IntSumReducer.class);
    job.setReducerClass(IntSumReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);

    // wordCount 的输出输出路径
    for (int i = 0; i < otherArgs.length - 1; ++i) {
        FileInputFormat.addInputPath(job, new Path(otherArgs[i]));
    }
    FileOutputFormat.setOutputPath(job,
        new Path(otherArgs[otherArgs.length - 1]));

    // 执行任务等待完成
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

```

/**
 * Submit the job to the cluster and wait for it to finish.
 *
 * @param verbose print the progress to the user
 * @return true if the job succeeded
 * @throws IOException thrown if the communication with the
 *         <code>JobTracker</code> is lost
 */
public boolean waitForCompletion(boolean verbose
) throws IOException, InterruptedException,
    ClassNotFoundException {
    // state 初始化状态为 DEFINE
    if (state == JobState.DEFINE) {

```

```

        // 提交任务
        submit();
    }
    // 是否打印任务进度给用户
    if (verbose) {
        monitorAndPrintJob();
    } else {
        // get the completion poll interval from the client.
        int completionPollIntervalMillis =
            Job.getCompletionPollInterval(cluster.getConf());
        while (!isComplete()) {
            try {
                Thread.sleep(completionPollIntervalMillis);
            } catch (InterruptedException ie) {
            }
        }
    }
    return isSuccessful();
}

```

```

/**
 * Submit the job to the cluster and return immediately.
 *
 * @throws IOException
 */
public void submit()
    throws IOException, InterruptedException, ClassNotFoundException {
    // 确保任务状态 state 为 DEFINE
    ensureState(JobState.DEFINE);

    // 设置任务使用新 API
    setUseNewAPI();

    // 连接 ResourceManager 获取 ClientRMService 代理 RPC 客户端
    /**
     * 1 创建 Cluster 其构造方法 创建 YARNRunner 服务
     * 2 YARNRunner 服务的构造方法创建 ResourceMgrDelegate 服务
     * 3 ResourceMgrDelegate 服务的构造方法创建 YarnClientImpl 服务
     * 4.1 ResourceMgrDelegate 服务的构造方法调用其 serviceInit() 进而调用
    YarnClientImpl.serviceInit()
     *     YarnClientImpl.serviceInit() 从配置文件获取对应的信息
     * 4.2 ResourceMgrDelegate 服务的构造方法调用其 serviceStart() 进而调用
    YarnClientImpl.serviceStart()
     *     获取 ResourceManager 组合服务的子服务 ClientRMService RPC Server (端口
    8032)
     *     的代理对象 RPC 通讯协议接口为 ApplicationClientProtocol
    */
    connect();

    // 获取任务提交对象 JobSubmitter
    final JobSubmitter submitter =

```



```

        getJobSubmitter(
            cluster.getFileSystem(),
            cluster.getClient());
    status = ugi.doAs(new PrivilegedExceptionAction<JobStatus>() {
        public JobStatus run() throws IOException, InterruptedException,
            ClassNotFoundException {
            // 提交任务
            return submitter.submitJobInternal(Job.this, cluster);
        }
    });
    state = JobState.RUNNING;
    // The url to track the job:
    http://hadoop103:8088/proxy/application_1684656010852_0001/
    LOG.info("The url to track the job: " + getTrackingURL());
}

```

3.1 连接 ResourceManager 获取 ClientRMService 代理 RPC 客户端

```

private synchronized void connect()
    throws IOException, InterruptedException, ClassNotFoundException {
    if (cluster == null) {
        cluster =
            ugi.doAs(new PrivilegedExceptionAction<Cluster>() {
                public Cluster run()
                    throws IOException, InterruptedException,
                    ClassNotFoundException {
                    // 创建 Cluster
                    return new Cluster(getConfiguration());
                }
            });
    }
}

```

3.1.1 创建 Cluster

```

/**
 * Provides a way to access information about the map/reduce cluster.
 */
@InterfaceAudience.Public
@InterfaceStability.Evolving
public class Cluster {
    public Cluster(Configuration conf) throws IOException {
        // 往下追
        this(null, conf);
    }
}

```

```

public Cluster(InetSocketAddress jobTrackAddr, Configuration conf)
    throws IOException {
    this.conf = conf;
    this.ugi = UserGroupInformation.getCurrentUser();
    // 初始化 Cluster
    initialize(jobTrackAddr, conf);
}

```

```

private void initialize(InetSocketAddress jobTrackAddr, Configuration conf)
    throws IOException {

    // 加载 ClientProtocolProvider 所有子类
    // 1 LocalClientProtocolProvider (本地)
    // 2 YarnClientProtocolProvider (yarn)
    initProviderList();
    final IOException initEx = new IOException(
        "Cannot initialize Cluster. Please check your configuration for "
        + MRConfig.FRAMEWORK_NAME
        + " and the correspond server addresses.");
    // jobTrackAddr 默认 null
    if (jobTrackAddr != null) {
        LOG.info(
            "Initializing cluster for Job Tracker=" +
jobTrackAddr.toString());
    }
    for (ClientProtocolProvider provider : providerList) {
        LOG.debug("Trying ClientProtocolProvider : "
            + provider.getClass().getName());
        ClientProtocol clientProtocol = null;
        try {
            if (jobTrackAddr == null) {
                // 根据 key = mapreduce.framework.name value = yarn (一般在
mapred-site.xml 配置)
                // 故直接调用 YarnClientProtocolProvider.create() 返回 YARNRunner
                clientProtocol = provider.create(conf);
            } else {
                clientProtocol = provider.create(jobTrackAddr, conf);
            }

            if (clientProtocol != null) {
                // 基于 yarn 方式提交任务 clientProtocolProvider =
YarnClientProtocolProvider
                clientProtocolProvider = provider;
                // YARNRunner
                client = clientProtocol;
                LOG.debug("Picked " + provider.getClass().getName()
                    + " as the ClientProtocolProvider");
                break;
            } else {
                LOG.debug("Cannot pick " + provider.getClass().getName()

```

```

        + " as the ClientProtocolProvider - returned null
protocol");
    }
    } catch (Exception e) {
        final String errMsg = "Failed to use " +
provider.getClass().getName()
        + " due to error: ";
        initEx.addSuppressed(new IOException(errMsg, e));
        LOG.info(errMsg, e);
    }
}

if (null == clientProtocolProvider || null == client) {
    throw initEx;
}
}
}

```

3.1.1.1 创建 YARNRunner

```

// 调用 YarnClientProtocolProvider.create()
public class YarnClientProtocolProvider extends ClientProtocolProvider {

    @Override
    public ClientProtocol create(Configuration conf) throws IOException {
        if (MRConfig.YARN_FRAMEWORK_NAME.equals(conf.get(MRConfig.FRAMEWORK_NAME))) {
            // 创建 YARNRunner
            return new YARNRunner(conf);
        }
        return null;
    }
}

```

```

/**
 * This class enables the current JobClient (0.22 hadoop) to run on YARN.
 */
@SuppressWarnings("unchecked")
public class YARNRunner implements ClientProtocol {
    /**
     * Yarn runner incapsulates the client interface of
     * yarn
     *
     * @param conf the configuration object for the client
     */
    public YARNRunner(Configuration conf) {
        // 往下追
        this(conf,
            // 创建资源管理委托服务 ResourceMgrDelegate (与 ResourceManager 打交道)
            new ResourceMgrDelegate(new YarnConfiguration(conf))
        );
    }
}

```

3.1.1.1.1 创建 ResourceMgrDelegate(真正 RPC 客户端 YarnClientImpl)

```
public class ResourceMgrDelegate extends YarnClient {
    /**
     * Delegate responsible for communicating with the Resource Manager's
     * {@link ApplicationClientProtocol}.
     *
     * @param conf the configuration object.
     */
    public ResourceMgrDelegate(YarnConfiguration conf) {
        super(ResourceMgrDelegate.class.getName());
        this.conf = conf;
        // 创建 Yarn 客户端实例服务 YarnClientImpl
        this.client = YarnClient.createYarnClient();
        // 初始化 (调用 this.serviceInit())
        init(conf);
        // 启动 (调用 this.serviceStart())
        start();
    }
}
```

```
@Override
protected void serviceInit(Configuration conf) throws Exception {
    // 初始化 Yarn 客户端实例服务 (调用 YarnClientImpl.serviceInit())
    client.init(conf);
    super.serviceInit(conf);
}
```

```
@Override
protected void serviceStart() throws Exception {
    // 启动 Yarn 客户端实例服务 (调用 YarnClientImpl.serviceStart())
    client.start();
    super.serviceStart();
}
```

```
// 调用 YarnClientImpl.serviceInit()
public class YarnClientImpl extends YarnClient {
    @Override
    protected void serviceInit(Configuration conf) throws Exception {
        // 获取相关配置信息

        // 默认 200ms
        asyncApiPollIntervalMillis =

        conf.getLong(YarnConfiguration.YARN_CLIENT_APPLICATION_CLIENT_PROTOCOL_POLL_INTERVAL_MS,

        YarnConfiguration.DEFAULT_YARN_CLIENT_APPLICATION_CLIENT_PROTOCOL_POLL_INTERVAL_MS)
        ;

        asyncApiPollTimeoutMillis =
```

```

    conf.getLong(YarnConfiguration.YARN_CLIENT_APPLICATION_CLIENT_PROTOCOL_POLL_TIMEOUT_MS,

    YarnConfiguration.DEFAULT_YARN_CLIENT_APPLICATION_CLIENT_PROTOCOL_POLL_TIMEOUT_MS);
    submitPollIntervalMillis = asyncApiPollIntervalMillis;
    if (conf.get(YarnConfiguration.YARN_CLIENT_APP_SUBMISSION_POLL_INTERVAL_MS)
        != null) {
        submitPollIntervalMillis = conf.getLong(
            YarnConfiguration.YARN_CLIENT_APP_SUBMISSION_POLL_INTERVAL_MS,

    YarnConfiguration.DEFAULT_YARN_CLIENT_APPLICATION_CLIENT_PROTOCOL_POLL_INTERVAL_MS)
;
    }

    if (YarnConfiguration.timelineServiceV1Enabled(conf)) {
        timelineV1ServiceEnabled = true;
        timelineDTRenewer = getTimelineDelegationTokenRenewer(conf);
        timelineService = TimelineUtils.buildTimelineTokenService(conf);
    }

    if (YarnConfiguration.timelineServiceV2Enabled(conf)) {
        timelineV2ServiceEnabled = true;
    }

    // The AHSClientService is enabled by default when we start the
    // TimelineServer which means we are able to get history information
    // for applications/applicationAttempts/containers by using ahsClient
    // when the TimelineServer is running.
    if (timelineV1ServiceEnabled || conf.getBoolean(
        YarnConfiguration.APPLICATION_HISTORY_ENABLED,
        YarnConfiguration.DEFAULT_APPLICATION_HISTORY_ENABLED)) {
        historyServiceEnabled = true;
        historyClient = AHSClient.createAHSClient();
        historyClient.init(conf);
    }

    if (timelineV2ServiceEnabled) {
        ahsV2Client = AHSClient.createAHSv2Client();
        ahsV2Client.init(conf);
    }

    timelineServiceBestEffort = conf.getBoolean(
        YarnConfiguration.TIMELINE_SERVICE_CLIENT_BEST_EFFORT,
        YarnConfiguration.DEFAULT_TIMELINE_SERVICE_CLIENT_BEST_EFFORT);

    loadResourceTypesFromServer = conf.getBoolean(
        YarnConfiguration.YARN_CLIENT_LOAD_RESOURCETYPES_FROM_SERVER,

    YarnConfiguration.DEFAULT_YARN_CLIENT_LOAD_RESOURCETYPES_FROM_SERVER);

    super.serviceInit(conf);

```

```

    }
}

```

```

// 调用 YarnClientImpl.serviceStart()
@Override
protected void serviceStart() throws Exception {
    try {
        // 获取 ResourceManager 组合服务的子服务 ClientRMService RPC Server (端口
8032)

        // 的代理对象 RPC 通讯协议接口为 ApplicationClientProtocol
rmClient = ClientRMProxy.createRMProxy(getConfig(),
        ApplicationClientProtocol.class);

        if (historyServiceEnabled) {
            historyClient.start();
        }
        if (timelineV2ServiceEnabled) {
            ahsv2Client.start();
        }
    } catch (IOException e) {
        throw new YarnRuntimeException(e);
    }

    // Reinitialize local resource types cache from list of resources pulled
from
    // RM.
    if (loadResourceTypesFromServer) {
        ResourceUtils.reinitializeResources(getResourceTypeInfo());
    }

    super.serviceStart();
}

```

3.1.1.2 初始化 YANRRunner

```

/**
 * Similar to {@link #YARNRunner(Configuration)} but allowing injecting
 * {@link ResourceMgrDelegate}. Enables mocking and testing.
 *
 * @param conf the configuration object for the client
 * @param resMgrDelegate the resourcemanager client handle.
 */
public YANRRunner(Configuration conf, ResourceMgrDelegate resMgrDelegate) {
    // 往下追
    this(
        conf,
        resMgrDelegate,
        // 创建客户端缓存 ClientCache
        new ClientCache(conf, resMgrDelegate));
}

```

```

/**
 * Similar to {@link YARNRunner#YARNRunner(Configuration, ResourceMgrDelegate)}
 * but allowing injecting {@link ClientCache}. Enable mocking and testing.
 *
 * @param conf the configuration object
 * @param resMgrDelegate the resource manager delegate
 * @param clientCache the client cache object.
 */
public YARNRunner(Configuration conf, ResourceMgrDelegate resMgrDelegate,
                  ClientCache clientCache) {
    this.conf = conf;
    try {
        // ResourceMgrDelegate
        this.resMgrDelegate = resMgrDelegate;
        // ClientCache
        this.clientCache = clientCache;
        // 获取文件系统上下文 FileContext
        this.defaultFileContext = FileContext.getFileContext(this.conf);
    } catch (UnsupportedFileSystemException ufe) {
        throw new RuntimeException("Error in instantiating YarnClient", ufe);
    }
}

```

3.2 提交任务 Job

```

@InterfaceAudience.Private
@InterfaceStability.Unstable
class JobSubmitter {
    /**
     * Internal method for submitting jobs to the system.
     *
     * <p>The job submission process involves:
     * <ol>
     * <li>
     *     Checking the input and output specifications of the job.
     * </li>
     * <li>
     *     Computing the {@link InputSplit}s for the job.
     * </li>
     * <li>
     *     Setup the requisite accounting information for the
     *     {@link DistributedCache} of the job, if necessary.
     * </li>
     * <li>
     *     Copying the job's jar and configuration to the map-reduce system
     *     directory on the distributed file-system.
     * </li>
     * <li>
     *     Submitting the job to the <code>JobTracker</code> and optionally
     *     monitoring it's status.
     * </li>
     *

```

```

* </ol></p>
*
* @param job      the configuration to submit
* @param cluster the handle to the Cluster
* @throws ClassNotFoundException
* @throws InterruptedException
* @throws IOException
*/
JobStatus submitJobInternal(Job job, Cluster cluster)
    throws ClassNotFoundException, InterruptedException, IOException {

    // validate the jobs output specs
    // 检查任务 Job 输出路径是否存在 如果存在则直接报错
    checkSpecs(job);

    Configuration conf = job.getConfiguration();
    addMRFrameworkToDistributedCache(conf);

    // 获取任务提交临时目录 格式为 /tmp/hadoop-yarn/staging/${user}/.staging
    Path jobStagingArea = JobSubmissionFiles.getStagingDir(cluster, conf);

    // configure the command line options correctly on the submitting dfs
    // 设置当前提交任务节点的主机地址和主机名
    InetAddress ip = InetAddress.getLocalHost();
    if (ip != null) {
        submitHostAddress = ip.getHostAddress();
        submitHostName = ip.getHostName();
        conf.set(MRJobConfig.JOB_SUBMITHOST, submitHostName);
        conf.set(MRJobConfig.JOB_SUBMITHOSTADDR, submitHostAddress);
    }

    // 向 ResourceManager 申请任务 ID
    JobID jobId = submitClient.getNewJobID();
    job.setJobID(jobId);

    // 任务的临时 HDFS 目录 比如 /tmp/hadoop-
    // yarn/staging/tanbs/.staging/job_1684656010852_0002
    Path submitJobDir = new Path(jobStagingArea, jobId.toString());
    JobStatus status = null;
    try {
        conf.set(MRJobConfig.USER_NAME,
            UserGroupInformation.getCurrentUser().getShortUserName());
        conf.set("hadoop.http.filter.initializers",
            "org.apache.hadoop.yarn.server.webproxy.amfilter.AmFilterInitializer");
        // 设置任务的临时目录
        conf.set(MRJobConfig.MAPREDUCE_JOB_DIR, submitJobDir.toString());
        LOG.debug("Configuring job " + jobId + " with " + submitJobDir
            + " as the submit dir");
        // get delegation token for the dir
        TokenCache.obtainTokensForNamenodes(job.getCredentials(),
            new Path[]{submitJobDir}, conf);
    }
}

```



```

populateTokenCache(conf, job.getCredentials());

// generate a secret to authenticate shuffle transfers
if (TokenCache.getShuffleSecretKey(job.getCredentials()) == null) {
    KeyGenerator keyGen;
    try {
        keyGen = KeyGenerator.getInstance(SHUFFLE_KEYGEN_ALGORITHM);
        keyGen.init(SHUFFLE_KEY_LENGTH);
    } catch (NoSuchAlgorithmException e) {
        throw new IOException("Error generating shuffle secret key", e);
    }
    SecretKey shuffleKey = keyGen.generateKey();
    TokenCache.setShuffleSecretKey(shuffleKey.getEncoded(),
        job.getCredentials());
}
if (Cryptoutils.isEncryptedSpillEnabled(conf)) {
    conf.setInt(MRJobConfig.MR_AM_MAX_ATTEMPTS, 1);
    LOG.warn("Max job attempts set to 1 since encrypted intermediate" +
        "data spill is enabled");
}

// 提交任务资源到临时目录 (核心提交 job.jar)
copyAndConfigureFiles(job, submitJobDir);

// 获取 ${任务提交临时目录}/job.xml 文件目录
Path submitJobFile = JobSubmissionFiles.getJobConfPath(submitJobDir);

// Create the splits for the job
LOG.debug("Creating splits at " + jtFs.makeQualified(submitJobDir));

// 获取 map 切片信息 (计算切片)
// 1 创建切片文件 job.split 并上传到 HDFS
// 2 创建切片元数据文件 job.splitmetainfo 并上传切片元数据信息
int maps = writeSplits(job, submitJobDir);
conf.setInt(MRJobConfig.NUM_MAPS, maps);
// 打印切片个数
LOG.info("number of splits:" + maps);

int maxMaps = conf.getInt(MRJobConfig.JOB_MAX_MAP,
    MRJobConfig.DEFAULT_JOB_MAX_MAP);
if (maxMaps >= 0 && maxMaps < maps) {
    throw new IllegalArgumentException("The number of map tasks " + maps
+
        " exceeded limit " + maxMaps);
}

// write "queue admins of the queue to which job is being submitted"
// to job file.
// 获取任务执行队列 默认 default 队列
String queue = conf.get(MRJobConfig.QUEUE_NAME,
    JobConf.DEFAULT_QUEUE_NAME);

```

```

AccessControlList acl = submitClient.getQueueAdmins(queue);
conf.set(toFullPropertyName(queue,
    QueueACL.ADMINISTER_JOBS.getAclName()), acl.getAclString());

// removing jobtoken referrals before copying the jobconf to HDFS
// as the tasks don't need this setting, actually they may break
// because of it if present as the referral will point to a
// different job.
TokenCache.cleanupTokenReferral(conf);

// 默认 false
if (conf.getBoolean(
    MRJobConfig.JOB_TOKEN_TRACKING_IDS_ENABLED,
    MRJobConfig.DEFAULT_JOB_TOKEN_TRACKING_IDS_ENABLED)) {
    // Add HDFS tracking ids
    ArrayList<String> trackingIds = new ArrayList<String>();
    for (Token<? extends TokenIdentifier> t :
        job.getCredentials().getAllTokens()) {
        trackingIds.add(t.decodeIdentifier().getTrackingId());
    }
    conf.setStrings(MRJobConfig.JOB_TOKEN_TRACKING_IDS,
        trackingIds.toArray(new String[trackingIds.size()]));
}

// Set reservation info if it exists
ReservationId reservationId = job.getReservationId();
if (reservationId != null) {
    conf.set(MRJobConfig.RESERVATION_ID, reservationId.toString());
}

// Write job file to submit dir
// 上传任务配置信息文件 job.xml
writeConf(conf, submitJobFile);

// 到这里才是真正提交任务 上面都是一些环境准备 (比如提交任务的相关信息到 HDFS
// 比如
// 1 任务运行包 job.jar
// 2 切片信息 job.split&job.splitmetainfo
// 3 任务运行配置信息 job.xml
// )
// Now, actually submit the job (using the submit name)
// 打印信息
// 1 Submitting tokens for job: job_1684656010852_0002
// 2 Executing with tokens: []
printTokens(jobId, job.getCredentials());

// 真正客户端提交任务到 Yarn (调用 YARNRunner.submitJob())
status = submitClient.submitJob(
    jobId, // 任务 ID
    submitJobDir.toString(), // 任务运行相关文件路径目录
    job.getCredentials() // 任务运行凭证

```

```

    );
    if (status != null) {
        return status;
    } else {
        throw new IOException("Could not launch job");
    }
} finally {
    if (status == null) {
        LOG.info("Cleaning up the staging area " + submitJobDir);
        if (jtFs != null && submitJobDir != null)
            jtFs.delete(submitJobDir, true);
    }
}
}
}
}

```

3.2.1 向 ResourceManager 的 ClientRMService RPC 申请任务 ApplicationId

```

// 发送 RPC 请求给 ClientRMService 服务返回一个 JobID
// 调用 YARNRunner.getNewJobID()
@Override
public JobID getNewJobID() throws IOException, InterruptedException {
    // 向 RM 申请任务 ID
    return resMgrDelegate.getNewJobID();
}

```

```

public JobID getNewJobID() throws IOException, InterruptedException {
    try {
        // 发送 RPC 请求给 RM 请求创建一个任务的 ApplicationId 封装成
        YarnClientApplication
        this.application = client.createApplication()
            .getApplicationSubmissionContext();
        this.applicationId = this.application.getApplicationId();
        return TypeConverter.fromYarn(applicationId);
    } catch (YarnException e) {
        throw new IOException(e);
    }
}
}

```

3.2.2 获取任务临时目录

```

// 该临时目录是任务提交文件上传路径

// 获取任务提交临时目录 格式为 /tmp/hadoop-yarn/staging/${user}/.staging
Path jobStagingArea = JobSubmissionFiles.getStagingDir(cluster, conf);

// 任务的临时 HDFS 目录 比如 /tmp/hadoop-
yarn/staging/tanbs/.staging/job_1684656010852_0002
Path submitJobDir = new Path(jobStagingArea, jobId.toString());

```

3.2.3 提交 job.jar 信息到临时目录

```
/**
 * configure the jobconf of the user with the command line options of
 * -libjars, -files, -archives.
 *
 * @param job
 * @throws IOException
 */
private void copyAndConfigureFiles(Job job, Path jobSubmitDir)
    throws IOException {
    Configuration conf = job.getConfiguration();
    // 默认 true
    boolean usewildcards = conf.getBoolean(Job.USE_WILDCARD_FOR_LIBJARS,
        Job.DEFAULT_USE_WILDCARD_FOR_LIBJARS);
    // 创建任务提交资源 JobResourceUploader
    JobResourceUploader ruploader = new JobResourceUploader(jtFs, usewildcards);

    // 提交任务资源到临时目录 (核心提交 job.jar)
    ruploader.uploadResources(job, jobSubmitDir);

    // Get the working directory. If not set, sets it to filesystem working dir
    // This code has been added so that working directory reset before running
    // the job. This is necessary for backward compatibility as other systems
    // might use the public API JobConf#setWorkingDirectory to reset the working
    // directory.
    job.getWorkingDirectory();
}
```

3.2.4 切片相关信息(job.split, job.splitmetainfo)

```
// 获取 map 切片信息 (计算切片)
// 1 创建切片文件 job.split 并上传到 HDFS
// 2 创建切片元数据文件 job.splitmetainfo 并上传切片元数据信息
int maps = writeSplits(job, submitJobDir);
conf.setInt(MRJobConfig.NUM_MAPS, maps);
```

```
private int writeSplits(org.apache.hadoop.mapreduce.JobContext job,
    Path jobSubmitDir) throws IOException,
    InterruptedException, ClassNotFoundException {
    JobConf jConf = (JobConf) job.getConfiguration();
    int maps;
    // 一般使用新的 Mapper API
    if (jConf.getUseNewMapper()) {
        // 计算 map 切片个数
        maps = writeNewSplits(job, jobSubmitDir);
    } else {
        maps = writeOldSplits(jConf, jobSubmitDir);
    }
    return maps;
}
```

```
}
```

```
@SuppressWarnings("unchecked")
private <T extends InputSplit>
int writeNewSplits(JobContext job, Path jobSubmitDir) throws IOException,
    InterruptedException, ClassNotFoundException {
    Configuration conf = job.getConfiguration();
    // 默认 MR 任务输入 TextInputFormat
    InputFormat<?, ?> input =
        ReflectionUtils.newInstance(job.getInputFormatClass(), conf);

    // 计算 map 切片信息 默认调用 TextInputFormat 的父类 FileInputFormat.getSplits()
    List<InputSplit> splits = input.getSplits(job);
    T[] array = (T[]) splits.toArray(new InputSplit[splits.size()]);

    // sort the splits into order based on size, so that the biggest
    // go first
    Arrays.sort(array, new SplitComparator());
    // 将切片信息上传到 HDFS
    // 1 创建切片文件 job.split 并上传到 HDFS
    // 2 创建切片元数据文件 job.splitmetainfo 并上传切片元数据信息
    JobSplitWriter.createSplitFiles(jobSubmitDir, conf,
        jobSubmitDir.getFileSystem(conf), array);
    // 返回切片个数
    return array.length;
}
```

3.2.4.1 计算切片

```
/**
 * Logically split the set of input files for the job.
 *
 * <p>Each {@link InputSplit} is then assigned to an individual {@link Mapper}
 * for processing.</p>
 *
 * <p><i>Note</i>: The split is a <i>logical</i> split of the inputs and the
 * input files are not physically split into chunks. For e.g. a split could
 * be <i>&lt;input-file-path, start, offset&gt;</i> tuple. The InputFormat
 * also creates the {@link RecordReader} to read the {@link InputSplit}.
 *
 * @param context job configuration.
 * @return an array of {@link InputSplit}s for the job.
 */
public abstract
    List<InputSplit> getSplits(JobContext context
        ) throws IOException, InterruptedException;
// 默认调用 TextInputFormat 的父类 FileInputFormat.getSplits()

/**
 * A base class for file-based {@link InputFormat}s.
 *
```

```

* <p><code>FileInputFormat</code> is the base class for all file-based
* <code>InputFormat</code>s. This provides a generic implementation of
* {@link #getSplits(JobContext)}.
* <p>
* Implementations of <code>FileInputFormat</code> can also override the
* {@link #isSplittable(JobContext, Path)} method to prevent input files
* from being split-up in certain situations. Implementations that may
* deal with non-splittable files must override this method, since
* the default implementation assumes splitting is always possible.
*/
@InterfaceAudience.Public
@InterfaceStability.Stable
public abstract class FileInputFormat<K, V> extends InputFormat<K, V> {
    /**
     * Generate the list of files and make them into FileSplits.
     *
     * @param job the job context
     * @throws IOException
     */
    public List<InputSplit> getSplits(JobContext job) throws IOException {
        // 记录切片计算开始时间
        Stopwatch sw = new Stopwatch().start();

        // 默认 1
        long minSize = Math.max(getFormatMinSplitsize(), getMinSplitsize(job));
        // 默认 Long.MAX_VALUE
        long maxSize = getMaxSplitsize(job);

        // generate splits
        List<InputSplit> splits = new ArrayList<InputSplit>();
        // 获取任务输入的路径 (可能有多个)
        List<FileStatus> files = listStatus(job);

        boolean ignoreDirs = !getInputDirRecursive(job)
            &&
            job.getConfiguration().getBoolean(INPUT_DIR_NONRECURSIVE_IGNORE_SUBDIRS, false);

        // 遍历每个任务输入路径
        for (FileStatus file : files) {
            if (ignoreDirs && file.isDirectory()) {
                continue;
            }
            Path path = file.getPath();
            // 获取文件字节长度
            long length = file.getLen();
            if (length != 0) {
                BlockLocation[] blkLocations;

                // 文件存在 HDFS 可能有多个 block
                if (file instanceof LocatedFileStatus) {
                    blkLocations = ((LocatedFileStatus) file).getBlockLocations();
                } else {

```

```

        FileSystem fs = path.getFileSystem(job.getConfiguration());
        blkLocations = fs.getFileBlockLocations(file, 0, length);
    }

    // 判断文件是否可以切片
    if (isSplittable(job, path)) {
        // 获取文件块大小 默认 128MB
        long blockSize = file.getBlockSize();
        // 计算一个 map 任务处理多少数据量 ( max(minSize,min(maxSize,
        blockSize) )

        long splitSize = computeSplitSize(blockSize, minSize, maxSize);

        // 遍历当前文件夹的切片之后残余数据字节数
        long bytesRemaining = length;
        // 残余数据字节数 / splitSize(默认 128MB) > 1.1
        while (((double) bytesRemaining) / splitSize > SPLIT_SLOP) {
            // 获取文件 block 索引 (第几个 block)
            int blkIndex = getBlockIndex(blkLocations, length -
            bytesRemaining);

            // 封装切片信息
            splits.add(makeSplit(
                path, // 哪个文件路径
                length - bytesRemaining, // 哪个文件读取数据偏移量
                splitSize, // block 大小
                blkLocations[blkIndex].getHosts(), // 哪个 block 存在
                // 在哪些主机上
                blkLocations[blkIndex].getCachedHosts()
            ));
            // 减少切片字节数
            bytesRemaining -= splitSize;
        }

        // 最后是否剩余
        if (bytesRemaining != 0) {
            int blkIndex = getBlockIndex(blkLocations, length -
            bytesRemaining);

            splits.add(makeSplit(path, length - bytesRemaining,
            bytesRemaining,
                blkLocations[blkIndex].getHosts(),
                blkLocations[blkIndex].getCachedHosts()));
        }
    } else { // not splittable
        if (LOG.isDebugEnabled()) {
            // Log only if the file is big enough to be splitted
            if (length > Math.min(file.getBlockSize(), minSize)) {
                LOG.debug("File is not splittable so no parallelization
                "
                    + "is possible: " + file.getPath());
            }
        }
        splits.add(makeSplit(path, 0, length,
        blkLocations[0].getHosts(),

```

```

        blkLocations[0].getCachedHosts()));
    }
} else {
    //Create empty hosts array for zero length files
    splits.add(makeSplit(path, 0, length, new String[0]));
}
}
// Save the number of input files for metrics/loadgen
job.getConfiguration().setLong(NUM_INPUT_FILES, files.size());
sw.stop();
if (LOG.isDebugEnabled()) {
    LOG.debug("Total # of splits generated by getSplits: " + splits.size()
        + ", TimeTaken: " + sw.now(TimeUnit.MILLISECONDS));
}
return splits;
}
}

```

3.2.4.2 上传切片信息到临时目录

```

// 1 创建切片文件 job.split 并上传到 HDFS
// 2 创建切片元数据文件 job.splitmetainfo 并上传切片元数据信息

/**
 * The class that is used by the Job clients to write splits (both the meta
 * and the raw bytes parts)
 */
@InterfaceAudience.Private
@InterfaceStability.Unstable
public class JobSplitWriter {
    public static <T extends InputSplit> void createSplitFiles(Path jobSubmitDir,
        Configuration conf,
        FileSystem fs, T[] splits)
        throws IOException, InterruptedException {
        // 创建切片文件 job.split 并上传到 HDFS
        FSDataOutputStream out = createFile(fs,
            JobSubmissionFiles.getJobSplitFile(jobSubmitDir), conf);
        SplitMetaInfo[] info = writeNewSplits(conf, splits, out);
        out.close();

        // 创建切片元数据文件 job.splitmetainfo 并上传切片元数据信息
        writeJobSplitMetaInfo(fs,
            JobSubmissionFiles.getJobSplitMetaFile(jobSubmitDir),
            new FsPermission(JobSubmissionFiles.JOB_FILE_PERMISSION),
            splitVersion,
            info);
    }
}

```


3.2.5 上传任务配置信息文件 job.xml

```
// 获取 ${任务提交临时目录}/job.xml 文件目录
Path submitJobFile = JobSubmissionFiles.getJobConfPath(submitJobDir);
.....
// write job file to submit dir
// 上传任务配置信息文件 job.xml
writeConf(conf, submitJobFile);
```

```
private void writeConf(Configuration conf, Path jobFile)
    throws IOException {
    // write job file to JobTracker's fs
    FSDataOutputStream out =
        FileSystem.create(jtFs, jobFile,
            new FsPermission(JobSubmissionFiles.JOB_FILE_PERMISSION));
    try {
        // 写任务配置信息 job.xml
        conf.writeXml(out);
    } finally {
        out.close();
    }
}
```

3.2.6 RPC 客户端提交任务到 Yarn (调用 YARNRunner.submitJob())

```
// RPC 客户端提交任务到 Yarn (调用 YARNRunner.submitJob())
status = submitClient.submitJob(
    jobId, // 任务 ID
    submitJobDir.toString(), // 任务运行相关文件路径目录
    job.getCredentials() // 任务运行凭证
);
```

```
@Override
public JobStatus submitJob(JobID jobId, String jobSubmitDir, Credentials ts)
    throws IOException, InterruptedException {

    addHistoryToken(ts);

    // 将提交任务信息 (任务配置文件路径) 封装成 ApplicationSubmissionContext
    ApplicationSubmissionContext appContext =
        createApplicationSubmissionContext(conf, jobSubmitDir, ts);

    // Submit to ResourceManager
    try {
        // 提交任务 (发送 RPC 请求)
        ApplicationId applicationId =
            resMgrDelegate.submitApplication(appContext);

        // 获取任务执行信息
```

```

        ApplicationReport appMaster = resMgrDelegate
            .getApplicationReport(applicationId);
        String diagnostics =
            (appMaster == null ?
                "application report is null" :
appMaster.getDiagnostics());
        if (appMaster == null
            || appMaster.getYarnApplicationState() ==
YarnApplicationState.FAILED
            || appMaster.getYarnApplicationState() ==
YarnApplicationState.KILLED) {
            throw new IOException("Failed to run job : " +
                diagnostics);
        }
        return clientCache.getClient(jobId).getJobStatus(jobId);
    } catch (YarnException e) {
        throw new IOException(e);
    }
}

```

3.2.6.1 提交任务信息封装成 ApplicationSubmissionContext (启动 MRAppMaster)

```

/**
 * Constructs all the necessary information to start the MR AM.
 *
 * @param jobConf      the configuration for the MR job
 * @param jobSubmitDir the directory path for the job
 * @param ts           the security credentials for the job
 * @return ApplicationSubmissionContext
 * @throws IOException on IO error (e.g. path resolution)
 */
public ApplicationSubmissionContext createApplicationSubmissionContext(
    Configuration jobConf, String jobSubmitDir, Credentials ts)
    throws IOException {
    // 获取任务 ApplicationId
    ApplicationId applicationId = resMgrDelegate.getApplicationId();

    // Setup LocalResources
    // 设置本地资源配置到任务
    // 1 job.xml -> LocalResource
    // 2 job.jar -> LocalResource
    // 3 jobSubmitDir/job.split -> LocalResource
    // 4 jobSubmitDir/job.splitmetainfo -> LocalResource
    Map<String, LocalResource> localResources =
        setupLocalResources(jobConf, jobSubmitDir);

    // Setup security tokens
    DataOutputStream dob = new DataOutputStream();
    ts.writeTokenStorageToStream(dob);
    ByteBuffer securityTokens =
        ByteBuffer.wrap(dob.getData(), 0, dob.getLength());
}

```

```

// Setup ContainerLaunchContext for AM container
// 构建启动 AM 命令
/**
 * 基本启动 AM 命令模版
 * ${JAVA_HOME}/bin/java
 * -Djava.io.tmpdir=./tmp
 * -Dlog4j.configuration=container-log4j.properties
 * -Dyarn.app.container.log.dir=<LOG_DIR>
 * -Dyarn.app.container.log.filesize=0
 * -Dhadoop.root.logger=INFO,CLA
 * -Dhadoop.root.logfile=syslog
 * -Xmx1024m
 * org.apache.hadoop.mapreduce.v2.app.MRAppMaster
 * 1><LOG_DIR>.stdout
 * 2><LOG_DIR>.stderr
 */
List<String> vargs = setupAMCommand(jobConf);
// 创建启动 AM 容器上下文 (也即拼接最终启动 AM 容器命令以及启动环境变量)
ContainerLaunchContext amContainer = setupContainerLaunchContextForAM(
    jobConf, localResources, securityTokens, vargs);

String regex = conf.get(MRJobConfig.MR_JOB_SEND_TOKEN_CONF);
if (regex != null && !regex.isEmpty()) {
    setTokenRenewerConf(amContainer, conf, regex);
}

Collection<String> tagsFromConf =
    jobConf.getTrimmedStringCollection(MRJobConfig.JOB_TAGS);

// Set up the ApplicationSubmissionContext
// 创建任务提交上下文对象 ApplicationSubmissionContext
ApplicationSubmissionContext appContext =
    recordFactory.newRecordInstance(ApplicationSubmissionContext.class);
// 任务提交上下文对象设置任务应用 ID
appContext.setApplicationId(applicationId); // ApplicationId
// 设置队列 默认 default
appContext.setQueue( // Queue name
    jobConf.get(JobContext.QUEUE_NAME,
        YarnConfiguration.DEFAULT_QUEUE_NAME));

// add reservationID if present
// 默认 null
ReservationId reservationID = null;
try {
    reservationID =
        ReservationId.parseReservationId(jobConf
            .get(JobContext.RESERVATION_ID));
} catch (NumberFormatException e) {
    // throw exception as reservationid as is invalid
    String errMsg =
        "Invalid reservationId: " +
        jobConf.get(JobContext.RESERVATION_ID)

```

```

        + " specified for the app: " + applicationId;
    LOG.warn(errMsg);
    throw new IOException(errMsg);
}
if (reservationID != null) {
    appContext.setReservationID(reservationID);
    LOG.info("SUBMITTING ApplicationSubmissionContext app:" + applicationId
        + " to queue:" + appContext.getQueue() + " with reservationId:"
        + appContext.getReservationID());
}

// 设置任务名称 (Job.getInstance(conf, "word count") word count)
appContext.setApplicationName( // Job name
    jobConf.get(JobContext.JOB_NAME,
        YarnConfiguration.DEFAULT_APPLICATION_NAME));
appContext.setCancelTokensWhenComplete(
    conf.getBoolean(MRJobConfig.JOB_CANCEL_DELEGATION_TOKEN, true));
// 设置 AM 启动容器
appContext.setAMContainerSpec(amContainer); // AM Container
// 设置最大重启次数 默认 2
appContext.setMaxAppAttempts(
    conf.getInt(MRJobConfig.MR_AM_MAX_ATTEMPTS,
        MRJobConfig.DEFAULT_MR_AM_MAX_ATTEMPTS));

// Setup the AM ResourceRequests
// 生成启动 AM 容器资源请求 (默认内存 1536 MB 1 CPU *)
List<ResourceRequest> amResourceRequests = generateResourceRequests();
appContext.setAMContainerResourceRequests(amResourceRequests);

// set labels for the AM container requests if present
String amNodeLabelExpression = conf.get(MRJobConfig.AM_NODE_LABEL_EXP);
if (null != amNodeLabelExpression
    && amNodeLabelExpression.trim().length() != 0) {
    for (ResourceRequest amResourceRequest : amResourceRequests) {
        amResourceRequest.setNodeLabelExpression(amNodeLabelExpression.trim());
    }
}
// set labels for the Job containers
appContext.setNodeLabelExpression(jobConf
    .get(JobContext.JOB_NODE_LABEL_EXP));

// 设置任务类型为 MAPREDUCE
appContext.setApplicationType(MRJobConfig.MR_APPLICATION_TYPE);
if (tagsFromConf != null && !tagsFromConf.isEmpty()) {
    appContext.setApplicationTags(new HashSet<>(tagsFromConf));
}

String jobPriority = jobConf.get(MRJobConfig.PRIORITY);
if (jobPriority != null) {
    int iPriority;
    try {

```

```

        iPriority = TypeConverter.toYarnApplicationPriority(jobPriority);
    } catch (IllegalArgumentException e) {
        iPriority = Integer.parseInt(jobPriority);
    }
    appContext.setPriority(Priority.newInstance(iPriority));
}

// 返回 ApplicationSubmissionContext
return appContext;
}

```

3.2.6.2 提交任务(调用 YarnClientImpl.submitApplication())

```

@Override
public ApplicationId
submitApplication(ApplicationSubmissionContext appContext)
    throws YarnException, IOException {
    // 提交任务到 YARN (调用 YarnClientImpl.submitApplication())
    return client.submitApplication(appContext);
}

```

```

@Override
public ApplicationId
submitApplication(ApplicationSubmissionContext appContext)
    throws YarnException, IOException {
    // 判断提交任务的 ApplicationId 是否为空
    ApplicationId applicationId = appContext.getApplicationId();
    if (applicationId == null) {
        throw new ApplicationIdNotProvidedException(
            "ApplicationId is not provided in
ApplicationSubmissionContext");
    }
    // 封装任务提交上下文对象为 RPC 请求对象
    SubmitApplicationRequest request =
        Records.newRecord(SubmitApplicationRequest.class);
    request.setApplicationSubmissionContext(appContext);

    // Automatically add the timeline DT into the CLC
    // Only when the security and the timeline service are both enabled
    if (isSecurityEnabled() && timelineV1ServiceEnabled) {
        addTimelineDelegationToken(appContext.getAMContainerSpec());
    }

    //TODO: YARN-1763:Handle RM failovers during the submitApplication call.
    // 提交任务 (调用 ClientRMService.submitApplication())
    rmClient.submitApplication(request);

    int pollCount = 0;
    long startTime = System.currentTimeMillis();
    EnumSet<YarnApplicationState> waitingStates =
        EnumSet.of(YarnApplicationState.NEW,

```

```

        YarnApplicationState.NEW_SAVING,
        YarnApplicationState.SUBMITTED);
EnumSet<YarnApplicationState> failToSubmitStates =
    EnumSet.of(YarnApplicationState.FAILED,
        YarnApplicationState.KILLED);
while (true) {
    try {
        // 持续循环判断提交的任务是否成功
        ApplicationReport appReport = getApplicationReport(applicationId);
        YarnApplicationState state = appReport.getYarnApplicationState();
        if (!waitingStates.contains(state)) {
            if (failToSubmitStates.contains(state)) {
                throw new YarnException("Failed to submit " + applicationId
+
                    " to YARN : " + appReport.getDiagnostics());
            }
            // Submitted application application_1684656010852_0002
            LOG.info("Submitted application " + applicationId);
            break;
        }

        long elapsedMillis = System.currentTimeMillis() - startTime;
        if (enforceAsyncAPITimeout() &&
            elapsedMillis >= asyncApiPollTimeoutMillis) {
            throw new YarnException("Timed out while waiting for application
" +
                applicationId + " to be submitted successfully");
        }

        // Notify the client through the log every 10 poll, in case the
client
        // is blocked here too long.
        if (++pollCount % 10 == 0) {
            LOG.info("Application submission is not finished, " +
                "submitted application " + applicationId +
                " is still in " + state);
        }
        try {
            Thread.sleep(submitPollIntervalMillis);
        } catch (InterruptedException ie) {
            String msg = "Interrupted while waiting for application "
                + applicationId + " to be successfully submitted.";
            LOG.error(msg);
            throw new YarnException(msg, ie);
        }
    } catch (ApplicationNotFoundException ex) {
        // Failover or RM restart happens before RMStateStore saves
        // ApplicationState
        LOG.info("Re-submit application " + applicationId + "with the " +
            "same ApplicationSubmissionContext");
        rmClient.submitApplication(request);
    }
}

```

```

    }

    return applicationId;
}

```

```

@Override
public ApplicationId
submitApplication(ApplicationSubmissionContext appContext)
    throws YarnException, IOException {
    // 判断提交任务的 ApplicationId 是否为空
    ApplicationId applicationId = appContext.getApplicationId();
    if (applicationId == null) {
        throw new ApplicationIdNotProvidedException(
            "ApplicationId is not provided in
ApplicationSubmissionContext");
    }
    // 封装任务提交上下文对象为 RPC 请求对象
    SubmitApplicationRequest request =
        Records.newRecord(SubmitApplicationRequest.class);
    request.setApplicationSubmissionContext(appContext);

    // Automatically add the timeline DT into the CLC
    // Only when the security and the timeline service are both enabled
    if (isSecurityEnabled() && timelineV1ServiceEnabled) {
        addTimelineDelegationToken(appContext.getAMContainerSpec());
    }

    //TODO: YARN-1763:Handle RM failovers during the submitApplication call.
    // 提交任务 (调用 ClientRMService.submitApplication())
    rmClient.submitApplication(request);

    int pollCount = 0;
    long startTime = System.currentTimeMillis();
    EnumSet<YarnApplicationState> waitingStates =
        EnumSet.of(YarnApplicationState.NEW,
            YarnApplicationState.NEW_SAVING,
            YarnApplicationState.SUBMITTED);
    EnumSet<YarnApplicationState> failToSubmitStates =
        EnumSet.of(YarnApplicationState.FAILED,
            YarnApplicationState.KILLED);
    while (true) {
        try {
            // 持续循环判断提交的任务是否成功
            ApplicationReport appReport = getApplicationReport(applicationId);
            YarnApplicationState state = appReport.getYarnApplicationState();
            if (!waitingStates.contains(state)) {
                if (failToSubmitStates.contains(state)) {
                    throw new YarnException("Failed to submit " + applicationId
+
                                " to YARN : " + appReport.getDiagnostics());
                }
            }
        }
    }
}

```

```

        // Submitted application application_1684656010852_0002
        LOG.info("Submitted application " + applicationId);
        break;
    }

    long elapsedMillis = System.currentTimeMillis() - startTime;
    if (enforceAsyncAPITimeout() &&
        elapsedMillis >= asyncApiPollTimeoutMillis) {
        throw new YarnException("Timed out while waiting for application
" +
                                applicationId + " to be submitted successfully");
    }

    // Notify the client through the log every 10 poll, in case the
client
    // is blocked here too long.
    if (++pollCount % 10 == 0) {
        LOG.info("Application submission is not finished, " +
            "submitted application " + applicationId +
            " is still in " + state);
    }
    try {
        Thread.sleep(submitPollIntervalMillis);
    } catch (InterruptedException ie) {
        String msg = "Interrupted while waiting for application "
            + applicationId + " to be successfully submitted.";
        LOG.error(msg);
        throw new YarnException(msg, ie);
    }
} catch (ApplicationNotFoundException ex) {
    // FailOver or RM restart happens before RMStateStore saves
    // ApplicationState
    LOG.info("Re-submit application " + applicationId + "with the " +
        "same ApplicationSubmissionContext");
    rmClient.submitApplication(request);
}

}

return applicationId;
}

```

四 ResourceManager 的 ClientRMService 接收到客户端发送的 RPC 请求 (启动 AM)

```

// 提交任务 (调用 ClientRMService.submitApplication())
@Override
public SubmitApplicationResponse submitApplication(
    SubmitApplicationRequest request) throws YarnException, IOException {
    // RPC 请求信息
    ApplicationSubmissionContext submissionContext = request

```



```

        .getApplicationSubmissionContext();
ApplicationId applicationId = submissionContext.getApplicationId();
CallerContext callerContext = CallerContext.getCurrent();

// ApplicationSubmissionContext needs to be validated for safety - only
// those fields that are independent of the RM's configuration will be
// checked here, those that are dependent on RM configuration are validated
// in RMAAppManager.

String user = null;
try {
    // safety
    user = UserGroupInformation.getCurrentUser().getShortUserName();
} catch (IOException ie) {
    LOG.warn("Unable to get the current user.", ie);
    RMAuditLogger.logFailure(user, AuditConstants.SUBMIT_APP_REQUEST,
        ie.getMessage(), "ClientRMService",
        "Exception in submitting application", applicationId,
callerContext,
        submissionContext.getQueue());
    throw RPCUtil.getRemoteException(ie);
}

if (timelineServiceV2Enabled) {
    // Sanity check for flow run
    String value = null;
    try {
        for (String tag : submissionContext.getApplicationTags()) {
            if (tag.startsWith(TimelineUtils.FLOW_RUN_ID_TAG_PREFIX + ":"))
||
                tag.startsWith(

TimelineUtils.FLOW_RUN_ID_TAG_PREFIX.toLowerCase() + ":")) {
                value =
tag.substring(TimelineUtils.FLOW_RUN_ID_TAG_PREFIX.length()
                    + 1);
                // In order to check the number format
                Long.valueOf(value);
            }
        }
    } catch (NumberFormatException e) {
        LOG.warn("Invalid to flow run: " + value +
            ". Flow run should be a long integer", e);
        RMAuditLogger.logFailure(user, AuditConstants.SUBMIT_APP_REQUEST,
            e.getMessage(), "ClientRMService",
            "Exception in submitting application", applicationId,
            submissionContext.getQueue());
        throw RPCUtil.getRemoteException(e);
    }
}

// Check whether app has already been put into rmContext,

```

```

// If it is, simply return the response
if (rmContext.getRMApps().get(applicationId) != null) {
    LOG.info("This is an earlier submitted application: " + applicationId);
    return SubmitApplicationResponse.newInstance();
}

ByteBuffer tokenConf =
    submissionContext.getAMContainerSpec()
        .getTokensConf();
if (tokenConf != null) {
    int maxSize = getConfig()
        .getInt(YarnConfiguration.RM_DELEGATION_TOKEN_MAX_CONF_SIZE,
YarnConfiguration.DEFAULT_RM_DELEGATION_TOKEN_MAX_CONF_SIZE_BYTES);
    LOG.info("Using app provided configurations for delegation token
renewal,"
        + " total size = " + tokenConf.capacity());
    if (tokenConf.capacity() > maxSize) {
        throw new YarnException(
            "Exceed " +
YarnConfiguration.RM_DELEGATION_TOKEN_MAX_CONF_SIZE
            + " = " + maxSize + " bytes, current conf size = "
            + tokenConf.capacity() + " bytes.");
    }
}

if (submissionContext.getQueue() == null) {
    submissionContext.setQueue(YarnConfiguration.DEFAULT_QUEUE_NAME);
}
if (submissionContext.getApplicationName() == null) {
    submissionContext.setApplicationName(
        YarnConfiguration.DEFAULT_APPLICATION_NAME);
}

if (submissionContext.getApplicationType() == null) {
    submissionContext
        .setApplicationType(YarnConfiguration.DEFAULT_APPLICATION_TYPE);
} else {
    if (submissionContext.getApplicationType().length() >
YarnConfiguration.APPLICATION_TYPE_LENGTH) {
        submissionContext.setApplicationType(submissionContext
            .getApplicationType().substring(0,
                YarnConfiguration.APPLICATION_TYPE_LENGTH));
    }
}

ReservationId reservationId = request.getApplicationSubmissionContext()
    .getReservationID();
checkReservationACLs(submissionContext.getQueue(), AuditConstants
    .SUBMIT_RESERVATION_REQUEST, reservationId);

try {

```

```

// call RMAppManager to submit application directly
// 直接提交任务到 RMAppManager
rmAppManager.submitApplication(
    submissionContext,
    System.currentTimeMillis(),
    user);

// Application with id 5 submitted by user tanbs
LOG.info("Application with id " + applicationId.getId() +
    " submitted by user " + user);
RMAuditLogger.logSuccess(user, AuditConstants.SUBMIT_APP_REQUEST,
    "ClientRMService", applicationId, callerContext,
    submissionContext.getQueue());
} catch (YarnException e) {
    LOG.info("Exception in submitting " + applicationId, e);
    RMAuditLogger.logFailure(user, AuditConstants.SUBMIT_APP_REQUEST,
        e.getMessage(), "ClientRMService",
        "Exception in submitting application", applicationId,
        callerContext,
        submissionContext.getQueue());
    throw e;
}

return recordFactory
    .newRecordInstance(SubmitApplicationResponse.class);
}

```

4.1 备注说明

// 由于提交启动 MRAppMaster RPC 请求涉及很多的状态机转换 故文档不暂时 需要的话请阅读源码进行理解
 // 这里直接来到准备启动 MRAppMaster 代码 (注意的是 此时启动 AM 已经被分配了容器 Container)

4.2 启动 AM (MRAppMaster) 调用 ApplicationMasterLauncher.handle()

```

public class ApplicationMasterLauncher extends AbstractService implements
    EventHandler<AMLancherEvent> {
    @Override
    public synchronized void handle(AMLancherEvent appEvent) {

        // event = AMLancherEvent
        // eventType = AMLancherEventType.LAUNCH

        AMLancherEventType event = appEvent.getType();
        RMAppAttempt application = appEvent.getAppAttempt();
        switch (event) {
            case LAUNCH:
                // 启动 AM
                launch(application);
                break;

```

```

        case CLEANUP:
            cleanup(application);
            break;
        default:
            break;
    }
}
}

```

```

private void launch(RMAppAttempt application) {

    // 创建 AMLauncher
    Runnable launcher = createRunnableLauncher(application,
        AMLauncherEventType.LAUNCH);
    // 将启动 AM 的 AMLauncher 任务存放到阻塞队列等待 LauncherThread 线程处理
    // 调用 LauncherThread.run()
    masterEvents.add(launcher);
}

```

```

private class LauncherThread extends Thread {

    public LauncherThread() {
        super("ApplicationMaster Launcher");
    }

    @Override
    public void run() {
        while (!this.isInterrupted()) {
            Runnable toLaunch;
            try {
                // 拉取需要启动 ApplicationMaster
                toLaunch = masterEvents.take();
                // 执行启动 AM 调用 AMLauncher.run()
                launcherPool.execute(toLaunch);
            } catch (InterruptedException e) {
                LOG.warn(this.getClass().getName() + " interrupted.
Returning.");
                return;
            }
        }
    }
}

```

4.2.1 执行启动 AM 调用 AMLauncher.run()

```

/**
 * The launch of the AM itself.
 */
public class AMLauncher implements Runnable {
    @SuppressWarnings("unchecked")

```

```

public void run() {
    switch (eventType) {
        // 启动 AM 事件
        case LAUNCH:
            try {
                // Launching masterattempt_1684656010852_0005_000001
                LOG.info("Launching master" + application.getAppAttemptId());
                // 启动 AM
                launch();
                // 调用
                RMAppAttemptEventDispatcher.handle(RMAppAttemptEventType.LAUNCHED)
                    handler.handle(new
RMAppAttemptEvent(application.getAppAttemptId(),
                        RMAppAttemptEventType.LAUNCHED,
                        System.currentTimeMillis()));
            } catch (Exception ie) {
                onAMLaunchFailed(masterContainer.getId(), ie);
            }
            break;
        case CLEANUP:
            try {
                LOG.info("Cleaning master " + application.getAppAttemptId());
                cleanup();
            } catch (IOException ie) {
                LOG.info("Error cleaning master ", ie);
            } catch (YarnException e) {
                StringBuilder sb = new StringBuilder("Container ");
                sb.append(masterContainer.getId().toString());
                sb.append(" is not handled by this NodeManager");
                if (!e.getMessage().contains(sb.toString())) {
                    // Ignoring if container is already killed by Node Manager.
                    LOG.info("Error cleaning master ", e);
                }
            }
            break;
        default:
            LOG.warn("Received unknown event-type " + eventType + ".
Ignoring.");
            break;
    }
}
}

```

```

private void launch() throws IOException, YarnException {
    // 根据 AM 分配的 NodeId( NodeManager 的 RPC 服务 ContainerManagerImpl)
    // 获取 ContainerManagerImpl 的 RPC 客户端代理
    // 通讯协议接口为 ContainerManagementProtocol
    connect();
    // 获取启动 AM 的容器信息
    ContainerId masterContainerID = masterContainer.getId();
    // 获取启动 AM 的上下文信息
}

```

```

ApplicationSubmissionContext applicationContext =
    application.getSubmissionContext();
// Setting up container Container:
// [ContainerId: container_1684656010852_0005_01_000001,
// AllocationRequestId: -1,
// Version: 0,
// NodeId: hadoop104:35189,
// NodeHttpAddress: hadoop104:8042,
// Resource: <memory:1536, vCores:1>,
// Priority: 0,
// Token: Token { kind: ContainerToken, service: 192.168.6.104:35189 },
// ExecutionType: GUARANTEED, ]
// for AM appattempt_1684656010852_0005_000001
LOG.info("Setting up container " + masterContainer
    + " for AM " + application.getAppAttemptId());

// 创建启动容器的上下文 ContainerLaunchContext
// (也即封装启动 AM 的上下文、容器信息、其他 Token、以及环境变量相关等等)
ContainerLaunchContext launchContext =
    createAMContainerLaunchContext(applicationContext,
masterContainerID);

// 创建启动容器请求体 (封装启动容器的上下文 ContainerLaunchContext )
StartContainerRequest scRequest =
    StartContainerRequest.newInstance(launchContext,
        masterContainer.getContainerToken());
List<StartContainerRequest> list = new ArrayList<StartContainerRequest>();
list.add(scRequest);
StartContainersRequest allRequests =
    StartContainersRequest.newInstance(list);

// 发送 RPC 请求到 NodeManager 的 ContainerManagerImpl RPC 服务
// 调用 ContainerManagerImpl.startContainers()
StartContainersResponse response =
    containerMgrProxy.startContainers(allRequests);
if (response.getFailedRequests() != null
    && response.getFailedRequests().containsKey(masterContainerID)) {
    Throwable t =

response.getFailedRequests().get(masterContainerID).deSerialize();
    parseAndThrowException(t);
} else {
    // Done launching container Container:
    // [ContainerId: container_1684656010852_0005_01_000001,
    // AllocationRequestId: -1,
    // Version: 0,
    // NodeId: hadoop104:35189,
    // NodeHttpAddress: hadoop104:8042,
    // Resource: <memory:1536, vCores:1>,
    // Priority: 0,
    // Token: Token { kind: ContainerToken, service: 192.168.6.104:35189 },
    // ExecutionType: GUARANTEED, ]

```

```

        // for AM appattempt_1684656010852_0005_000001
        LOG.info("Done launching container " + masterContainer + " for AM "
            + application.getAppAttemptId());
    }
}

```

4.2.1.1 获取 ContainerManagerImpl 的 RPC 客户端代理

```

private void connect() throws IOException {
    // 获取启动 AM 的容器信息 (也即 NodeManager 信息)
    ContainerId masterContainerID = masterContainer.getId();
    // 创建 NodeManager 服务的 ContainerManagerImpl RPC Server 的客户端代理
    // 通讯协议接口为 ContainerManagementProtocol
    containerMgrProxy = getContainerMgrProxy(masterContainerID);
}

```

4.2.1.2 发送 RPC 请求到 NodeManager 的 ContainerManagerImpl RPC 服务启动容器

```

// 调用 ContainerManagerImpl.startContainers()
/**
 * start a list of containers on this NodeManager.
 */
@Override
public StartContainersResponse startContainers(
    StartContainersRequest requests) throws YarnException, IOException {
    // 获取 RPC 客户端用户组信息并执行认证相关工作
    UserGroupInformation remoteUgi = getRemoteUgi();
    NMTokenIdentifier nmTokenIdentifier = selectNMTokenIdentifier(remoteUgi);
    authorizeUser(remoteUgi, nmTokenIdentifier);

    List<ContainerId> succeededContainers = new ArrayList<ContainerId>();
    Map<ContainerId, SerializedException> failedContainers =
        new HashMap<ContainerId, SerializedException>();
    // Synchronize with NodeStatusUpdaterImpl#registerWithRM
    // to avoid race condition during NM-RM resync (due to RM restart) while a
    // container is being started, in particular when the container has not yet
    // been added to the containers map in NMContext.
    synchronized (this.context) {
        // 针对启动 AM 而言 StartContainerRequest 只有一个
        for (StartContainerRequest request : requests
            .getStartContainerRequests()) {
            ContainerId containerId = null;
            try {
                if (request.getContainerToken() == null
                    || request.getContainerToken().getIdentifier() == null)
                {
                    throw new IOException(INVALID_CONTAINERTOKEN_MSG);
                }

                ContainerTokenIdentifier containerTokenIdentifier = BuilderUtils

```

```

.newContainerTokenIdentifier(request.getContainerToken());

verifyAndGetContainerTokenIdentifier(request.getContainerToken(),
    containerTokenIdentifier);

    // 获取容器信息
    containerId = containerTokenIdentifier.getContainerID();

    // Initialize the AMRMPProxy service instance only if the
container is of
    // type AM and if the AMRMPProxy service is enabled
    if (amrmProxyEnabled &&
containerTokenIdentifier.getContainerType()
        .equals(ContainerType.APPLICATION_MASTER)) {

this.getAMRMPProxyService().processApplicationStartRequest(request);
    }
    performContainerPreStartChecks(nmTokenIdentifier, request,
        containerTokenIdentifier);

    // 启动容器
    startContainerInternal(containerTokenIdentifier, request);
    // 表示启动容器成功
    succeededContainers.add(containerId);
} catch (YarnException e) {
    failedContainers.put(containerId,
SerializedException.newInstance(e));
} catch (InvalidToken ie) {
    failedContainers
        .put(containerId, SerializedException.newInstance(ie));
    throw ie;
} catch (IOException e) {
    throw RPCUtil.getRemoteException(e);
}
}
// 返回 RPC 请求响应
return StartContainersResponse
    .newInstance(getAuxServiceMetaData(), succeededContainers,
        failedContainers);
}
}

```

```

@SuppressWarnings("unchecked")
protected void startContainerInternal(
    ContainerTokenIdentifier containerTokenIdentifier,
    StartContainerRequest request) throws YarnException, IOException {

    ContainerId containerId = containerTokenIdentifier.getContainerID();
    String containerIdStr = containerId.toString();
    String user = containerTokenIdentifier.getApplicationSubmitter();

```



```

// start request for container_1684656010852_0005_01_000001 by user tanbs
LOG.info("Start request for " + containerIdStr + " by user " + user);

// 获取启动容器上下文
ContainerLaunchContext launchContext = request.getContainerLaunchContext();

// Sanity check for local resources
// 对本地资源进行完整性检查
for (Map.Entry<String, LocalResource> rsrc : launchContext
    .getLocalResources().entrySet()) {
    if (rsrc.getValue() == null || rsrc.getValue().getResource() == null) {
        throw new YarnException(
            "Null resource URL for local resource " + rsrc.getKey() + "
: " + rsrc.getValue());
    } else if (rsrc.getValue().getType() == null) {
        throw new YarnException(
            "Null resource type for local resource " + rsrc.getKey() + "
: " + rsrc.getValue());
    } else if (rsrc.getValue().getVisibility() == null) {
        throw new YarnException(
            "Null resource visibility for local resource " +
rsrc.getKey() + " : " + rsrc.getValue());
    }
}

Credentials credentials =
    YarnServerSecurityUtils.parseCredentials(launchContext);

long containerStartTime = SystemClock.getInstance().getTime();
// 创建 ContainerImpl (再次封装容器)
Container container =
    new ContainerImpl(
        getConfig(),
        this.dispatcher,
        launchContext,
        credentials,
        metrics,
        containerTokenIdentifier,
        context,
        containerStartTime);
// 获取启动容器的任务 ApplicationId
ApplicationId applicationID =
    containerId.getApplicationAttemptId().getApplicationId();
// NodeManager 的上下文对象缓存容器 (ContainerId -> ContainerImpl)
if (context.getContainers().putIfAbsent(containerId, container) != null) {
    NMAuditLogger.logFailure(user, AuditConstants.START_CONTAINER,
        "ContainerManagerImpl", "Container already running on this
node!",
        applicationID, containerId);
    throw RPCUtil.getRemoteException("Container " + containerIdStr
        + " already is running on this node!!");
}

```

```

    }

    this.readLock.lock();
    try {
        if (!isServiceStopped()) {
            // 判断当前 NodeManager 是否已经启动过 ApplicationId 对应的任务
            // 对于启动 AM 而言 还没有启动 ApplicationId 对应的任务
            if (!context.getApplications().containsKey(applicationID)) {
                // Create the application
                // populate the flow context from the launch context if the
                timeline
                // service v.2 is enabled
                FlowContext flowContext =
                    getFlowContext(launchContext, applicationID);

                // 创建 ApplicationImpl
                Application application =
                    new ApplicationImpl(dispatcher, user, flowContext,
                        applicationID, credentials, context);
                if (context.getApplications().putIfAbsent(applicationID,
                    application) == null) {
                    // Creating a new application reference for app
                    // application_1684656010852_0005
                    LOG.info("Creating a new application reference for app "
                        + applicationID);
                    LogAggregationContext logAggregationContext =
                        containerTokenIdentifier.getLogAggregationContext();
                    Map<ApplicationAccessType, String> appAcls =
                        container.getLaunchContext().getApplicationACLs();
                    // 存储 ApplicationId (ApplicationId ->
                    ContainerManagerApplicationProto)
                    context.getNMStateStore().storeApplication(applicationID,
                        buildAppProto(applicationID, user, credentials,
                        appAcls,
                        logAggregationContext, flowContext));

                    // 调用
                    ApplicationEventDispatcher.handle(ApplicationEventType.INIT_APPLICATION)
                    // .....
                    // 最终调用 ContainerScheduler.handle(SCHEDULE_CONTAINER) 将启
                    动容器放入缓存队列

                    // 判断是否启动容器 (强制启动或者资源充足)
                    dispatcher.getEventHandler()
                        .handle(
                            // 创建
                            ApplicationInitEvent(ApplicationEventType.INIT_APPLICATION)
                                new ApplicationInitEvent(applicationID,
                                    appAcls, logAggregationContext));
                }
            } else if (containerTokenIdentifier.getContainerType()
                == ContainerType.APPLICATION_MASTER) {
                FlowContext flowContext =
                    getFlowContext(launchContext, applicationID);

```

```

        if (flowContext != null) {
            ApplicationImpl application =
                (ApplicationImpl)
context.getApplications().get(applicationID);

            // update flowContext reference in ApplicationImpl
            application.setFlowContext(flowContext);

            // Required to update state store for recovery.
            context.getNMStateStore().storeApplication(applicationID,
                buildAppProto(applicationID, user, credentials,

container.getLaunchContext().getApplicationACLs(),

containerTokenIdentifier.getLogAggregationContext(),
                flowContext));

            LOG.info(
                "Updated application reference with flowContext " +
flowContext
                + " for app " + applicationID);
        } else {
            LOG.info("TimelineService v2.0 is not enabled. Skipping
updating "
                + "flowContext for application " + applicationID);
        }
    }

    this.context.getNMStateStore().storeContainer(containerId,
        containerTokenIdentifier.getVersion(), containerStartTime,
request);

    // 调用
ApplicationEventDispatcher.handle(ApplicationEventType.INIT_CONTAINER)
    dispatcher.getEventHandler().handle(
        new ApplicationContainerInitEvent(container));

    this.context.getContainerTokenSecretManager().startContainerSuccessful(
        containerTokenIdentifier);
    NMAuditLogger.logSuccess(user, AuditConstants.START_CONTAINER,
        "ContainerManageImpl", applicationID, containerId);
    // TODO launchedContainer misplaced -> doesn't necessarily mean a
container

    // launch. A finished Application will not launch containers.
    metrics.launchedContainer();
    metrics.allocateContainer(containerTokenIdentifier.getResource());
} else {
    throw new YarnException(
        "Container start failed as the NodeManager is " +
        "in the process of shutting down");
}
} finally {

```

```
        this.readLock.unlock();
    }
}
```

4.2.1.2.1 备注说明

// 由于 **NodeManager** 启动 **MRAppMaster** 容器涉及很多的状态机转换 故文档不暂时 需要的话请阅读源码进行理解
// 所以直接来到 **ContainersScheduler.handle(SCHEDULE_CONTAINER)** 方法执行调度容器

4.2.1.2.2 调度执行 AM 容器(ContainerScheduler.handle(SCHEDULE_CONTAINER))

```
/**
 * The ContainersScheduler manages a collection of runnable containers. It
 * ensures that a container is launched only if all its launch criteria are
 * met. It also ensures that OPPORTUNISTIC containers are killed to make
 * room for GUARANTEED containers.
 */
public class ContainersScheduler extends AbstractService implements
    EventHandler<ContainersSchedulerEvent> {
    /**
     * Handle ContainersSchedulerEvents.
     *
     * @param event ContainersSchedulerEvent.
     */
    @Override
    public void handle(ContainersSchedulerEvent event) {
        switch (event.getType()) {
            // eventType = SCHEDULE_CONTAINER
            case SCHEDULE_CONTAINER:
                // 执行调度容器
                scheduleContainer(event.getContainer());
                break;
            // NOTE: Is sent only after container state has changed to PAUSED...
            case CONTAINER_PAUSED:
                // NOTE: Is sent only after container state has changed to DONE...
            case CONTAINER_COMPLETED:
                onResourcesReclaimed(event.getContainer());
                break;
            case UPDATE_CONTAINER:
                if (event instanceof UpdateContainersSchedulerEvent) {
                    onUpdateContainer((UpdateContainersSchedulerEvent) event);
                } else {
                    LOG.error("Unknown event type on UpdateContainer: " +
event.getType());
                }
                break;
            case SHED_QUEUED_CONTAINERS:
                shedQueuedOpportunisticContainers();
                break;
            case RECOVERY_COMPLETED:
```

```

        startPendingContainers(maxOppQueueLength <= 0);
        break;
    default:
        LOG.error("Unknown event arrived at ContainersScheduler: "
            + event.toString());
    }
}
}

```

```

@VisibleForTesting
protected void scheduleContainer(Container container) {
    // true
    boolean isGuaranteedContainer = container.getContainerTokenIdentifier().
        getExecutionType() == ExecutionType.GUARANTEED;

    // Given a guaranteed container, we enqueue it first and then try to start
    // as many queuing guaranteed containers as possible followed by queuing
    // opportunistic containers based on remaining resources available. If the
    // container still stays in the queue afterwards, we need to preempt just
    // enough number of opportunistic containers.
    if (isGuaranteedContainer) {
        // 将调度启动容器放入容器队列等待调度执行
        enqueueContainer(container);

        // When opportunistic container not allowed (which is determined by
        // max-queue length of pending opportunistic containers <= 0), start
        // guaranteed containers without looking at available resources.
        boolean forceStartGuaranteedContainers = (maxOppQueueLength <= 0);
        // 启动容器 (是否强制启动容器 默认 true)
        startPendingContainers(forceStartGuaranteedContainers);

        // if the guaranteed container is queued, we need to preempt
        opportunistic
        // containers for make room for it
        if (queuedGuaranteedContainers.containsKey(container.getContainerId()))
        {
            reclaimOpportunisticContainerResources(container);
        }
    } else {
        // Given an opportunistic container, we first try to start as many
        queuing
        // guaranteed containers as possible followed by queuing opportunistic
        // containers based on remaining resource available, then enqueue the
        // opportunistic container. If the container is enqueued, we do another
        // pass to try to start the newly enqueued opportunistic container.
        startPendingContainers(false);
        boolean containerQueued = enqueueContainer(container);
        // container may not get queued because the max opportunistic container
        // queue length is reached. If so, there is no point doing another pass
        if (containerQueued) {
            startPendingContainers(false);
        }
    }
}

```

```

    }
}
}

```

```

/**
 * Start pending containers in the queue.
 *
 * @param forceStartGuaranteedContainers When this is true, start guaranteed
 *                                       container without looking at available
resource
 */
private void startPendingContainers(boolean forceStartGuaranteedContainers) {
    // Start guaranteed containers that are paused, if resources available.
    // 判断 NodeManager 资源是否够用 (forceStartGuaranteedContainers = true)
    boolean resourcesAvailable = startContainers(
        queuedGuaranteedContainers.values(),
        forceStartGuaranteedContainers);
    // Start opportunistic containers, if resources available.
    if (resourcesAvailable) {
        startContainers(queuedOpportunisticContainers.values(), false);
    }
}

```

```

private boolean startContainers(
    Collection<Container> containersToBeStarted, boolean force) {
    // 获取待启动容器
    Iterator<Container> cIter = containersToBeStarted.iterator();
    boolean resourcesAvailable = true;
    while (cIter.hasNext() && resourcesAvailable) {
        // 变量容器
        Container container = cIter.next();
        // 尝试启动容器
        if (tryStartContainer(container, force)) {
            cIter.remove();
        } else {
            resourcesAvailable = false;
        }
    }
    return resourcesAvailable;
}

```

```

private boolean tryStartContainer(Container container, boolean force) {
    boolean containerStarted = false;
    // call startContainer without checking available resource when force==true
    if (force // 是否强制启动容器
        ||
        resourceAvailableToStartContainer(container) // 判断 NodeManager 资源
        是否够用启动容器
    ) {
        // 启动容器
        startContainer(container);
        containerStarted = true;
    }
    return containerStarted;
}

```

```

private void startContainer(Container container) {
    // Starting container [container_1684656010852_0005_01_000001]
    LOG.info("Starting container [" + container.getContainerId() + "]");
    // Skip to put into runningContainers and addUtilization when recover
    if (!runningContainers.containsKey(container.getContainerId())) {
        runningContainers.put(container.getContainerId(), container);
        this.utilizationTracker.addContainerResources(container);
    }
    if (container.getContainerTokenIdentifier().getExecutionType() ==
        ExecutionType.OPPORTUNISTIC) {
        this.metrics.startOpportunisticContainer(container.getResource());
    }
    // 发送启动容器事件 调用 ContainerImpl.sendLaunchEvent()
    container.sendLaunchEvent();
}

```

```

public class ContainerImpl implements Container {
    @SuppressWarnings("unchecked") // dispatcher not typed
    @Override
    public void sendLaunchEvent() {
        if (ContainerState.PAUSED == getContainerState()) {
            dispatcher.getEventHandler().handle(
                new ContainerResumeEvent(containerId,
                    "Container Resumed as some resources freed up"));
        } else {
            ContainersLauncherEventType launcherEvent =
                ContainersLauncherEventType.LAUNCH_CONTAINER;
            if (recoveredStatus == RecoveredContainerStatus.LAUNCHED) {
                // try to recover a container that was previously launched
                launcherEvent = ContainersLauncherEventType.RECOVER_CONTAINER;
            } else if (recoveredStatus == RecoveredContainerStatus.PAUSED) {
                launcherEvent =
                ContainersLauncherEventType.RECOVER_PAUSED_CONTAINER;
            }
        }
    }
}

```

```

        containerLaunchStartTime = clock.getTime();
        // 调用
ContainersLauncher.handle(ContainersLauncherEventType.LAUNCH_CONTAINER)
        dispatcher.getEventHandler().handle(
            new ContainersLauncherEvent(this, launcherEvent));
    }

}

}

```

4.2.1.2.3 真正启动 AM 容器(调用ContainersLauncher.handle(LAUNCH_CONTAINER))

```

/**
 * The launcher for the containers. This service should be started only after
 * the {@link ResourceLocalizationService} is started as it depends on creation
 * of system directories on the local file-system.
 */
public class ContainersLauncher extends AbstractService
    implements AbstractContainersLauncher {
    @Override
    public void handle(ContainersLauncherEvent event) {
        // TODO: ContainersLauncher launches containers one by one!!
        Container container = event.getContainer();
        ContainerId containerId = container.getContainerId();
        // eventType = LAUNCH_CONTAINER
        switch (event.getType()) {
            case LAUNCH_CONTAINER:
                // 获取启动容器对应的任务 ApplicationId
                Application app =
                    context.getApplications().get(
                        containerId.getApplicationAttemptId().getApplicationId());

                // 封装启动容器 ContainerLaunch (是一个 Callable 接口)
                ContainerLaunch launch =
                    new ContainerLaunch(context, getConfig(), dispatcher, exec,
app,
                    event.getContainer(), dirsHandler,
containerManager);

                // 调用 ContainerLaunch.call() 执行启动容器
                containerLauncher.submit(launch);
                running.put(containerId, launch);
                break;
            case RELAUNCH_CONTAINER:
                app = context.getApplications().get(
                    containerId.getApplicationAttemptId().getApplicationId());

                ContainerRelaunch relaunch =
                    new ContainerRelaunch(context, getConfig(), dispatcher,
exec, app,

```



```

        event.getContainer(), dirHandler,
containerManager);
        containerLauncher.submit(relaunch);
        running.put(containerId, relaunch);
        break;
    case RECOVER_CONTAINER:
        app = context.getApplications().get(
            containerId.getApplicationAttemptId().getApplicationId());
        launch = new RecoveredContainerLaunch(context, getConfig(),
dispatcher,
            exec, app, event.getContainer(), dirHandler,
containerManager);
        containerLauncher.submit(launch);
        running.put(containerId, launch);
        break;
    case RECOVER_PAUSED_CONTAINER:
        app = context.getApplications().get(
            containerId.getApplicationAttemptId().getApplicationId());
        launch = new RecoverPausedContainerLaunch(context, getConfig(),
dispatcher, exec, app, event.getContainer(), dirHandler,
            containerManager);
        containerLauncher.submit(launch);
        break;
    case CLEANUP_CONTAINER:
    case CLEANUP_CONTAINER_FOR_REINIT:
        ContainerLaunch launcher = running.remove(containerId);
        if (launcher == null) {
            // Container not launched.
            // triggering KILLING to CONTAINER_CLEANEDUP_AFTER_KILL
transition.
            dispatcher.getEventHandler().handle(
                new ContainerExitEvent(containerId,
                    ContainerEventType.CONTAINER_KILLED_ON_REQUEST,
                    Shell.WINDOWS ?
ContainerExecutor.ExitCode.FORCE_KILLED.getExitCode() :
ContainerExecutor.ExitCode.TERMINATED.getExitCode(),
                    "Container terminated before launch."));
            return;
        }

        // Cleanup a container whether it is running/killed/completed, so
that
        // no sub-processes are alive.
        try {
            launcher.cleanupContainer();
        } catch (IOException e) {
            LOG.warn("Got exception while cleaning container " + containerId
                + ". Ignoring.");
        }
        break;
    case SIGNAL_CONTAINER:

```

```

        SignalContainersLauncherEvent signalEvent =
            (SignalContainersLauncherEvent) event;
        ContainerLaunch runningContainer = running.get(containerId);
        if (runningContainer == null) {
            // Container not launched. So nothing needs to be done.
            LOG.info("Container " + containerId + " not running, nothing to
signal.");
            return;
        }

        try {
            runningContainer.signalContainer(signalEvent.getCommand());
        } catch (IOException e) {
            LOG.warn("Got exception while signaling container " +
containerId
                    + " with command " + signalEvent.getCommand());
        }
        break;
    case PAUSE_CONTAINER:
        ContainerLaunch launchedContainer = running.get(containerId);
        if (launchedContainer == null) {
            // Container not launched. So nothing needs to be done.
            return;
        }

        // Pause the container
        try {
            launchedContainer.pauseContainer();
        } catch (Exception e) {
            LOG.info("Got exception while pausing container: " +
                StringUtils.stringifyException(e));
        }
        break;
    case RESUME_CONTAINER:
        ContainerLaunch launchCont = running.get(containerId);
        if (launchCont == null) {
            // Container not launched. So nothing needs to be done.
            return;
        }

        // Resume the container.
        try {
            launchCont.resumeContainer();
        } catch (Exception e) {
            LOG.info("Got exception while resuming container: " +
                StringUtils.stringifyException(e));
        }
        break;
    }
}
}
}

```

```

public class ContainerLaunch implements Callable<Integer> {
    @Override
    @SuppressWarnings("unchecked") // dispatcher not typed
    public Integer call() {
        if (!validateContainerState()) {
            return 0;
        }

        // 获取启动容器的上下文对象
        final ContainerLaunchContext launchContext = container.getLaunchContext();
        // 获取启动容器信息
        ContainerId containerID = container.getContainerId();
        String containerIdStr = containerID.toString();
        // 获取启动容器的命令
        /**
         * 针对 MR 的 AM而言 基本启动 AM 命令模版
         * ${JAVA_HOME}/bin/java
         * -Djava.io.tmpdir=./tmp
         * -Dlog4j.configuration=container-log4j.properties
         * -Dyarn.app.container.log.dir=<LOG_DIR>
         * -Dyarn.app.container.log.filesize=0
         * -Dhadoop.root.logger=INFO,CLA
         * -Dhadoop.root.logfile=syslog
         * -Xmx1024m
         * org.apache.hadoop.mapreduce.v2.app.MRAppMaster
         * 1><LOG_DIR>.stdout
         * 2><LOG_DIR>.stderr
         */
        final List<String> command = launchContext.getCommands();
        int ret = -1;

        Path containerLogDir;
        try {
            Map<Path, List<String>> localResources = getLocalizedResources();

            final String user = container.getUser();
            // ////////////////////////////////// variable expansion
            // Before the container script gets written out.
            List<String> newCmds = new ArrayList<String>(command.size());
            String appIdStr = app.getAppId().toString();
            // appIdStr/containerIdStr (相对路径)
            String relativeContainerLogDir = ContainerLaunch
                .getRelativeContainerLogDir(appIdStr, containerIdStr);
            // 容器日志目录 (也即在本地创建日志目录)
            containerLogDir =
                dirsHandler.getLogPathForWrite(relativeContainerLogDir, false);
            recordContainerLogDir(containerID, containerLogDir.toString());

            for (String str : command) {
                // TODO: Should we instead work via symlinks without this grammar?
                // 扩展环境命令 (也即把 <LOG_DIR> 替换成容器日志目录 containerLogDir)
                newCmds.add(expandEnvironment(str, containerLogDir));
            }
        }
    }
}

```

```

}
// 重新更为启动容器命令
launchContext.setCommands(newCmds);

// 环境所有的环境变量
Map<String, String> environment = expandAllEnvironmentVars(
    launchContext, containerLogDir);
// ////////////////////////////////// End of variable expansion

// Use this to track variables that are added to the environment by nm.
LinkedHashSet<String> nmEnvVars = new LinkedHashSet<String>();

// 获取本地文件系统
FileContext lfs = FileContext.getLocalFSFileContext();

// 获取容器脚本路径 (xxx/launch_container)
Path nmPrivateContainerScriptPath = dirHandler.getLocalPathForWrite(
    getContainerPrivateDir(appIdStr, containerIdStr) +
Path.SEPARATOR
    + CONTAINER_SCRIPT);

// 获取私有 Token 路径
Path nmPrivateTokensPath = dirHandler.getLocalPathForWrite(
    getContainerPrivateDir(appIdStr, containerIdStr) +
Path.SEPARATOR
    + String.format(ContainerLocalizer.TOKEN_FILE_NAME_FMT,
        containerIdStr));

// 获取类JAR路径
Path nmPrivateClasspathJarDir = dirHandler.getLocalPathForWrite(
    getContainerPrivateDir(appIdStr, containerIdStr));

// Select the working directory for the container
// 获取容器工作目录
Path containerWorkDir = deriveContainerWorkDir();
recordContainerWorkDir(containerID, containerWorkDir.toString());

// 获取启动容器进程 PID 文件路径
String pidFileSubpath = getPidFileSubpath(appIdStr, containerIdStr);
// pid file should be in nm private dir so that it is not
// accessible by users
pidFilePath = dirHandler.getLocalPathForWrite(pidFileSubpath);

// 获取相关文件路径目录
List<String> localDirs = dirHandler.getLocalDirs();
List<String> localDirsForRead = dirHandler.getLocalDirsForRead();
List<String> logDirs = dirHandler.getLogDirs();
List<String> filecachedDirs = getNMFilecachedDirs(localDirsForRead);
List<String> userLocalDirs = getUserLocalDirs(localDirs);
List<String> containerLocalDirs = getContainerLocalDirs(localDirs);
List<String> containerLogDirs = getContainerLogDirs(logDirs);
List<String> userFilecachedDirs = getUserFilecachedDirs(localDirsForRead);
List<String> applicationLocalDirs = getApplicationLocalDirs(localDirs,
appIdStr);

```

```

if (!dirsHandler.areDisksHealthy()) {
    ret = ContainerExitStatus.DISKS_FAILED;
    throw new IOException("Most of the disks failed. "
        + dirsHandler.getDisksHealthReport(false));
}
List<Path> appDirs = new ArrayList<Path>(localDirs.size());
for (String localDir : localDirs) {
    Path usersdir = new Path(localDir, ContainerLocalizer.USERCACHE);
    Path userdir = new Path(usersdir, user);
    Path appsdire = new Path(userdir, ContainerLocalizer.APPCACHE);
    appDirs.add(new Path(appsdire, appIdStr));
}

// Set the token location too.
// 设置本地 Token (xxx/container_tokens)
addToEnvMap(environment, nmEnvVars,
    ApplicationConstants.CONTAINER_TOKEN_FILE_ENV_NAME,
    new Path(containerWorkDir,
        FINAL_CONTAINER_TOKENS_FILE).toUri().getPath());

// ////////// write out the container-script in the nmPrivate space.
// 写启动容器脚本到文件 (launch_container)
try (DataOutputStream containerScriptOutputStream =
    lfs.create(nmPrivateContainersScriptPath,
        EnumSet.of(CREATE, OVERWRITE))) {
    // Sanitize the container's environment
    sanitizeEnv(environment, containerWorkDir, appDirs, userLocalDirs,
        containerLogDirs, localResources, nmPrivateClasspathJarDir,
        nmEnvVars);

    // 准备容器
    prepareContainer(localResources, containerLocalDirs);

    // write out the environment
    // 写相关环境信息到 launch_container 脚本
    exec.writeLaunchEnv(containerScriptOutputStream, environment,
        localResources, launchContext.getCommands(),
        containerLogDir, user, nmEnvVars);
}
// ////////// End of writing out container-script

// ////////// write out the container-tokens in the nmPrivate space.
// 写容器 Token 消息到文件
try (DataOutputStream tokensOutputStream =
    lfs.create(nmPrivateTokensPath, EnumSet.of(CREATE,
OVERWRITE))) {
    Credentials creds = container.getCredentials();
    creds.writeTokenStorageToStream(tokensOutputStream);
}
// ////////// End of writing out container-tokens

```

```

// 执行脚本启动容器
ret = launchContainer(new ContainerStartContext.Builder()
    .setContainer(container)
    .setLocalizedResources(localResources)
    .setNmPrivateContainersScriptPath(nmPrivateContainersScriptPath)
    .setNmPrivateTokensPath(nmPrivateTokensPath)
    .setUser(user)
    .setAppId(appIdStr)
    .setContainerWorkDir(containerWorkDir)
    .setLocalDirs(localDirs)
    .setLogDirs(logDirs)
    .setFilecachedDirs(filecachedDirs)
    .setUserLocalDirs(userLocalDirs)
    .setContainerLocalDirs(containerLocalDirs)
    .setContainerLogDirs(containerLogDirs)
    .setUserFilecachedDirs(userFilecachedDirs)
    .setApplicationLocalDirs(applicationLocalDirs).build());
} catch (ConfigurationException e) {
    LOG.error("Failed to launch container due to configuration error.", e);
    dispatcher.getEventHandler().handle(new ContainerExitEvent(
        containerID, ContainerEventType.CONTAINER_EXITED_WITH_FAILURE,
ret,
        e.getMessage()));
    // Mark the node as unhealthy
    context.getNodeStatusUpdater().reportException(e);
    return ret;
} catch (Throwable e) {
    LOG.warn("Failed to launch container.", e);
    dispatcher.getEventHandler().handle(new ContainerExitEvent(
        containerID, ContainerEventType.CONTAINER_EXITED_WITH_FAILURE,
ret,
        e.getMessage()));
    return ret;
} finally {
    setContainerCompletedStatus(ret);
}

handleContainerExitCode(ret, containerLogDir);
return ret;
}
}

```

五 启动 MRAppMaster 进程 (MR 任务的 ApplicationMaster)

```

/**
 * The Map-Reduce Application Master.
 * The state machine is encapsulated in the implementation of Job interface.
 * All state changes happens via Job interface. Each event
 * results in a Finite State Transition in Job.

```

```

*
* MR AppMaster is the composition of loosely coupled services. The services
* interact with each other via events. The components resembles the
* Actors model. The component acts on received event and send out the
* events to other components.
* This keeps it highly concurrent with no or minimal synchronization needs.
*
* The events are dispatched by a central Dispatch mechanism. All components
* register to the Dispatcher.
*
* The information is shared across different components using AppContext.
*/

```

```

@SuppressWarnings("rawtypes")
public class MRAppMaster extends CompositeService {
    public static void main(String[] args) {
        try {
            mainStarted = true;
            Thread.setDefaultUncaughtExceptionHandler(new
YarnUncaughtExceptionHandler());
            // 从当前 NodeManager 的节点获取当前启动容器的环境变量 key = CONTAINER_ID
            /**
             * 这些环境变量存储在 launch_container 脚本里
             * 以 export key = value 形式暴露
             */
            String containerIdStr =
                System.getenv(Environment.CONTAINER_ID.name());
            String nodeHostString = System.getenv(Environment.NM_HOST.name());
            String nodePortString = System.getenv(Environment.NM_PORT.name());
            String nodeHttpPortString =
                System.getenv(Environment.NM_HTTP_PORT.name());
            String appSubmitTimeStr =
                System.getenv(ApplicationConstants.APP_SUBMIT_TIME_ENV);
            validateInputParam(containerIdStr,
                Environment.CONTAINER_ID.name());
            validateInputParam(nodeHostString, Environment.NM_HOST.name());
            validateInputParam(nodePortString, Environment.NM_PORT.name());
            validateInputParam(nodeHttpPortString,
                Environment.NM_HTTP_PORT.name());
            validateInputParam(appSubmitTimeStr,
                ApplicationConstants.APP_SUBMIT_TIME_ENV);

            ContainerId containerId = ContainerId.fromString(containerIdStr);
            ApplicationAttemptId applicationAttemptId =
                containerId.getApplicationAttemptId();
            if (applicationAttemptId != null) {
                CallerContext.setCurrent(new CallerContext.Builder(
                    "mr_appmaster_" + applicationAttemptId.toString()).build());
            }
            long appSubmitTime = Long.parseLong(appSubmitTimeStr);

            // 创建 MRAppMaster

```

```

MRAppMaster appMaster =
    new MRAppMaster(
        applicationAttemptId,
        containerId,
        nodeHostString,
        Integer.parseInt(nodePortString),
        Integer.parseInt(nodeHttpPortString),
        appSubmitTime);

ShutdownHookManager.get().addShutdownHook(
    new MRAppMasterShutdownHook(appMaster), SHUTDOWN_HOOK_PRIORITY);
JobConf conf = new JobConf(new YarnConfiguration());
conf.addResource(new Path(MRJobConfig.JOB_CONF_FILE));

MRWebAppUtil.initialize(conf);
// log the system properties
String systemPropsToLog = MRApps.getSystemPropertiesToLog(conf);
if (systemPropsToLog != null) {
    /**
     * /*****
     * [system properties]
     * os.name: Linux
     * os.version: 3.10.0-1062.el7.x86_64
     * java.home: /opt/app/jdk1.8.0_212/jre
     * java.runtime.version: 1.8.0_212-b10
     * java.vendor: Oracle Corporation
     * java.version: 1.8.0_212
     * java.vm.name: Java HotSpot(TM) 64-Bit Server VM

```



```
* java.class.path: /opt/app/hadoop-3.1.3/data/nm-local-  
dir/usercache/tanbs/appcache/application_1684656010852_0006/container_1684656010852_  
0006_01_000001:/opt/app/hadoop-3.1.3/etc/hadoop:/opt/app/hadoop-  
3.1.3/share/hadoop/common/hadoop-common-3.1.3-tests.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/hadoop-common-3.1.3.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/hadoop-kms-3.1.3.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/hadoop-nfs-3.1.3.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/jetty-io-9.3.24.v20180605.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/accessors-smart-1.2.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/jcip-annotations-1.0-1.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/animal-sniffer-annotations-1.17.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/netty-3.10.5.Final.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/asm-5.0.4.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/jersey-core-1.19.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/audience-annotations-0.5.0.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/nimbus-jose-jwt-4.41.1.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/avro-1.7.7.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/jetty-security-9.3.24.v20180605.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/checker-qual-2.5.2.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/jetty-server-9.3.24.v20180605.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/commons-beanutils-1.9.3.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/paranamer-2.3.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/commons-cli-1.2.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/jsr311-api-1.1.1.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/commons-codec-1.11.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/jersey-json-1.19.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/commons-collections-3.2.2.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/jul-to-slf4j-1.7.25.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/commons-compress-1.18.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/jersey-servlet-1.19.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/commons-configuration2-2.1.1.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/protobuf-java-2.5.0.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/commons-io-2.5.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/re2j-1.1.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/commons-lang-2.6.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/kerb-admin-1.0.1.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/commons-lang3-3.4.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/kerb-client-1.0.1.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/commons-logging-1.1.3.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/kerb-common-1.0.1.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/commons-math3-3.1.1.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/slf4j-api-1.7.25.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/commons-net-3.6.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/kerb-core-1.0.1.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/curator-client-2.13.0.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/kerb-crypto-1.0.1.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/curator-framework-2.13.0.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/kerb-identity-1.0.1.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/curator-recipes-2.13.0.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/jersey-server-1.19.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/error_prone_annotations-2.2.0.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/kerb-server-1.0.1.jar:/opt/app/hadoop-
```

3.1.3/share/hadoop/common/lib/failureaccess-1.0.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/gson-2.2.4.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/snappy-java-1.0.5.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/guava-27.0-jre.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/kerb-simplekdc-1.0.1.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/hadoop-annotations-3.1.3.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/kerb-util-1.0.1.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/hadoop-auth-3.1.3.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jsch-0.1.54.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/htrace-core4-4.1.0-incubating.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/stax2-api-3.1.4.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/httpclient-4.5.2.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/token-provider-1.0.1.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/httpcore-4.4.4.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/kerby-asn1-1.0.1.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/j2objc-annotations-1.1.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jettison-1.1.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jackson-annotations-2.7.8.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/kerby-config-1.0.1.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jackson-core-2.7.8.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/kerby-pkix-1.0.1.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jackson-core-asl-1.9.13.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/kerby-util-1.0.1.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jackson-databind-2.7.8.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/kerby-xdr-1.0.1.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jackson-jaxrs-1.9.13.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jetty-http-9.3.24.v20180605.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jackson-mapper-asl-1.9.13.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jackson-xc-1.9.13.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/log4j-1.2.17.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/javax.servlet-api-3.1.0.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/woodstox-core-5.0.3.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jaxb-api-2.2.11.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/metrics-core-3.2.4.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jaxb-impl-2.2.3-1.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jetty-servlet-9.3.24.v20180605.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/json-smart-2.3.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jetty-util-9.3.24.v20180605.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jsp-api-2.1.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jetty-webapp-9.3.24.v20180605.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jsr305-3.0.0.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jetty-xml-9.3.24.v20180605.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/listenablefuture-9999.0-empty-to-avoid-conflict-with-
guava.jar:/opt/app/hadoop-3.1.3/share/hadoop/common/lib/zookeeper-
3.4.13.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/hadoop-hdfs-3.1.3-
tests.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/hadoop-hdfs-
3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/hadoop-hdfs-client-3.1.3-
tests.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/hadoop-hdfs-client-
3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/hadoop-hdfs-httpfs-
3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/hadoop-hdfs-native-client-3.1.3-
tests.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/hadoop-hdfs-native-client-

3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/hadoop-hdfs-nfs-
3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/hadoop-hdfs-rbf-3.1.3-
tests.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/hadoop-hdfs-rbf-
3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jetty-http-
9.3.24.v20180605.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/accessors-smart-
1.2.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jaxb-impl-2.2.3-
1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/animal-sniffer-annotations-
1.17.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/log4j-
1.2.17.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/asm-
5.0.4.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jcip-annotations-1.0-
1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/audience-annotations-
0.5.0.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/netty-
3.10.5.Final.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/avro-
1.7.7.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jetty-io-
9.3.24.v20180605.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/checker-qual-
2.5.2.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jetty-security-
9.3.24.v20180605.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/commons-beanutils-
1.9.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/netty-all-
4.0.52.Final.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/commons-cli-
1.2.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/json-smart-
2.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/commons-codec-
1.11.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jersey-core-
1.19.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/commons-collections-
3.2.2.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jsr305-
3.0.0.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/commons-compress-
1.18.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jersey-server-
1.19.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/commons-configuration2-
2.1.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jsr311-api-
1.1.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/commons-daemon-
1.0.13.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/nimbus-jose-jwt-
4.41.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/commons-io-
2.5.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/okhttp-
2.7.5.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/commons-lang-
2.6.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerb-admin-
1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/commons-lang3-
3.4.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerb-client-
1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/commons-logging-
1.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerb-common-
1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/commons-math3-
3.1.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/okio-
1.6.0.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/commons-net-
3.6.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerb-core-
1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/curator-client-
2.13.0.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerb-crypto-
1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/curator-framework-
2.13.0.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerb-identity-
1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/curator-recipes-
2.13.0.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jersey-json-
1.19.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/error_prone_annotations-
2.2.0.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerb-server-
1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/failureaccess-
1.0.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/paranamer-

2.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/gson-2.2.4.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/protobuf-java-2.5.0.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/guava-27.0-jre.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerb-simplekdc-1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/hadoop-annotations-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerb-util-1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/hadoop-auth-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jetty-webapp-9.3.24.v20180605.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/htrace-core4-4.1.0-incubating.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/re2j-1.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/httpclient-4.5.2.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/snappy-java-1.0.5.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/httpcore-4.4.4.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerby-asn1-1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/j2objc-annotations-1.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jersey-servlet-1.19.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jackson-annotations-2.7.8.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerby-config-1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jackson-core-2.7.8.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerby-pkix-1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jackson-core-asl-1.9.13.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerby-util-1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jackson-databind-2.7.8.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerby-xdr-1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jackson-jaxrs-1.9.13.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jettison-1.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jackson-mapper-asl-1.9.13.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/leveldbjni-all-1.8.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jackson-xc-1.9.13.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/javax.servlet-api-3.1.0.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/stax2-api-3.1.4.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jaxb-api-2.2.11.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jetty-server-9.3.24.v20180605.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jetty-xml-9.3.24.v20180605.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jetty-servlet-9.3.24.v20180605.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jsch-0.1.54.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jetty-util-9.3.24.v20180605.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/json-simple-1.1.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jetty-util-ajax-9.3.24.v20180605.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/listenablefuture-9999.0-empty-to-avoid-conflict-with-guava.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/token-provider-1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/woodstox-core-5.0.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/zookeeper-3.4.13.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-api-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-applications-distributedshell-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-applications-unmanaged-am-launcher-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-client-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-common-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-registry-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-server-applicationhistoryservice-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-server-common-

3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-server-nodemanager-
3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-server-
resourcemanager-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-
server-router-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-server-
sharedcachemanager-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-
server-tests-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-server-
timeline-pluginstorage-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-
yarn-server-web-proxy-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-
services-api-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-services-
core-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/HikariCP-java7-
2.4.12.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/aopalliance-
1.0.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/dnsjava-
2.1.7.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/ehcache-
3.3.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/fst-2.50.jar:/opt/app/hadoop-
3.1.3/share/hadoop/yarn/lib/geronimo-jcache_1.0_spec-1.0-alpha-
1.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/guice-4.0.jar:/opt/app/hadoop-
3.1.3/share/hadoop/yarn/lib/guice-servlet-4.0.jar:/opt/app/hadoop-
3.1.3/share/hadoop/yarn/lib/jackson-jaxrs-base-2.7.8.jar:/opt/app/hadoop-
3.1.3/share/hadoop/yarn/lib/jackson-jaxrs-json-provider-2.7.8.jar:/opt/app/hadoop-
3.1.3/share/hadoop/yarn/lib/jackson-module-jaxb-annotations-
2.7.8.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/java-util-
1.9.0.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/javax.inject-
1.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/jersey-client-
1.19.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/jersey-guice-
1.19.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/json-io-
2.5.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/metrics-core-
3.2.4.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/mssql-jdbc-
6.2.1.jre7.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/objenesis-
1.0.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/sakeyam1-
1.16.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/swagger-annotations-
1.5.4.jar:/opt/app/hadoop-3.1.3/share/hadoop/mapreduce/hadoop-mapreduce-client-app-
3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/mapreduce/hadoop-mapreduce-client-
common-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/mapreduce/hadoop-mapreduce-
client-core-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/mapreduce/hadoop-mapreduce-
client-hs-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/mapreduce/hadoop-mapreduce-
client-hs-plugins-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/mapreduce/hadoop-
mapreduce-client-jobclient-3.1.3-tests.jar:/opt/app/hadoop-
3.1.3/share/hadoop/mapreduce/hadoop-mapreduce-client-jobclient-
3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/mapreduce/hadoop-mapreduce-client-
nativetask-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/mapreduce/hadoop-mapreduce-
client-shuffle-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/mapreduce/hadoop-
mapreduce-client-uploader-3.1.3.jar:/opt/app/hadoop-
3.1.3/share/hadoop/mapreduce/hadoop-mapreduce-examples-3.1.3.jar:/opt/app/hadoop-
3.1.3/share/hadoop/mapreduce/lib/hamcrest-core-1.3.jar:/opt/app/hadoop-
3.1.3/share/hadoop/mapreduce/lib/junit-
4.11.jar:/opt/app/hadoop-3.1.3/share/hadoop/mapreduce/lib/job.jar:/opt/app/hadoop-
3.1.3/data/nm-local-
dir/usercache/tanbs/appcache/application_1684656010852_0006/container_1684656010852_0006_01_000001/job.jar
* java.io.tmpdir: /opt/app/hadoop-3.1.3/data/nm-local-
dir/usercache/tanbs/appcache/application_1684656010852_0006/container_1684656010852_0006_01_000001/tmp

```

        * user.dir: /opt/app/hadoop-3.1.3/data/nm-local-
dir/usercache/tanbs/appcache/application_1684656010852_0006/container_1684656010852_
0006_01_000001

        * user.name: tanbs
        * *****

        */
    LOG.info(systemPropsToLog);
}

String jobUserName = System
    .getenv(ApplicationConstants.Environment.USER.name());
conf.set(MRJobConfig.USER_NAME, jobUserName);

// 初始化并启动 AM
initAndStartAppMaster(appMaster, conf, jobUserName);
} catch (Throwable t) {
    LOG.error("Error starting MRAppMaster", t);
    ExitUtil.terminate(1, t);
}
}
}
}

```

5.1 创建 MRAppMaster 组合服务

```

/**
 * The Map-Reduce Application Master.
 * The state machine is encapsulated in the implementation of Job interface.
 * All state changes happens via Job interface. Each event
 * results in a Finite State Transition in Job.
 * <p>
 * MR AppMaster is the composition of loosely coupled services. The services
 * interact with each other via events. The components resembles the
 * Actors model. The component acts on received event and send out the
 * events to other components.
 * This keeps it highly concurrent with no or minimal synchronization needs.
 * <p>
 * The events are dispatched by a central Dispatch mechanism. All components
 * register to the Dispatcher.
 * <p>
 * The information is shared across different components using AppContext.
 */

@SuppressWarnings("rawtypes")
public class MRAppMaster extends CompositeService {
    public MRAppMaster(ApplicationAttemptId applicationAttemptId,
        ContainerId containerId, String nmHost, int nmPort, int
nmHttpPort,
        long appSubmitTime) {
        // 往下追
        this(applicationAttemptId, containerId, nmHost, nmPort, nmHttpPort,
            SystemClock.getInstance(), appSubmitTime);
    }
}

```



```

    }

    public MRAppMaster(ApplicationAttemptId applicationAttemptId,
                        ContainerId containerId, String nmHost, int nmPort, int
nmHttpPort,
                        Clock clock, long appSubmitTime) {
        super(MRAppMaster.class.getName());
        this.clock = clock;
        this.startTime = clock.getTime();
        this.appSubmitTime = appSubmitTime;
        this.appAttemptID = applicationAttemptId;
        this.containerID = containerId;
        this.nmHost = nmHost;
        this.nmPort = nmPort;
        this.nmHttpPort = nmHttpPort;
        this.metrics = MRAppMetrics.create();
        logSyncer = TaskLog.createLogSyncer();
        // Created MRAppMaster for application appattempt_1684656010852_0006_000001
        LOG.info("Created MRAppMaster for application " + applicationAttemptId);
    }
}

```

5.2 初始化并启动 AM

```

protected static void initAndStartAppMaster(final MRAppMaster appMaster,
                                              final JobConf conf, String
jobUserName) throws IOException,
                          InterruptedException {
    UserGroupInformation.setConfiguration(conf);
    // MAPREDUCE-6565: need to set configuration for SecurityUtil.
    SecurityUtil.setConfiguration(conf);
    // Security framework already loaded the tokens into current UGI, just use
    // them
    Credentials credentials =
        UserGroupInformation.getCurrentUser().getCredentials();

    // Executing with tokens:
    // [Kind: YARN_AM_RM_TOKEN, Service: ,
    // Idnt: (appAttemptId { application_id
    // { id: 6 cluster_timestamp: 1684656010852 } attemptId: 1 } keyId:
-1993740727)]
    LOG.info("Executing with tokens: {}", credentials.getAllTokens());

    UserGroupInformation appMasterUgi = UserGroupInformation
        .createRemoteUser(jobUserName);
    appMasterUgi.addCredentials(credentials);

    // Now remove the AM->RM token so tasks don't have it
    Iterator<Token<?>> iter = credentials.getAllTokens().iterator();
    while (iter.hasNext()) {
        Token<?> token = iter.next();
    }
}

```

```

        if (token.getKind().equals(AMRMTokenIdentifier.KIND_NAME)) {
            iter.remove();
        }
    }
    conf.getCredentials().addAll(credentials);
    appMasterUgi.doAs(new PrivilegedExceptionAction<Object>() {
        @Override
        public Object run() throws Exception {
            // 初始化 MRAppMaster (调用 MRAppMaster.serviceInit())
            appMaster.init(conf);
            // 启动 MRAppMaster (调用 MRAppMaster.serviceStart())
            appMaster.start();
            if (appMaster.errorHappenedShutDown) {
                throw new IOException("was asked to shut down.");
            }
            return null;
        }
    });
}

```

5.2.1 初始化 MRAppMaster (调用 MRAppMaster.serviceInit())

```

@Override
protected void serviceInit(final Configuration conf) throws Exception {
    // create the job classloader if enabled
    // 如果开启了任务类加载器 默认不开启
    createJobClassLoader(conf);

    initJobCredentialsAndUGI(conf);

    // 创建并添加 AsyncDispatcher 服务
    dispatcher = createDispatcher();
    addIfService(dispatcher);

    // 创建并添加任务完成监控服务 TaskAttemptFinishingMonitor
    taskAttemptFinishingMonitor =
    createTaskAttemptFinishingMonitor(dispatcher.getEventHandler());
    addIfService(taskAttemptFinishingMonitor);

    // 创建 AM 运行上下文对象 RunningAppContext
    context = new RunningAppContext(conf, taskAttemptFinishingMonitor);

    // Job name is the same as the app name until we support DAG of jobs
    // for an app later
    // 获取任务名称
    appName = conf.get(MRJobConfig.JOB_NAME, "<missing app name>");

    // 设置 key = mapreduce.job.application.attempt.id value =
    ApplicationAttemptId
    conf.setInt(MRJobConfig.APPLICATION_ATTEMPT_ID,
    appAttemptID.getAttemptId());
}

```



```

newApiCommitter = false;
jobId = MRBuilderUtils.newJobId(appAttemptID.getApplicationId(),
    appAttemptID.getApplicationId().getId());
int numReduceTasks = conf.getInt(MRJobConfig.NUM_REDUCES, 0);
if ((numReduceTasks > 0 &&
    conf.getBoolean("mapred.reducer.new-api", false)) ||
    (numReduceTasks == 0 &&
    conf.getBoolean("mapred.mapper.new-api", false))) {
    newApiCommitter = true;
    // Using mapred newApiCommitter.
    LOG.info("Using mapred newApiCommitter.");
}

boolean copyHistory = false;
// 创建 FileOutputCommitter
committer = createOutputCommitter(conf);
try {
    String user = UserGroupInformation.getCurrentUser().getShortUserName();
    Path stagingDir = MRApps.getStagingAreaDir(conf, user);
    FileSystem fs = getFileSystem(conf);

    boolean stagingExists = fs.exists(stagingDir);
    Path startCommitFile = MRApps.getStartJobCommitFile(conf, user, jobId);
    boolean commitStarted = fs.exists(startCommitFile);
    Path endCommitSuccessFile = MRApps.getEndJobCommitSuccessFile(conf,
user, jobId);
    boolean commitSuccess = fs.exists(endCommitSuccessFile);
    Path endCommitFailureFile = MRApps.getEndJobCommitFailureFile(conf,
user, jobId);
    boolean commitFailure = fs.exists(endCommitFailureFile);

    if (!stagingExists) {
        isLastAMRetry = true;
        LOG.info("Attempt num: " + appAttemptID.getAttemptId() +
            " is last retry: " + isLastAMRetry +
            " because the staging dir doesn't exist.");
        errorHappenedShutDown = true;
        forcedState = JobStateInternal.ERROR;
        shutDownMessage = "Staging dir does not exist " + stagingDir;
        LOG.error(shutDownMessage);
    } else if (commitStarted) {
        //A commit was started so this is the last time, we just need to
know

        // what result we will use to notify, and how we will unregister
errorHappenedShutDown = true;
        isLastAMRetry = true;
        LOG.info("Attempt num: " + appAttemptID.getAttemptId() +
            " is last retry: " + isLastAMRetry +
            " because a commit was started.");
        copyHistory = true;
        if (commitSuccess) {

```

```

        shutDownMessage =
            "Job commit succeeded in a prior MRAppMaster attempt " +
            "before it crashed. Recovering.";
        forcedState = JobStateInternal.SUCCEEDED;
    } else if (commitFailure) {
        shutDownMessage =
            "Job commit failed in a prior MRAppMaster attempt " +
            "before it crashed. Not retrying.";
        forcedState = JobStateInternal.FAILED;
    } else {
        if (isCommitJobRepeatable()) {
            // cleanup previous half done commits if committer supports
            // repeatable job commit.
            errorHappenedShutDown = false;
            cleanupInterruptedCommit(conf, fs, startCommitFile);
        } else {
            //The commit is still pending, commit error
            shutDownMessage =
                "Job commit from a prior MRAppMaster attempt is " +
                "potentially in progress. Preventing
multiple commit executions";
            forcedState = JobStateInternal.ERROR;
        }
    }
}

} catch (IOException e) {
    throw new YarnRuntimeException("Error while initializing", e);
}

// 默认 false
if (errorHappenedShutDown) {
    NoopEventHandler eater = new NoopEventHandler();
    //We do not have a JobEventDispatcher in this path
    dispatcher.register(JobEventType.class, eater);

    EventHandler<JobHistoryEvent> historyService = null;
    if (copyHistory) {
        historyService =
            createJobHistoryHandler(context);

        dispatcher.register(org.apache.hadoop.mapreduce.jobhistory.EventType.class,
            historyService);
    } else {
        dispatcher.register(org.apache.hadoop.mapreduce.jobhistory.EventType.class,
            eater);
    }

    if (copyHistory) {
        // Now that there's a FINISHING state for application on RM to give

```

```

        // plenty of time to clean up after unregister it's safe to clean
staging
        // directory after unregistering with RM. So, we start the staging-
dir
        // cleaner BEFORE the ContainerAllocator so that on shut-down,
        // ContainerAllocator unregisters first and then the staging-dir
cleaner
        // deletes staging directory.
        addService(createStagingDirCleaningService());
    }

    // service to allocate containers from RM (if non-uber) or to fake it
(uber)
    containerAllocator = createContainerAllocator(null, context);
    addIfService(containerAllocator);
    dispatcher.register(ContainerAllocator.EventType.class,
containerAllocator);

    if (copyHistory) {
        // Add the JobHistoryEventHandler last so that it is properly
stopped first.
        // This will guarantee that all history-events are flushed before AM
goes
        // ahead with shutdown.
        // Note: Even though JobHistoryEventHandler is started last, if any
        // component creates a JobHistoryEvent in the meanwhile, it will be
just be
        // queued inside the JobHistoryEventHandler
        addIfService(historyService);

        JobHistoryCopyService cpHist = new
JobHistoryCopyService(appAttemptID,
            dispatcher.getEventHandler());
        addIfService(cpHist);
    }
    else {

        //service to handle requests from JobClient
        // 创建 MRClientService 服务 该服务是与 JobClient 客户端交互(也即提交任务的客户
端)
        // JobClient 持续获取提交任务的状态信息
        // 通讯协议接口为 MRClientProtocol
        clientService = createClientService(context);
        // Init ClientService separately so that we stop it separately, since
this
        // service needs to wait some time before it stops so clients can know
the
        // final states
        // 初始化 MRClientService
        // (调用 MRClientService的 父类 AbstractService.serviceInit())
        clientService.init(conf);

```

```

// 创建容器申请器 ContainerAllocatorRouter
containerAllocator = createContainerAllocator(clientService, context);

//service to handle the output committer
// 创建并添加 CommitterEventHandler 服务
committerEventHandler = createCommitterEventHandler(context, committer);
addIfService(committerEventHandler);

//policy handling preemption requests from RM
callWithJobClassLoader(conf, new Action<Void>() {
    public void call(Configuration conf) {
        preemptionPolicy = createPreemptionPolicy(conf);
        preemptionPolicy.init(context);
        return null;
    }
});

//service to handle requests to TaskUmbilicalProtocol
// 创建并添加任务监听服务 TaskAttemptListenerImpl
taskAttemptListener = createTaskAttemptListener(context,
preemptionPolicy);
addIfService(taskAttemptListener);

//service to log job history events
// 创建并注册任务历史事件处理器 JobHistoryEventHandler
EventHandler<JobHistoryEvent> historyService =
    createJobHistoryHandler(context);

dispatcher.register(org.apache.hadoop.mapreduce.jobhistory.EventType.class,
    historyService);

// 创建任务事件分发器 JobEventDispatcher
this.jobEventDispatcher = new JobEventDispatcher();

//register the event dispatchers
// 注册相关事件类型
dispatcher.register(JobEventType.class, jobEventDispatcher);
dispatcher.register(TaskEventType.class, new TaskEventDispatcher());
dispatcher.register(TaskAttemptEventType.class,
    new TaskAttemptEventDispatcher());
dispatcher.register(CommitterEventType.class, committerEventHandler);

// 是否开启任务的推测执行
if (conf.getBoolean(MRJobConfig.MAP_SPECULATIVE, false)
    || conf.getBoolean(MRJobConfig.REDUCE_SPECULATIVE, false)) {
    //optional service to speculate on task attempts' progress
    // 创建并注册任务推测执行 DefaultSpeculator 服务
    speculator = createSpeculator(conf, context);
    addIfService(speculator);
}

// 注册推测执行事件类型

```

```

speculatorEventDispatcher = new SpeculatorEventDispatcher(conf);
dispatcher.register(Speculator.EventType.class,
    speculatorEventDispatcher);

// Now that there's a FINISHING state for application on RM to give AMS
// plenty of time to clean up after unregister it's safe to clean
staging
// directory after unregistering with RM. So, we start the staging-dir
// cleaner BEFORE the ContainerAllocator so that on shut-down,
// ContainerAllocator unregisters first and then the staging-dir cleaner
// deletes staging directory.
// 创建任务清除目录服务 StagingDirCleaningService
addService(createStagingDirCleaningService());

// service to allocate containers from RM (if non-uber) or to fake it
(uber)
// 添加容器申请器 ContainerAllocatorRouter
addIfService(containerAllocator);
dispatcher.register(ContainerAllocator.EventType.class,
containerAllocator);

// corresponding service to launch allocated containers via NodeManager
// 创建并注册启动容器服务 ContainerLauncherRouter
containerLauncher = createContainerLauncher(context);
addIfService(containerLauncher);
dispatcher.register(ContainerLauncher.EventType.class,
containerLauncher);

// Add the JobHistoryEventHandler last so that it is properly stopped
first.
// This will guarantee that all history-events are flushed before AM
goes
// ahead with shutdown.
// Note: Even though JobHistoryEventHandler is started last, if any
// component creates a JobHistoryEvent in the meanwhile, it will be just
be
// queued inside the JobHistoryEventHandler
// 注册任务历史事件处理器 JobHistoryEventHandler
addIfService(historyService);
}
/**
 * MRAppMaster 组合服务的 AsyncDispatcher 异步事件分发器注册了哪些事件注册表
 * 1 EventType          -> JobHistoryEventHandler (任务历史事件处理器)
 * 2 JobEventType        -> JobEventDispatcher (任务事件分发器)
 * 3 TaskEventType        -> TaskEventDispatcher (Task事件分发器)
 * 4 TaskAttemptEventType -> TaskAttemptEventDispatcher (Task 尝试事件分发
器)
 * 5 CommitterEventType   -> CommitterEventHandler (提交事件处理器)
 * 6 Speculator.EventType -> SpeculatorEventDispatcher (推测事件处理器)
 * 7 ContainerAllocator.EventType -> ContainerAllocatorRouter (容器申请事件处理
器)

```

```

    * 8 ContainerLauncher.EventType -> ContainerLauncherRouter (启动容器事件处理
器)

    */

/**
 * MRAppMaster 组合服务有哪些子服务?
 * 1 异步事件分发器 -> AsyncDispatcher
 * 2 任务完成监控服务 -> TaskAttemptFinishingMonitor
 * 3 提交者事件处理 -> CommitterEventHandler
 * 4 Task监听服务 -> TaskAttemptListenerImpl
 * 5 任务推测执行服务 -> DefaultSpeculator
 * 6 任务清除目录服务 -> StagingDirCleaningService
 * 7 容器申请服务 -> ContainerAllocatorRouter
 * 8 启动容器服务 -> ContainerLauncherRouter
 * 9 任务历史事件处理器服务 -> JobHistoryEventHandler
 */

// 调用 MRAppMaster 组合服务的所有子服务 serviceInit()
super.serviceInit(conf);
} // end of init()

```

5.2.1.1 异步事件分发器 AsyncDispatcher.serviceInit()

```
// 调用其父类 AbstractService.serviceInit()
```

5.2.1.2 任务完成监控服务 TaskAttemptFinishingMonitor.serviceInit()

```
// 调用其父类 AbstractService.serviceInit()
```

5.2.1.3 提交者事件处理 CommitterEventHandler.serviceInit()

```

public class CommitterEventHandler extends AbstractService
    implements EventHandler<CommitterEvent> {
    @Override
    protected void serviceInit(Configuration conf) throws Exception {
        super.serviceInit(conf);
        // 默认 60s
        commitThreadCancelTimeoutMs = conf.getInt(
            MRJobConfig.MR_AM_COMMITTER_CANCEL_TIMEOUT_MS,
            MRJobConfig.DEFAULT_MR_AM_COMMITTER_CANCEL_TIMEOUT_MS);
        // 默认 10s
        commitWindowMs = conf.getLong(MRJobConfig.MR_AM_COMMIT_WINDOW_MS,
            MRJobConfig.DEFAULT_MR_AM_COMMIT_WINDOW_MS);
        try {
            fs = FileSystem.get(conf);
            JobID id = TypeConverter.fromYarn(context.getApplicationID());
            JobId jobId = TypeConverter.toYarn(id);
            String user = UserGroupInformation.getCurrentUser().getShortUserName();
            startCommitFile = MRApps.getStartJobCommitFile(conf, user, jobId);
            endCommitSuccessFile = MRApps.getEndJobCommitSuccessFile(conf, user,
jobId);

```

```

        endCommitFailureFile = MRApps.getEndJobCommitFailureFile(conf, user,
jobId);
    } catch (IOException e) {
        throw new YarnRuntimeException(e);
    }
}
}

```

5.2.1.4 Task监听服务 TaskAttemptListenerImpl.serviceInit()

```

/**
 * This class is responsible for talking to the task umbilical.
 * It also converts all the old data structures
 * to yarn data structures.
 * <p>
 * This class HAS to be in this package to access package private
 * methods/classes.
 */
public class TaskAttemptListenerImpl extends CompositeService
    implements TaskUmbilicalProtocol, TaskAttemptListener {
    @Override
    protected void serviceInit(Configuration conf) throws Exception {
        // 注册 Task 心跳处理器 (也即 AM 与 Tasks 心跳)
        registerHeartbeatHandler(conf);
        // 默认 10s
        commitwindowMs = conf.getLong(MRJobConfig.MR_AM_COMMIT_WINDOW_MS,
            MRJobConfig.DEFAULT_MR_AM_COMMIT_WINDOW_MS);
        super.serviceInit(conf);
    }
}

```

```

    protected void registerHeartbeatHandler(Configuration conf) {
        // 创建 TaskHeartbeatHandler
        taskHeartbeatHandler = new TaskHeartbeatHandler(context.getEventHandler(),
            context.getClock(),
            conf.getInt(MRJobConfig.MR_AM_TASK_LISTENER_THREAD_COUNT,
                MRJobConfig.DEFAULT_MR_AM_TASK_LISTENER_THREAD_COUNT));
        // 添加 TaskHeartbeatHandler 服务
        addService(taskHeartbeatHandler);
    }
}

```

5.2.1.5 任务推测执行服务 DefaultSpeculator.serviceInit()

```

public class DefaultSpeculator extends AbstractService implements
    Speculator {
    // 调用其父类 AbstractService.serviceInit()
}

```

5.2.1.6 任务清除目录服务 StagingDirCleaningService.serviceInit()

```
private final class StagingDirCleaningService extends AbstractService {  
    // 调用其父类 AbstractService.serviceInit()  
}
```

5.2.1.7 容器申请者服务 ContainerAllocatorRouter.serviceInit()

```
/**  
 * By the time life-cycle of this router starts, job-init would have already  
 * happened.  
 */  
private final class ContainerAllocatorRouter extends AbstractService  
    implements ContainerAllocator, RMHeartbeatHandler {  
    // 调用其父类 AbstractService.serviceInit()  
}
```

5.2.1.8 启动容器服务 ContainerLauncherRouter.serviceInit()

```
/**  
 * By the time life-cycle of this router starts, job-init would have already  
 * happened.  
 */  
private final class ContainerLauncherRouter extends AbstractService  
    implements ContainerLauncher {  
    // 调用其父类 AbstractService.serviceInit()  
}
```

5.2.2 启动 MRAppMaster (调用 MRAppMaster.serviceStart())

```
@SuppressWarnings("unchecked")  
@Override  
protected void serviceStart() throws Exception {  
  
    /**  
     * MRAppMaster 组合服务的 AsyncDispatcher 异步事件分发器注册了哪些事件注册表  
     * 1 jobhistory.EventType      -> JobHistoryEventHandler (任务历史事件处理器)  
     * 2 JobEventType              -> JobEventDispatcher (任务事件分发器)  
     * 3 TaskEventType             -> TaskEventDispatcher (Task事件分发器)  
     * 4 TaskAttemptEventType      -> TaskAttemptEventDispatcher (Task 尝试事件分发  
器)  
     * 5 CommitterEventType        -> CommitterEventHandler (提交事件处理器)  
     * 6 Speculator.EventType      -> SpeculatorEventDispatcher (推测事件处理器)  
     * 7 ContainerAllocator.EventType -> ContainerAllocatorRouter (容器申请事件处理  
器)  
     * 8 ContainerLauncher.EventType -> ContainerLauncherRouter (启动容器事件处理  
器)  
     */  
  
    amInfos = new LinkedList<AMInfo>();
```



```

completedTasksFromPreviousRun = new HashMap<TaskId, TaskInfo>();
processRecovery();
cleanUpPreviousJobOutput();

// Current an AMInfo for the current AM generation.
// 创建当前 AM 信息 AMInfo
AMInfo amInfo =
    MRBuilderUtils.newAMInfo(appAttemptID, startTime, containerID,
nmHost,
        nmPort, nmHttpPort);

// ////////////////////////////////// Create the job itself.
job = createJob(getConfig(), forcedState, shutDownMessage);

// End of creating the job.

// Send out an MR AM init event for all previous AMs.
// 恢复上一个 AM (默认 amInfos 为空)
for (AMInfo info : amInfos) {
    dispatcher.getEventHandler().handle(
        new JobHistoryEvent(job.getID(), new AMStartedEvent(info
            .getAppAttemptId(), info.getStartTime(),
info.getContainerId(),
            info.getNodeManagerHost(), info.getNodeManagerPort(),
info
            .getNodeManagerHttpPort(), appSubmitTime)));
}

// Send out an MR AM init event for this AM.
// 调用 JobHistoryEventHandler.handle(AMStartedEvent)
dispatcher.getEventHandler().handle(
    // 创建 JobHistoryEvent
    new JobHistoryEvent(
        job.getID(),
        new AMStartedEvent(
            amInfo.getAppAttemptId(),
            amInfo.getStartTime(),
            amInfo.getContainerId(),
            amInfo.getNodeManagerHost(),
            amInfo.getNodeManagerPort(), amInfo
                .getNodeManagerHttpPort(),
            this.forcedState == null ? null :
this.forcedState.toString(), appSubmitTime)
        ));
    amInfos.add(amInfo);

// metrics system init is really init & start.
// It's more test friendly to put it here.
DefaultMetricsSystem.initialize("MRAppMaster");

boolean initFailed = false;
if (!errorHappenedShutDown) {

```

```

        // create a job event for job initialization
        JobEvent initJobEvent = new JobEvent(job.getID(),
JobEventType.JOB_INIT);
        // Send init to the job (this does NOT trigger job execution)
        // This is a synchronous call, not an event through dispatcher. We want
        // job-init to be done completely here.
        // 调用 JobEventDispatcher.handle(JobEventType.JOB_INIT)
        jobEventDispatcher.handle(initJobEvent);

        // If job is still not initialized, an error happened during
        // initialization. Must complete starting all of the services so failure
        // events can be processed.
        initFailed = (((JobImpl) job).getInternalState() !=
JobStateInternal.INITED);

        // JobImpl's InitTransition is done (call above is synchronous), so the
        // "uber-decision" (MR-1220) has been made. Query job and switch to
        // ubermode if appropriate (by registering different container-allocator
        // and container-launcher services/event-handlers).

        // 默认 false
        if (job.isUber()) {
            speculatorEventDispatcher.disableSpeculation();
            LOG.info("MRAppMaster uberizing job " + job.getID()
                + " in local container (\\"uber-AM\\") on node "
                + nmHost + ":" + nmPort + ".");
        } else {
            // send init to speculator only for non-uber jobs.
            // This won't yet start as dispatcher isn't started yet.
            // 调用
SpeculatorEventDispatcher.handle(Speculator.EventType.JOB_CREATE)
            dispatcher.getEventHandler().handle(
                new SpeculatorEvent(job.getID(), clock.getTime()));
            // MRAppMaster launching normal, non-uberized, multi-container
            // job job_1684656010852_0006
            LOG.info("MRAppMaster launching normal, non-uberized, multi-
container "
                + "job " + job.getID() + ".");
        }
        // Start ClientService here, since it's not initialized if
        // errorHappenedShutDown is true
        // 启动 MRClientService RPC 服务 其调用 serviceStart()
        clientService.start();
    }
    //start all the components
    /**
     * MRAppMaster 组合服务有哪些子服务? 调用其 serviceStart()
     * 1 异步事件分发器      -> AsyncDispatcher
     * 2 任务完成监控服务      -> TaskAttemptFinishingMonitor
     * 3 提交者事件处理        -> CommitterEventHandler
     * 4 Task监听服务          -> TaskAttemptListenerImpl
     * 5 任务推测执行服务      -> DefaultSpeculator

```

```

        * 6 任务清除目录服务      -> StagingDirCleaningService
        * 7 容器申请服务          -> ContainerAllocatorRouter
        * 8 启动容器服务          -> ContainerLauncherRouter
        * 9 任务历史事件处理器服务 -> JobHistoryEventHandler
    */
    super.serviceStart();

    // finally set the job classloader
    MRApps.setClassLoader(jobClassLoader, getConfig());

    if (initFailed) {
        JobEvent initFailedEvent = new JobEvent(job.getID(),
        JobEventType.JOB_INIT_FAILED);
        jobEventDispatcher.handle(initFailedEvent);
    } else {
        // All components have started, start the job.
        // 全部的组件已经启动 开始启动任务
        startJobs();
    }
}

```

5.2.2.1 启动 MRClientService RPC 服务 其调用 serviceStart()

```

// 该服务在 RMAppMaster.serviceInit() 创建
// service to handle requests from JobClient
// 创建 MRClientService 服务 该服务是与 JobClient 客户端交互(也即提交任务的客户端)
// JobClient 持续获取提交任务的状态信息
// 通讯协议接口为 MRClientProtocol
/**
 * This module is responsible for talking to the
 * jobclient (user facing).
 */
public class MRClientService extends AbstractService implements ClientService {
    public MRClientService(AppContext appContext) {
        super(MRClientService.class.getName());
        this.appContext = appContext;
        // 创建 MRClientProtocolHandler
        this.protocolHandler = new MRClientProtocolHandler();
    }

    protected void serviceStart() throws Exception {
        Configuration conf = getConfig();
        YarnRPC rpc = YarnRPC.create(conf);
        InetSocketAddress address = new InetSocketAddress(0);

        // 创建 RPC Server 通讯协议为 MRClientProtocol
        // 端口随机
        server =
            rpc.getServer(MRClientProtocol.class, protocolHandler, address,
                conf, appContext.getClientToAMTokenSecretManager(),
                conf.getInt(MRJobConfig.MR_AM_JOB_CLIENT_THREAD_COUNT,
                    MRJobConfig.DEFAULT_MR_AM_JOB_CLIENT_THREAD_COUNT),

```

```

MRJobConfig.MR_AM_JOB_CLIENT_PORT_RANGE);

// Enable service authorization?
if (conf.getBoolean(
    CommonConfigurationKeysPublic.HADOOP_SECURITY_AUTHORIZATION,
    false)) {
    refreshServiceAcls(conf, new MRAMPolicyProvider());
}

// 启动 RPC 服务
server.start();
this.bindAddress =
NetUtils.createSocketAddrForHost(appContext.getNMHostname(),
    server.getListenerAddress().getPort());
// Instantiated MRClientService at hadoop103/192.168.6.103:37605
LOG.info("Instantiated MRClientService at " + this.bindAddress);
try {
    // Explicitly disabling SSL for map reduce task as we can't allow MR
    users
    // to gain access to keystore file for opening SSL listener. We can
    trust
    // RM/NM to issue SSL certificates but definitely not MR-AM as it is
    // running in user-land.
    // 创建并启动 webApp
    webApp =
        WebApps.$for("mapreduce", AppContext.class, appContext, "ws")
            .withHttpPolicy(conf, Policy.HTTP_ONLY)
            .withPortRange(conf,
MRJobConfig.MR_AM_WEBAPP_PORT_RANGE)
            .start(new AMWebApp());
} catch (Exception e) {
    LOG.error("webapps failed to start. Ignoring for now:", e);
}
super.serviceStart();
}
}

```

5.2.2.2 异步事件分发器 AsyncDispatcher.serviceStart()

```

@Override
protected void serviceStart() throws Exception {
    //start all the components
    // 针对 ResourceManager 的 AsyncDispatcher 服务来说 啥也不干
    super.serviceStart();
    // 创建 EventHandler 线程
    eventHandlingThread = new Thread(createThread());
    eventHandlingThread.setName(dispatcherThreadName);
    eventHandlingThread.start();
}

```

5.2.2.3 任务完成监控服务 TaskAttemptFinishingMonitor.serviceStart()

```
// 调用其父类 AbstractLivelinessMonitor.serviceStart()
@Override
protected void serviceStart() throws Exception {
    assert !stopped : "starting when already stopped";
    // 重置 Timer
    resetTimer();
    // 创建并启动 PingChecker 线程
    checkerThread = new Thread(new PingChecker());
    checkerThread.setName("Ping Checker");
    checkerThread.start();
    super.serviceStart();
}
```

5.2.2.4 提交者事件处理 CommitterEventHandler.serviceStart()

```
@Override
protected void serviceStart() throws Exception {
    ThreadFactoryBuilder tfBuilder = new ThreadFactoryBuilder()
        .setNameFormat("CommitterEvent Processor #%d");
    // 默认为 null
    if (jobClassLoader != null) {
        // if the job classloader is enabled, we need to use the job classloader
        // as the thread context classloader (TCCL) of these threads in case the
        // committer needs to load another class via TCCL
        ThreadFactory backingTf = new ThreadFactory() {
            @Override
            public Thread newThread(Runnable r) {
                Thread thread = new Thread(r);
                thread.setContextClassLoader(jobClassLoader);
                return thread;
            }
        };
        tfBuilder.setThreadFactory(backingTf);
    }
    ThreadFactory tf = tfBuilder.build();
    // 创建线程池
    launcherPool = new HadoopThreadPoolExecutor(5, 5, 1,
        TimeUnit.HOURS, new LinkedBlockingQueue<Runnable>(), tf);
    // 创建并启动线程
    eventHandlingThread = new Thread(new Runnable() {
        @Override
        public void run() {
            CommitterEvent event = null;
            while (!stopped.get() && !Thread.currentThread().isInterrupted()) {
                try {
                    event = eventQueue.take();
                } catch (InterruptedException e) {
                    if (!stopped.get()) {
                        LOG.error("Returning, interrupted : " + e);
                    }
                }
            }
        }
    });
    eventHandlingThread.start();
}
```

```

        }
        return;
    }
    // the events from the queue are handled in parallel
    // using a thread pool
    // 执行
    launcherPool.execute(new EventProcessor(event));
}
}
});
eventHandlingThread.setName("CommitterEvent Handler");
eventHandlingThread.start();
super.serviceStart();
}

```

5.2.2.5 Task监听服务 TaskAttemptListenerImpl.serviceStart()

```

@Override
protected void serviceStart() throws Exception {
    // 启动 RPC Server
    startRpcServer();
    super.serviceStart();
}

```

```

protected void startRpcServer() {
    Configuration conf = getConfig();
    try {
        // 构建 RPC Server 通讯协议接口为 TaskUmbilicalProtocol 端口随机
        server = new RPC.Builder(conf).setProtocol(TaskUmbilicalProtocol.class)
            .setInstance(this).setBindAddress("0.0.0.0")
            .setPortRangeConfig(MRJobConfig.MR_AM_JOB_CLIENT_PORT_RANGE)
            .setNumHandlers(

conf.getInt(MRJobConfig.MR_AM_TASK_LISTENER_THREAD_COUNT,

MRJobConfig.DEFAULT_MR_AM_TASK_LISTENER_THREAD_COUNT))

.setVerbose(false).setSecretManager(jobTokenSecretManager).build();

        // Enable service authorization?
        if (conf.getBoolean(
            CommonConfigurationKeysPublic.HADOOP_SECURITY_AUTHORIZATION,
            false)) {
            refreshServiceAcls(conf, new MRAMPolicyProvider());
        }

        // 启动 RPC 服务
        server.start();

        // 解析 RPC Server 的 HOST PORT
        this.address = NetUtils.createSocketAddrForHost(

```

```

        context.getNMHostname(),
        server.getListenerAddress().getPort());
    } catch (IOException e) {
        throw new YarnRuntimeException(e);
    }
}

```

5.2.2.6 任务推测执行服务 DefaultSpeculator.serviceStart()

```

@Override
protected void serviceStart() throws Exception {
    // 创建并运行推测执行后台线程
    Runnable speculationBackgroundCore
        = new Runnable() {
        @Override
        public void run() {
            while (!stopped && !Thread.currentThread().isInterrupted()) {
                long backgroundRunStartTime = clock.getTime();
                try {
                    int speculations = computeSpeculations();
                    long mininumRecomp
                        = speculations > 0 ? soonestRetryAfterSpeculate
                        : soonestRetryAfterNoSpeculate;

                    long wait = Math.max(mininumRecomp,
                        clock.getTime() - backgroundRunStartTime);

                    if (speculations > 0) {
                        LOG.info("We launched " + speculations
                            + " speculations. sleeping " + wait + "
milliseconds.");
                    }

                    Object pollResult
                        = scanControl.poll(wait, TimeUnit.MILLISECONDS);
                } catch (InterruptedException e) {
                    if (!stopped) {
                        LOG.error("Background thread returning, interrupted",
e);
                    }
                    return;
                }
            }
        }
    };
    speculationBackgroundThread = new Thread
        (speculationBackgroundCore, "DefaultSpeculator background
processing");
    speculationBackgroundThread.start();

    super.serviceStart();
}

```

5.2.2.7 任务清除目录服务 StagingDirCleaningService.serviceStart()

```
// 调用其父类 AbstractService.serviceStart()
```

5.2.2.8 容器申请器服务 ContainerAllocatorRouter.serviceStart()

```
@Override
protected void serviceStart() throws Exception {
    if (job.isuber()) {
        MRApps.setupDistributedCacheLocal(getConfig());
        this.containerAllocator = new LocalContainerAllocator(
            this.clientService, this.context, nmHost, nmPort, nmHttpPort
            , containerID);
    } else {
        // 创建从 RM 申请容器服务 RMContainerAllocator
        this.containerAllocator = new RMContainerAllocator(
            this.clientService, this.context, preemptionPolicy);
    }
    // 调用 RMContainerAllocator.serviceInit()
    ((Service) this.containerAllocator).init(getConfig());
    // 调用 RMContainerAllocator.serviceStart()
    ((Service) this.containerAllocator).start();
    super.serviceStart();
}
```

5.2.2.8.1 创建从 RM 申请容器服务 RMContainerAllocator

```
/**
 * Allocates the container from the ResourceManager scheduler.
 */
public class RMContainerAllocator extends RMContainerRequestor
    implements ContainerAllocator {
    public RMContainerAllocator(ClientService clientService, AppContext context,
        AMPreemptionPolicy preemptionPolicy) {
        super(clientService, context);
        this.preemptionPolicy = preemptionPolicy;
        this.stopped = new AtomicBoolean(false);
        this.clock = context.getClock();
        // 创建 AssignedRequests
        this.assignedRequests = createAssignedRequests();
    }
}
```



```

/**
 * Keeps the data structures to send container requests to RM.
 */
public abstract class RMContainerRequestor extends RMCommunicator {
    public RMContainerRequestor(ClientService clientService, AppContext context) {
        super(clientService, context);
    }
}

```

```

/**
 * Registers/unregisters to RM and sends heartbeats to RM.
 */
public abstract class RMCommunicator extends AbstractService
    implements RMHeartbeatHandler {
    public RMCommunicator(ClientService clientService, AppContext context) {
        super("RMCommunicator");
        this.clientService = clientService;
        this.context = context;
        this.eventHandler = context.getEventHandler();
        this.applicationId = context.getApplicationID();
        this.stopped = new AtomicBoolean(false);
        this.heartbeatCallbacks = new ConcurrentLinkedQueue<Runnable>();
        this.schedulerResourceTypes = EnumSet.of(SchedulerResourceTypes.MEMORY);
    }
}

```

5.2.2.8.2 初始化调用 RMContainerAllocator.serviceInit()

```

@Override
protected void serviceInit(Configuration conf) throws Exception {
    // 调用父类
    super.serviceInit(conf);
    reduceSlowStart = conf.getFloat(
        MRJobConfig.COMPLETED_MAPS_FOR_REDUCE_SLOWSTART,
        DEFAULT_COMPLETED_MAPS_PERCENT_FOR_REDUCE_SLOWSTART);
    maxReduceRampupLimit = conf.getFloat(
        MRJobConfig.MR_AM_JOB_REDUCE_RAMPUP_UP_LIMIT,
        MRJobConfig.DEFAULT_MR_AM_JOB_REDUCE_RAMP_UP_LIMIT);
    maxReducePreemptionLimit = conf.getFloat(
        MRJobConfig.MR_AM_JOB_REDUCE_PREEMPTION_LIMIT,
        MRJobConfig.DEFAULT_MR_AM_JOB_REDUCE_PREEMPTION_LIMIT);
    reducerUnconditionalPreemptionDelayMs = 1000 * conf.getInt(
        MRJobConfig.MR_JOB_REDUCER_UNCONDITIONAL_PREEMPT_DELAY_SEC,
        MRJobConfig.DEFAULT_MR_JOB_REDUCER_UNCONDITIONAL_PREEMPT_DELAY_SEC);
    reducerNoHeadroomPreemptionDelayMs = conf.getInt(
        MRJobConfig.MR_JOB_REDUCER_PREEMPT_DELAY_SEC,
        MRJobConfig.DEFAULT_MR_JOB_REDUCER_PREEMPT_DELAY_SEC) * 1000; //sec -
    > ms
    maxRunningMaps = conf.getInt(MRJobConfig.JOB_RUNNING_MAP_LIMIT,
        MRJobConfig.DEFAULT_JOB_RUNNING_MAP_LIMIT);
}

```

```

        maxRunningReduces = conf.getInt(MRJobConfig.JOB_RUNNING_REDUCE_LIMIT,
            MRJobConfig.DEFAULT_JOB_RUNNING_REDUCE_LIMIT);
        RackResolver.init(conf);
        retryInterval =
getConfig().getLong(MRJobConfig.MR_AM_TO_RM_WAIT_INTERVAL_MS,
            MRJobConfig.DEFAULT_MR_AM_TO_RM_WAIT_INTERVAL_MS);
        mapNodeLabelExpression = conf.get(MRJobConfig.MAP_NODE_LABEL_EXP);
        reduceNodeLabelExpression = conf.get(MRJobConfig.REDUCE_NODE_LABEL_EXP);
        // Init startTime to current time. If all goes well, it will be reset after
        // first attempt to contact RM.
        retrystartTime = System.currentTimeMillis();
        this.scheduledRequests.setNumOpportunisticMapsPercent(
            conf.getInt(MRJobConfig.MR_NUM_OPPORTUNISTIC_MAPS_PERCENT,
                MRJobConfig.DEFAULT_MR_NUM_OPPORTUNISTIC_MAPS_PERCENT));
        // 0% of the mappers will be scheduled using OPPORTUNISTIC containers
        LOG.info(this.scheduledRequests.getNumOpportunisticMapsPercent() +
            "% of the mappers will be scheduled using OPPORTUNISTIC
containers");
    }

```

```

@Override
protected void serviceInit(Configuration conf) throws Exception {
    // 调用父类
    super.serviceInit(conf);
    nodeBlacklistingEnabled =
        conf.getBoolean(MRJobConfig.MR_AM_JOB_NODE_BLACKLISTING_ENABLE,
true);
    // nodeBlacklistingEnabled:true
    LOG.info("nodeBlacklistingEnabled:" + nodeBlacklistingEnabled);
    maxTaskFailuresPerNode =
        conf.getInt(MRJobConfig.MAX_TASK_FAILURES_PER_TRACKER, 3);
    blacklistDisablePercent =
        conf.getInt(
MRJobConfig.MR_AM_IGNORE_BLACKLISTING_BLACKLISTED_NODE_PERCENT,
MRJobConfig.DEFAULT_MR_AM_IGNORE_BLACKLISTING_BLACKLISTED_NODE_PERCENT);
    // maxTaskFailuresPerNode is 3
    LOG.info("maxTaskFailuresPerNode is " + maxTaskFailuresPerNode);
    if (blacklistDisablePercent < -1 || blacklistDisablePercent > 100) {
        throw new YarnRuntimeException("Invalid blacklistDisablePercent: "
            + blacklistDisablePercent
            + ". Should be an integer between 0 and 100 or -1 to disabled");
    }
    // blacklistDisablePercent is 33
    LOG.info("blacklistDisablePercent is " + blacklistDisablePercent);
}

```

5.2.2.8.3 启动调用 RMContainerAllocator.serviceStart()

```
@Override
protected void serviceStart() throws Exception {
    // 创建线程并启动
    this.eventHandlingThread = new Thread() {
        @SuppressWarnings("unchecked")
        @Override
        public void run() {

            ContainerAllocatorEvent event;

            while (!stopped.get() && !Thread.currentThread().isInterrupted()) {
                try {
                    // 拉取 event
                    event = RMContainerAllocator.this.eventQueue.take();
                } catch (InterruptedException e) {
                    if (!stopped.get()) {
                        LOG.error("Returning, interrupted : " + e);
                    }
                    return;
                }

                try {
                    // 处理 event
                    handleEvent(event);
                } catch (Throwable t) {
                    LOG.error("Error in handling event type " + event.getType()
                        + " to the ContainerAllocator", t);
                    // Kill the AM
                    eventHandler.handle(new JobEvent(getJob().getID(),
                        JobEventType.INTERNAL_ERROR));
                    return;
                }
            }
        }
    };
    this.eventHandlingThread.start();
    // 调用父类
    super.serviceStart();
}
```

```
@Override
protected void serviceStart() throws Exception {
    // 创建 ResourceManager 的 ApplicationMasterService RPC 服务的客户端代理对象
    // 通讯协议接口为 ApplicationMasterProtocol
    scheduler = createSchedulerProxy();
    JobID id = TypeConverter.fromYarn(this.applicationId);
    JobId jobId = TypeConverter.toYarn(id);
    job = context.getJob(jobId);
    // 向 ResourceManager 的 ApplicationMasterService RPC 服务注册 AM
}
```

```

register();
// 启动容器申请线程 AllocatorRunnable 也即发送心跳
startAllocatorThread();
super.serviceStart();
}

```

5.2.2.9 启动容器服务 ContainerLauncherRouter.serviceStart()

```

@Override
protected void serviceStart() throws Exception {
    if (job.isuber()) {
        this.containerLauncher = new LocalContainerLauncher(context,
            (TaskUmbilicalProtocol) taskAttemptListener,
            jobClassLoader);
        ((LocalContainerLauncher) this.containerLauncher)
            .setEncryptedSpillKey(encryptedSpillKey);
    } else {
        // 创建启动容器服务 ContainerLauncherImpl
        this.containerLauncher = new ContainerLauncherImpl(context);
    }
    // 调用 ContainerLauncherImpl.serviceInit()
    ((Service) this.containerLauncher).init(getConfig());
    // 调用 ContainerLauncherImpl.serviceStart()
    ((Service) this.containerLauncher).start();
    // 调用父类
    super.serviceStart();
}

```

5.2.2.9.1 创建启动容器服务 ContainerLauncherImpl

```

/**
 * This class is responsible for launching of containers.
 */
public class ContainerLauncherImpl extends AbstractService implements
    ContainerLauncher {
    public ContainerLauncherImpl(AppContext context) {
        super(ContainerLauncherImpl.class.getName());
        this.context = context;
        this.stopped = new AtomicBoolean(false);
    }
}

```

5.2.2.9.2 初始化调用 ContainerLauncherImpl.serviceInit()

```

@Override
protected void serviceInit(Configuration conf) throws Exception {
    // 默认 500
    this.limitOnPoolSize = conf.getInt(
        MRJobConfig.MR_AM_CONTAINERLAUNCHER_THREAD_COUNT_LIMIT,
        MRJobConfig.DEFAULT_MR_AM_CONTAINERLAUNCHER_THREAD_COUNT_LIMIT);
    // Upper limit on the thread pool size is 500
}

```

```

        LOG.info("Upper limit on the thread pool size is " + this.limitOnPoolSize);

        this.initialPoolSize = conf.getInt(
            MRJobConfig.MR_AM_CONTAINERLAUNCHER_THREADPOOL_INITIAL_SIZE,
            MRJobConfig.DEFAULT_MR_AM_CONTAINERLAUNCHER_THREADPOOL_INITIAL_SIZE);
        // The thread pool initial size is 10
        LOG.info("The thread pool initial size is " + this.initialPoolSize);

        super.serviceInit(conf);
        // 创建 ContainerManagementProtocolProxy
        cmProxy = new ContainerManagementProtocolProxy(conf);
    }

```

```

/**
 * Helper class to manage container manager proxies
 */
@LimitedPrivate({"MapReduce", "YARN"})
public class ContainerManagementProtocolProxy {
    public ContainerManagementProtocolProxy(Configuration conf,
                                           NMTokenCache nmTokenCache) {

        this.conf = new Configuration(conf);
        this.nmTokenCache = nmTokenCache;

        // 默认 0
        maxConnectedNMs =
            conf.getInt(YarnConfiguration.NM_CLIENT_MAX_NM_PROXIES,
                YarnConfiguration.DEFAULT_NM_CLIENT_MAX_NM_PROXIES);
        if (maxConnectedNMs < 0) {
            throw new YarnRuntimeException(
                YarnConfiguration.NM_CLIENT_MAX_NM_PROXIES
                + " (" + maxConnectedNMs + ") can not be less than 0.");
        }

        if (LOG.isDebugEnabled()) {
            LOG.debug(YarnConfiguration.NM_CLIENT_MAX_NM_PROXIES + " : " +
                maxConnectedNMs);
        }

        if (maxConnectedNMs > 0) {
            cmProxy =
                new LinkedHashMap<String, ContainerManagementProtocolProxyData>
();
        } else {
            cmProxy = Collections.emptyMap();
            // Connections are not being cached so ensure connections close quickly
            // to avoid creating thousands of RPC client threads on large clusters.
            this.conf.setInt(
                CommonConfigurationKeysPublic.IPC_CLIENT_CONNECTION_MAXIDLETIME_KEY,
                0);

```



```

        // later is just a buffer so we are not always
        increasing the
        // pool-size
        int newPoolSize = Math.min(limitOnPoolSize,
        idealPoolSize
        + initialPoolSize);
        LOG.info("Setting ContainerLauncher pool size to " +
        newPoolSize
        + " as number-of-nodes to talk to is " +
        numNodes);
        launcherPool.setCorePoolSize(newPoolSize);
    }
}

// the events from the queue are handled in parallel
// using a thread pool
launcherPool.execute(createEventProcessor(event));

// TODO: Group launching of multiple containers to a single
// NodeManager into a single connection
    }
}
};
eventHandlingThread.setName("ContainerLauncher Event Handler");
eventHandlingThread.start();
super.serviceStart();
}

```

5.4 AM 的所有组件基本启动完成开始执行任务

```

// All components have started, start the job.
// 全部的组件已经启动 开始启动任务
startJobs();

```

```

/**
 * This can be overridden to instantiate multiple jobs and create a
 * workflow.
 * <p>
 * TODO: Rework the design to actually support this. Currently much of the
 * job stuff has been moved to init() above to support uberization (MR-1220).
 * In a typical workflow, one presumably would want to uberize only a subset
 * of the jobs (the "small" ones), which is awkward with the current design.
 */
@SuppressWarnings("unchecked")
protected void startJobs() {
    /** create a job-start event to get this ball rolling */
    // 创建 job-start 事件
    // eventType = JobEventType.JOB_START
    JobEvent startJobEvent = new JobStartEvent(job.getID(),
        recoveredJobStartTime);
}

```

```

    /** send the job-start event. this triggers the job execution. */
    // 调用 JobEventDispatcher.handle(JobEventType.JOB_START)
    dispatcher.getEventHandler().handle(startJobEvent);
}

```

5.4.1 启动 MR 任务创建待执行 MapTask、ReduceTask 的源码分析

```

// 在启动 MRAppMaster.serviceStart() 时 触发了一个 JobEventType.JOB_INIT
// 核心源码入口如下：
if (!errorHappenedShutDown) {
    // create a job event for job initialization
    // 初始化 Job (也即创建 MapTask ReduceTask 等待后续被执行调度)
    JobEvent initJobEvent = new JobEvent(job.getID(),
JobEventType.JOB_INIT);
    // Send init to the job (this does NOT trigger job execution)
    // This is a synchronous call, not an event through dispatcher. We want
    // job-init to be done completely here.
    // 调用 JobEventDispatcher.handle(JobEventType.JOB_INIT)
    jobEventDispatcher.handle(initJobEvent);
.....}

```

```

// 经过各种状态转换 最终来到 InitTransition.doTransition(JobEventType.JOB_INIT)
public static class InitTransition
    implements MultipleArcTransition<JobImpl, JobEvent, JobStateInternal> {

    /**
     * Note that this transition method is called directly (and synchronously)
     * by MRAppMaster's init() method (i.e., no RPC, no thread-switching;
     * just plain sequential call within AM context), so we can trigger
     * modifications in AM state from here (at least, if AM is written that
     * way; MR version is).
     */
    @Override
    public JobStateInternal transition(JobImpl job, JobEvent event) {
        // eventType = JobEventType.JOB_INIT
        job.metrics.submittedJob(job);
        job.metrics.preparingJob(job);

        if (job.newApiCommitter) {
            // 创建 JobImpl 上下文 JobContextImpl
            job.jobContext = new JobContextImpl(job.conf,
                job.oldJobId);
        } else {
            job.jobContext = new org.apache.hadoop.mapred.JobContextImpl(
                job.conf, job.oldJobId);
        }

        try {
            setup(job);
            job.fs = job.getFileSystem(job.conf);

```



```

// log to job history
JobSubmittedEvent jse = new JobSubmittedEvent(job.oldJobId,
    job.conf.get(MRJobConfig.JOB_NAME, "test"),
    job.conf.get(MRJobConfig.USER_NAME, "mapred"),
    job.appSubmitTime,
    job.remoteJobConfFile.toString(),
    job.jobACLs, job.queueName,
    job.conf.get(MRJobConfig.WORKFLOW_ID, ""),
    job.conf.get(MRJobConfig.WORKFLOW_NAME, ""),
    job.conf.get(MRJobConfig.WORKFLOW_NODE_NAME, ""),
    getWorkflowAdjacencies(job.conf),
    job.conf.get(MRJobConfig.WORKFLOW_TAGS, ""), job.conf);
job.eventHandler.handle(new JobHistoryEvent(job.jobId, jse));
//TODO JH Verify jobACLs, UserName via UGI?

// 获取 MR 任务切片元数据信息（也即读取 HDFS 的切片信息）
TaskSplitMetaInfo[] taskSplitMetaInfo = createsplits(job,
job.jobId);

// 切片有多少个 MapTask 就有几个
job.numMapTasks = taskSplitMetaInfo.length;
// 默认 ReduceTask 为 0
job.numReduceTasks = job.conf.getInt(MRJobConfig.NUM_REDUCES, 0);

if (job.numMapTasks == 0 && job.numReduceTasks == 0) {
    job.addDiagnostic("No of maps and reduces are 0 " + job.jobId);
} else if (job.numMapTasks == 0) {
    job.reduceweight = 0.9f;
} else if (job.numReduceTasks == 0) {
    job.mapweight = 0.9f;
} else {
    job.mapweight = job.reduceweight = 0.45f;
}

checkTaskLimits();

long inputLength = 0;
for (int i = 0; i < job.numMapTasks; ++i) {
    inputLength += taskSplitMetaInfo[i].getInputDataLength();
}

job.makeUberDecision(inputLength);

job.taskAttemptCompletionEvents =
    new ArrayList<TaskAttemptCompletionEvent>(
        job.numMapTasks + job.numReduceTasks + 10);
job.mapAttemptCompletionEvents =
    new ArrayList<TaskCompletionEvent>(job.numMapTasks + 10);
job.taskCompletionIdxToMapCompletionIdx = new ArrayList<Integer>(
    job.numMapTasks + job.numReduceTasks + 10);

job.allowedMapFailuresPercent =

```

```

        job.conf.getInt(MRJobConfig.MAP_FAILURES_MAX_PERCENT, 0);
        job.allowedReduceFailuresPercent =
            job.conf.getInt(MRJobConfig.REDUCE_FAILURES_MAXPERCENT, 0);

        cleanupSharedCacheUploadPolicies(job.conf);

        // create the Tasks but don't start them yet
        // 创建 MapTask、ReduceTask 任务 (重要、重要、重要)
        createMapTasks(job, inputLength, taskSplitMetaInfo);
        createReduceTasks(job);

        job.metrics.endPreparingJob(job);
        // 返回 JobStateInternal.INITED
        return JobStateInternal.INITED;
    } catch (Exception e) {
        LOG.warn("Job init failed", e);
        job.metrics.endPreparingJob(job);
        job.addDiagnostic("Job init failed : "
            + StringUtils.stringifyException(e));
        // Leave job in the NEW state. The MR AM will detect that the state
is
        // not INITED and send a JOB_INIT_FAILED event.
        return JobStateInternal.NEW;
    }
}

```

```

// 创建 MapTaskImpl 并缓存
private void createMapTasks(JobImpl job, long inputLength,
    TaskSplitMetaInfo[] splits) {
    for (int i = 0; i < job.numMapTasks; ++i) {
        // 创建 MapTaskImpl (也即 MapTask)
        TaskImpl task =
            new MapTaskImpl(job.jobId, i,
                job.eventHandler,
                job.remoteJobConfFile,
                job.conf, splits[i],
                job.taskAttemptListener,
                job.jobToken, job.jobCredentials,
                job.clock,
                job.applicationAttemptId.getAttemptId(),
                job.metrics, job.appContext);

        // 添加待执行 MapTask
        job.addTask(task);
    }
    // Input size for job job_1684656010852_0006 = 147145. Number of splits
= 1
    LOG.info("Input size for job " + job.jobId + " = " + inputLength
        + ". Number of splits = " + splits.length);
}

```

```

// 创建 ReduceTaskImpl 并缓存

```

```

private void createReduceTasks(JobImpl job) {
    for (int i = 0; i < job.numReduceTasks; i++) {
        TaskImpl task =
            new ReduceTaskImpl(job.jobId, i,
                job.eventHandler,
                job.remoteJobConfFile,
                job.conf, job.numMapTasks,
                job.taskAttemptListener, job.jobToken,
                job.jobCredentials, job.clock,
                job.applicationAttemptId.getAttemptId(),
                job.metrics, job.appContext);
        job.addTask(task);
    }
    // Number of reduces for job job_1684656010852_0006 = 1
    LOG.info("Number of reduces for job " + job.jobId + " = "
        + job.numReduceTasks);
}

```

5.4.2 开始执行任务经过多种状态转换最终调用

SetupCompletedTransition.doTransition() 执行 MapTask或者 ReduceTask

```

private static class SetupCompletedTransition
    implements SingleArcTransition<JobImpl, JobEvent> {
    @Override
    public void transition(JobImpl job, JobEvent event) {
        // eventType = JobEventType.JOB_SETUP_COMPLETED
        job.setupProgress = 1.0f;
        // 调度执行 MapTask (缓存 MapTasks 在 JobImpl.addTask())
        job.scheduleTasks(job.mapTasks, job.numReduceTasks == 0);
        // 调度执行 ReduceTask (缓存 ReduceTasks 在 JobImpl.addTask())
        job.scheduleTasks(job.reduceTasks, true);

        // If we have no tasks, just transition to job completed
        // 如果没有 Task 则表示任务执行完成
        if (job.numReduceTasks == 0 && job.numMapTasks == 0) {
            // 调用 JobEventDispatcher.handle(JobEventType.JOB_COMPLETED)
            job.eventHandler.handle(new JobEvent(job.jobId,
                JobEventType.JOB_COMPLETED));
        }
    }
}

```

```

protected void scheduleTasks(Set<TaskId> taskIDs,
    boolean recoverTaskOutput) {
    for (TaskId taskID : taskIDs) {
        TaskInfo taskInfo = completedTasksFromPreviousRun.remove(taskID);
        if (taskInfo != null) {
            eventHandler.handle(new TaskRecoverEvent(taskID, taskInfo,
                committer, recoverTaskOutput));
        } else {

```

```

        // 处理启动 MapTask 任务
        // 处理启动 ReduceTask 任务
        // 调用 TaskEventDispatcher.handle(TaskEventType.T_SCHEDULE))
        eventHandler.handle(new TaskEvent(taskID,
TaskEventType.T_SCHEDULE));
    }
}
}

```

5.4.3 通过各种状态转换之后向 RM 的 ContainerManagerImpl RPC 服务申请容器启动MapTask或者ReduceTask源码分析

```

// 最终调用 RMContainerAllocator.handle(ContainerAllocator.EventType.CONTAINER_REQ) 处
// 理容器申请
@Override
public void handle(ContainerAllocatorEvent event) {
    // eventType = ContainerAllocator.EventType.CONTAINER_REQ
    int qSize = eventQueue.size();
    if (qSize != 0 && qSize % 1000 == 0) {
        LOG.info("Size of event-queue in RMContainerAllocator is " + qSize);
    }
    int remCapacity = eventQueue.remainingCapacity();
    if (remCapacity < 1000) {
        LOG.warn("Very low remaining capacity in the event-queue "
            + "of RMContainerAllocator: " + remCapacity);
    }
    try {
        // 添加申请容器事件
        // 调用 serviceStart() 启动线程的 run()
        eventQueue.put(event);
    } catch (InterruptedException e) {
        throw new YarnRuntimeException(e);
    }
}
}

```

```

@Override
protected void serviceStart() throws Exception {
    // 创建线程并启动
    this.eventHandlingThread = new Thread() {
        @SuppressWarnings("unchecked")
        @Override
        public void run() {

            ContainerAllocatorEvent event;
            // eventType =

            while (!stopped.get() && !Thread.currentThread().isInterrupted()) {
                try {
                    // 拉取 event (申请容器启动 MapTask 或者 ReduceTask)
                    event = RMContainerAllocator.this.eventQueue.take();

```

```

    } catch (InterruptedException e) {
        if (!stopped.get()) {
            LOG.error("Returning, interrupted : " + e);
        }
        return;
    }

    try {
        // 处理申请容器 event
        handleEvent(event);
    } catch (Throwable t) {
        LOG.error("Error in handling event type " + event.getType()
            + " to the ContainerAllocator", t);
        // Kill the AM
        eventHandler.handle(new JobEvent(getJob().getID(),
            JobEventType.INTERNAL_ERROR));
        return;
    }
}

};
this.eventHandlingThread.start();
// 调用父类
super.serviceStart();
}

```

```

protected synchronized void handleEvent(ContainerAllocatorEvent event) {
    recalculateReduceSchedule = true;
    // eventType = ContainerAllocator.EventType.CONTAINER_REQ
    if (event.getType() == ContainerAllocator.EventType.CONTAINER_REQ) {
        ContainerRequestEvent reqEvent = (ContainerRequestEvent) event;
        boolean isMap = reqEvent.getAttemptID().getTaskId().getTaskType().
            equals(TaskType.MAP);
        if (isMap) {
            // 处理申请启动 MapTask 容器请求
            handleMapContainerRequest(reqEvent);
        } else {
            // 处理申请启动 ReduceTask 容器请求
            handleReduceContainerRequest(reqEvent);
        }

    } else if (
        event.getType() ==
        ContainerAllocator.EventType.CONTAINER_DEALLOCATE) {

        LOG.info("Processing the event " + event.toString());

        TaskAttemptId aId = event.getAttemptID();

        boolean removed = scheduledRequests.remove(aId);
        if (!removed) {

```

```

        ContainerId containerId = assignedRequests.get(aId);
        if (containerId != null) {
            removed = true;
            assignedRequests.remove(aId);
            containersReleased++;
            pendingRelease.add(containerId);
            release(containerId);
        }
    }
    if (!removed) {
        LOG.error("Could not deallocate container for task attemptId " +
            aId);
    }
    preemptionPolicy.handleCompletedContainer(event.getAttemptID());
} else if (
    event.getType() == ContainerAllocator.EventType.CONTAINER_FAILED) {
    ContainerFailedEvent fEv = (ContainerFailedEvent) event;
    String host = getHost(fEv.getContMgrAddress());
    containerFailedOnHost(host);
    // propagate failures to preemption policy to discard checkpoints for
    // failed tasks
    preemptionPolicy.handleFailedContainer(event.getAttemptID());
}
}

```

5.4.3.1 处理 MapTask 容器请求

```

@SuppressWarnings({"unchecked"})
private void handleMapContainerRequest(ContainerRequestEvent reqEvent) {
    assert (reqEvent.getAttemptID().getTaskId().getTaskType().equals(
        TaskType.MAP));

    // 获取容器可以申请的最大资源配置
    Resource supportedMaxContainerCapability = getMaxContainerCapability();
    JobId jobId = getJob().getID();

    if (mapResourceRequest.equals(Resources.none())) {
        mapResourceRequest = reqEvent.getCapability();
        eventHandler.handle(new JobHistoryEvent(jobId,
            new NormalizedResourceEvent(
                org.apache.hadoop.mapreduce.TaskType.MAP,
                mapResourceRequest.getMemorySize())));
        // 默认容器申请资源 <memory:1024, vCores:1>
        LOG.info("mapResourceRequest:" + mapResourceRequest);
    }

    boolean mapContainerRequestAccepted = true;
    // 判断启动 MapTask 容器的资源是否大于容器的最大申请资源
    if (mapResourceRequest.getMemorySize() >
        supportedMaxContainerCapability.getMemorySize()
        ||
        mapResourceRequest.getVirtualCores() >

```

```

        supportedMaxContainerCapability.getVirtualCores()) {
    mapContainerRequestAccepted = false;
}

// mapContainerRequestAccepted = true 表示申请启动 MapTask 容器
// 的资源没有大于容器的最大申请资源
// 默认容器申请资源为 <memory:1024, vCores:1>
if (mapContainerRequestAccepted) {
    // set the resources
    // 申请容器请求设置资源
    reqEvent.getCapability().setMemorySize(
        mapResourceRequest.getMemorySize());
    reqEvent.getCapability().setVirtualCores(
        mapResourceRequest.getVirtualCores());
    // 添加请求
    scheduledRequests.addMap(reqEvent); //maps are immediately scheduled
} else {
    String diagMsg = "The required MAP capability is more than the " +
        "supported max container capability in the cluster. Killing" +
        " the Job. mapResourceRequest: " + mapResourceRequest +
        " maxContainerCapability:" + supportedMaxContainerCapability;
    LOG.info(diagMsg);
    eventHandler.handle(new JobDiagnosticsUpdateEvent(jobId, diagMsg));
    eventHandler.handle(new JobEvent(jobId, JobEventType.JOB_KILL));
}
}

```

```

void addMap(ContainerRequestEvent event) {
    ContainerRequest request = null;

    if (event.getEarlierAttemptFailed()) {
        earlierFailedMaps.add(event.getAttemptID());
        request =
            new ContainerRequest(event, PRIORITY_FAST_FAIL_MAP,
                mapNodeLabelExpression);
        LOG.info("Added " + event.getAttemptID() + " to list of failed
maps");

        // If its an earlier Failed attempt, do not retry as OPPORTUNISTIC
        maps.put(event.getAttemptID(), request);
        addContainerReq(request);
    } else {
        if (mapsMod100 < numOpportunisticMapsPercent) {
            request =
                new ContainerRequest(event, PRIORITY_OPPORTUNISTIC_MAP,
                    mapNodeLabelExpression);
            maps.put(event.getAttemptID(), request);
            addOpportunisticResourceRequest(request.priority,
request.capability);
        } else {
            // 创建容器请求 ContainerRequest
            request =

```

```

        new ContainerRequest(event, PRIORITY_MAP,
mapNodeLabelExpression);
    // 申请容器优先级 本地优先
    for (String host : event.getHosts()) {
        LinkedList<TaskAttemptId> list = mapsHostMapping.get(host);
        if (list == null) {
            list = new LinkedList<TaskAttemptId>();
            mapsHostMapping.put(host, list);
        }
        list.add(event.getAttemptID());
        if (LOG.isDebugEnabled()) {
            LOG.debug("Added attempt req to host " + host);
        }
    }
    // 申请容器优先级 其次机架
    for (String rack : event.getRacks()) {
        LinkedList<TaskAttemptId> list = mapsRackMapping.get(rack);
        if (list == null) {
            list = new LinkedList<TaskAttemptId>();
            mapsRackMapping.put(rack, list);
        }
        list.add(event.getAttemptID());
        if (LOG.isDebugEnabled()) {
            LOG.debug("Added attempt req to rack " + rack);
        }
    }
    // 缓存 AttemptID -> ContainerRequest
    maps.put(event.getAttemptID(), request);
    // 添加申请容器请求
    addContainerReq(request);
}
mapsMod100++;
mapsMod100 %= 100;
}
}

```

```

protected void addContainerReq(ContainerRequest req) {
    // Create resource requests
    for (String host : req.hosts) {
        // Data-local
        // 数据本地性 (判断启动 Task 容器的 host 不在 NM 黑名单)
        if (!isNodeBlacklisted(host)) {
            addResourceRequest(req.priority, host, req.capability,
                null);
        }
    }

    // Nothing Rack-local for now
    // 机架
    for (String rack : req.racks) {
        addResourceRequest(req.priority, rack, req.capability,

```



```

        null);
    }

    // Off-switch
    addResourceRequest(req.priority, ResourceRequest.ANY, req.capability,
        req.nodeLabelExpression);
}

```

```

private void addResourceRequest(Priority priority, String resourceName,
                                Resource capability, String nodeLabelExpression)
{
    // 添加资源(容器)申请
    addResourceRequest(priority, resourceName, capability, nodeLabelExpression,
        ExecutionType.GUARANTEED);
}

```

```

private void addResourceRequest(Priority priority, String resourceName,
                                Resource capability, String nodeLabelExpression,
                                ExecutionType executionType) {
    Map<String, Map<Resource, ResourceRequest>> remoteRequests =
        this.remoteRequestsTable.get(priority);
    if (remoteRequests == null) {
        remoteRequests = new HashMap<String, Map<Resource, ResourceRequest>>();
        this.remoteRequestsTable.put(priority, remoteRequests);
        if (LOG.isDebugEnabled()) {
            LOG.debug("Added priority=" + priority);
        }
    }
    Map<Resource, ResourceRequest> reqMap = remoteRequests.get(resourceName);
    if (reqMap == null) {
        reqMap = new HashMap<Resource, ResourceRequest>();
        remoteRequests.put(resourceName, reqMap);
    }
    ResourceRequest remoteRequest = reqMap.get(capability);
    if (remoteRequest == null) {
        remoteRequest = recordFactory.newRecordInstance(ResourceRequest.class);
        remoteRequest.setPriority(priority);
        remoteRequest.setResourceName(resourceName);
        remoteRequest.setCapability(capability);
        remoteRequest.setNumContainers(0);
        remoteRequest.setNodeLabelExpression(nodeLabelExpression);
        remoteRequest.setExecutionTypeRequest(
            ExecutionTypeRequest.newInstance(executionType, true));
        reqMap.put(capability, remoteRequest);
    }
    remoteRequest.setNumContainers(remoteRequest.getNumContainers() + 1);

    // Note this down for next interaction with ResourceManager
    // 申请容器请求 (将 ResourceRequest 请求缓存等待其他线程拉取)
    // 调用 AllocatorRunnable.run() 发送心跳给 RM 的 ApplicationMasterService 申请容

```

```

// 最终调用 RMContainerRequestor.makeRemoteRequest()
addResourceRequestToAsk(remoteRequest);
if (LOG.isDebugEnabled()) {
    LOG.debug("addResourceRequest:" + " applicationId="
        + applicationId.getId() + " priority=" + priority.getPriority()
        + " resourceName=" + resourceName + " numContainers="
        + remoteRequest.getNumContainers() + " #asks=" + ask.size());
}
}

```

```

private void addResourceRequestToAsk(ResourceRequest remoteRequest) {
    // because objects inside the resource map can be deleted ask can end up
    // containing an object that matches new resource object but with different
    // numContainers. So existing values must be replaced explicitly
    ask.remove(remoteRequest);
    ask.add(remoteRequest);
}

```

5.4.3.1.1 AM 申请容器启动 MapTask 容器是通过发送心跳向 RM 的 ContainerManagerImpl RPC 服务申请资源(调用ApplicationMasterService.allocate())

```

// 调用 AllocatorRunnable.run() 发送心跳给 RM 的 ApplicationMasterService 申请容器
@VisibleForTesting
public class AllocatorRunnable implements Runnable {
    @Override
    public void run() {
        while (!stopped.get() && !Thread.currentThread().isInterrupted()) {
            try {
                // 默认睡眠 1s
                Thread.sleep(rmPollInterval);
                try {
                    // 心跳 调用 RMContainerAllocator.heartbeat()
                    heartbeat();
                } catch (RMContainerAllocationException e) {
                    LOG.error("Error communicating with RM: " + e.getMessage(),
e);

                    return;
                } catch (Exception e) {
                    LOG.error("ERROR IN CONTACTING RM. ", e);
                    continue;
                    // TODO: for other exceptions
                }

                lastHeartbeatTime = context.getClock().getTime();

                // 执行心跳回调
                executeHeartbeatCallbacks();
            } catch (InterruptedException e) {
                if (!stopped.get()) {
                    LOG.warn("Allocated thread interrupted. Returning.");
                }
            }
        }
    }
}

```

```

        return;
    }
}
}
}

```

```

@Override
protected synchronized void heartbeat() throws Exception {
    scheduleStats.updateAndLogIfChanged("Before Scheduling: ");
    // 获取 AM 基于分配的容器 (里面发送 RPC 请求申请启动 Task 容器)
    // 也即 allocatedContainers 是已经申请好的容器资源
    List<Container> allocatedContainers = getResources();
    if (allocatedContainers != null && allocatedContainers.size() > 0) {
        // 分配容器启动 Task
        scheduledRequests.assign(allocatedContainers);
    }

    // 获取已完成 MAP 任务数
    int completedMaps = getJob().getCompletedMaps();
    // 获取已经完成任务数
    int completedTasks = completedMaps + getJob().getCompletedReduces();

    if ((lastCompletedTasks != completedTasks) ||
        (scheduledRequests.maps.size() > 0)) {
        lastCompletedTasks = completedTasks;
        recalculateReduceSchedule = true;
    }

    if (recalculateReduceSchedule) {
        boolean reducerPreempted = preemptReducesIfNeeded();

        if (!reducerPreempted) {
            // Only schedule new reducers if no reducer preemption happens for
            // this heartbeat
            schedulerReduces(getJob().getTotalMaps(), completedMaps,
                             scheduledRequests.maps.size(),
                             scheduledRequests.reduces.size(),
                             assignedRequests.maps.size(),
                             assignedRequests.reduces.size(),
                             mapResourceRequest, reduceResourceRequest,
                             pendingReduces.size(),
                             maxReduceRampupLimit, reducesSlowStart);
        }

        recalculateReduceSchedule = false;
    }

    scheduleStats.updateAndLogIfChanged("After Scheduling: ");
}

```

```

@SuppressWarnings("unchecked")

```

```

private List<Container> getResources() throws Exception {
    applyConcurrentTaskLimits();

    // will be null the first time
    Resource headRoom = Resources.clone(getAvailableResources());
    AllocateResponse response;
    /*
     * If contact with RM is lost, the AM will wait MR_AM_TO_RM_WAIT_INTERVAL_MS
     * milliseconds before aborting. During this interval, AM will still try
     * to contact the RM.
     */
    try {
        // 发送 RPC 请求申请容器并返回
        response = makeRemoteRequest();
        // Reset retry count if no exception occurred.
        retrystartTime = System.currentTimeMillis();
    } catch (ApplicationAttemptNotFoundException e) {
        // This can happen if the RM has been restarted. If it is in that state,
        // this application must clean itself up.
        eventHandler.handle(new JobEvent(this.getJob().getID(),
            JobEventType.JOB_AM_REBOOT));
        throw new RMContainerAllocationException(
            "Resource Manager doesn't recognize AttemptId: "
                + this.getContext().getApplicationAttemptId(), e);
    } catch (ApplicationMasterNotRegisteredException e) {
        LOG.info("ApplicationMaster is out of sync with ResourceManager,"
            + " hence resync and send outstanding requests.");
        // RM may have restarted, re-register with RM.
        lastResponseID = 0;
        register();
        addOutstandingRequestOnResync();
        return null;
    } catch (InvalidLabelResourceRequestException e) {
        // If Invalid label exception is received means the requested label
doesnt
        // have access so killing job in this case.
        String diagMsg = "Requested node-label-expression is invalid: "
            + StringUtils.stringifyException(e);
        LOG.info(diagMsg);
        JobId jobId = this.getJob().getID();
        eventHandler.handle(new JobDiagnosticsUpdateEvent(jobId, diagMsg));
        eventHandler.handle(new JobEvent(jobId, JobEventType.JOB_KILL));
        throw e;
    } catch (Exception e) {
        // This can happen when the connection to the RM has gone down. Keep
        // re-trying until the retryInterval has expired.
        if (System.currentTimeMillis() - retrystartTime >= retryInterval) {
            LOG.error("Could not contact RM after " + retryInterval + "
milliseconds.");
            eventHandler.handle(new JobEvent(this.getJob().getID(),
                JobEventType.JOB_AM_REBOOT));

```

```

        throw new RMContainerAllocationException("Could not contact RM after
" +
        retryInterval + " milliseconds.");
    }
    // Throw this up to the caller, which may decide to ignore it and
    // continue to attempt to contact the RM.
    throw e;
}
Resource newHeadRoom = getAvailableResources();
// 获取申请容器集合
List<Container> newContainers = response.getAllocatedContainers();
// Setting NMTokens
if (response.getNMTokens() != null) {
    for (NMToken nmToken : response.getNMTokens()) {
        NMTokenCache.setNMToken(nmToken.getNodeId().toString(),
            nmToken.getToken());
    }
}

// Setting AMRMToken
if (response.getAMRMToken() != null) {
    updateAMRMToken(response.getAMRMToken());
}

List<ContainerStatus> finishedContainers =
    response.getCompletedContainersStatuses();

// propagate preemption requests
final PreemptionMessage preemptReq = response.getPreemptionMessage();
if (preemptReq != null) {
    preemptionPolicy.preempt(
        new PreemptionContext(assignedRequests), preemptReq);
}

if (newContainers.size() + finishedContainers.size() > 0
    || !headRoom.equals(newHeadRoom)) {
    //something changed
    recalculateReduceSchedule = true;
    if (LOG.isDebugEnabled() && !headRoom.equals(newHeadRoom)) {
        LOG.debug("headroom=" + newHeadRoom);
    }
}

if (LOG.isDebugEnabled()) {
    for (Container cont : newContainers) {
        LOG.debug("Received new Container :" + cont);
    }
}

//Called on each allocation. Will know about newly blacklisted/added hosts.
computeIgnoreBlacklisting();

```

```

handleUpdatedNodes(response);
handleJobPriorityChange(response);
// Handle receiving the timeline collector address and token for this app.
MRAppMaster.RunningAppContext appContext =
    (MRAppMaster.RunningAppContext) this.getContext();
if (appContext.getTimelineV2Client() != null) {
    appContext.getTimelineV2Client().
        setTimelineCollectorInfo(response.getCollectorInfo());
}
for (ContainerStatus cont : finishedContainers) {
    processFinishedContainer(cont);
}
// 返回申请好的容器资源
return newContainers;
}

```

```

protected AllocateResponse makeRemoteRequest() throws YarnException,
    IOException {
    // 请求容器资源限制检查
    applyRequestLimits();
    // 申请容器的 NM 黑名单
    ResourceBlacklistRequest blacklistRequest =
        ResourceBlacklistRequest.newInstance(new ArrayList<String>
(blacklistAdditions),
            new ArrayList<String>(blacklistRemovals));
    // 创建申请容器请求 AllocateRequest
    AllocateRequest allocateRequest =
        AllocateRequest.newInstance(lastResponseID,
            super.getApplicationProgress(),
            // ask 就是申请目标容器的集合
            new ArrayList<ResourceRequest>(ask),
            new ArrayList<ContainerID>(release), blacklistRequest);
    // 发送 RPC 请求申请容器 (调用 ApplicationMasterService.allocate())
    // 返回申请容器结果
    AllocateResponse allocateResponse = scheduler.allocate(allocateRequest);
    lastResponseID = allocateResponse.getResponseID();
    availableResources = allocateResponse.getAvailableResources();
    lastClusterNmCount = clusterNmCount;
    clusterNmCount = allocateResponse.getNumClusterNodes();
    int numCompletedContainers =
        allocateResponse.getCompletedContainersStatuses().size();

    if (ask.size() > 0 || release.size() > 0) {
        /**
         * getResources() for application_1684656010852_0006:
         * ask=3 release= 0
         * newContainers=0 finishedContainers=0
         * resourceLimit=<memory:10752, vCores:23> knownNMs=3
         */
        LOG.info("getResources() for " + applicationId + ":" + " ask="
            + ask.size() + " release= " + release.size() + " newContainers="

```

```

        + allocateResponse.getAllocatedContainers().size()
        + " finishedContainers=" + numCompletedContainers
        + " resourceLimit=" + availableResources + " knownNMs="
        + clusterNmCount);
    }

    ask.clear();
    release.clear();

    if (numCompletedContainers > 0) {
        // re-send limited requests when a container completes to trigger asking
        // for more containers
        requestLimitsToUpdate.addAll(requestLimits.keySet());
    }

    if (blacklistAdditions.size() > 0 || blacklistRemovals.size() > 0) {
        LOG.info("Update the blacklist for " + applicationId +
            ": blacklistAdditions=" + blacklistAdditions.size() +
            " blacklistRemovals=" + blacklistRemovals.size());
    }
    blacklistAdditions.clear();
    blacklistRemovals.clear();
    return allocateResponse;
}

```

5.4.3.1.2 容器申请好之后开始做状态准备以及分配容器

```

// 分配启动容器 调用 ScheduledRequests.assign(allocatedContainers)
// this method will change the list of allocatedContainers.
private void assign(List<Container> allocatedContainers) {
    // allocatedContainers 是已经申请好的容器资源
    Iterator<Container> it = allocatedContainers.iterator();
    // Got allocated containers 1
    LOG.info("Got allocated containers " + allocatedContainers.size());
    containersAllocated += allocatedContainers.size();
    int reducePending = reduces.size();

    // 校验申请好的容器信息是否合法 (比如启动容器的 NM 信息是否在黑名单中)
    while (it.hasNext()) {
        Container allocated = it.next();
        if (LOG.isDebugEnabled()) {
            LOG.debug("Assigning container " + allocated.getId()
                + " with priority " + allocated.getPriority() + " to NM
                "
                + allocated.getNodeId());
        }

        // check if allocated container meets memory requirements
        // and whether we have any scheduled tasks that need
        // a container to be assigned
        boolean isAssignable = true;
        Priority priority = allocated.getPriority();
    }
}

```

```

// 获取申请好的容器资源
Resource allocatedResource = allocated.getResource();

if (PRIORITY_FAST_FAIL_MAP.equals(priority)
    || PRIORITY_MAP.equals(priority)
    || PRIORITY_OPPORTUNISTIC_MAP.equals(priority)) {
    if
(ResourceCalculatorUtils.computeAvailableContainers(allocatedResource,
    mapResourceRequest, getSchedularResourceTypes()) <= 0
    || maps.isEmpty()) {
        LOG.info("Cannot assign container " + allocated
            + " for a map as either "
            + " container memory less than required " +
mapResourceRequest
            + " or no pending map tasks - maps.isEmpty="
            + maps.isEmpty());
        isAssignable = false;
    }
} else if (PRIORITY_REDUCE.equals(priority)) {
    if
(ResourceCalculatorUtils.computeAvailableContainers(allocatedResource,
    reduceResourceRequest, getSchedularResourceTypes()) <= 0
    || (reducePending <= 0)) {
        LOG.info("Cannot assign container " + allocated
            + " for a reduce as either "
            + " container memory less than required " +
reduceResourceRequest
            + " or no pending reduce tasks.");
        isAssignable = false;
    } else {
        reducePending--;
    }
} else {
    LOG.warn("Container allocated at unwanted priority: " + priority
+
        ". Returning to RM...");
    isAssignable = false;
}

if (!isAssignable) {
    // release container if we could not assign it
    containerNotAssigned(allocated);
    it.remove();
    continue;
}

// do not assign if allocated container is on a
// blacklisted host
// 获取启动容器在 NM 的主机信息
String allocatedHost = allocated.getNodeId().getHost();
// 如果启动容器在 NM 的主机信息是在 NM 的黑名单 一般情况都不在黑名单
if (isNodeBlacklisted(allocatedHost)) {

```



```

        // we need to request for a new container
        // and release the current one
        LOG.info("Got allocated container on a blacklisted "
            + " host " + allocatedHost
            + ". Releasing container " + allocated);

        // find the request matching this allocated container
        // and replace it with a new one
        ContainerRequest toBeReplacedReq =
            getContainerReqToReplace(allocated);
        if (toBeReplacedReq != null) {
            LOG.info("Placing a new container request for task attempt "
                + toBeReplacedReq.attemptID);
            ContainerRequest newReq =
                getFilteredContainerRequest(toBeReplacedReq);
            decContainerReq(toBeReplacedReq);
            if (toBeReplacedReq.attemptID.getTaskId().getTaskType() ==
                TaskType.MAP) {
                maps.put(newReq.attemptID, newReq);
            } else {
                reduces.put(newReq.attemptID, newReq);
            }
            addContainerReq(newReq);
        } else {
            LOG.info("Could not map allocated container to a valid
request."
                + " Releasing allocated container " + allocated);
        }

        // release container if we could not assign it
        containerNotAssigned(allocated);
        it.remove();
        continue;
    }
}

// 到这里说明申请好的容器（也即启动容器在 NM）是合法可用
// 接下来就是根据容器的 NM 信息去连接 NM 启动对应的 Task
assignContainers(allocatedContainers);

// release container if we could not assign it
// 一般情况下 allocatedContainers 在 assignContainers() 已经使用移除啦
// 所以 it.hasNext() = false
it = allocatedContainers.iterator();
while (it.hasNext()) {
    Container allocated = it.next();
    LOG.info("Releasing unassigned container " + allocated);
    containerNotAssigned(allocated);
}
}

```

```

private void assignContainers(List<Container> allocatedContainers) {
    // 分配容器
    Iterator<Container> it = allocatedContainers.iterator();
    while (it.hasNext()) {
        // 分配好的容器 (向 RM 申请返回的)
        Container allocated = it.next();
        // 尝试分配容器 正常情况下是分配不成功的 (看里面的逻辑)
        ContainerRequest assigned = assignWithoutLocality(allocated);
        if (assigned != null) {
            // 分配容器 (正常情况下)
            containerAssigned(allocated, assigned);
            it.remove();
        }
    }

    // 正常情况下 到这里基本完成容器的分配 也即 allocatedContainers 为空
    // 这种情况是数据本地性 也即数据在当前节点 那么启动容器也在当前节点 避免数据的网络传输
    assignMapsWithLocality(allocatedContainers);
}

```

```

@SuppressWarnings("unchecked")
private void containerAssigned(Container allocated,
                               ContainerRequest assigned) {
    // Update resource requests
    // 更新容器资源
    decContainerReq(assigned);

    // send the container-assigned event to task attempt
    // 调用
    TaskAttemptEventDispatcher.handle(TaskAttemptEventType.TA_ASSIGNED)
    eventHandler.handle(new TaskAttemptContainerAssignedEvent(
        assigned.attemptID, allocated, applicationACLs));

    assignedRequests.add(allocated, assigned.attemptID);

    if (LOG.isDebugEnabled()) {
        LOG.debug("Assigned container (" + allocated + ") "
            + " to task " + assigned.attemptID + " on node "
            + allocated.getNodeId().toString());
    }
}

```

5.4.3.1.3 分配容器的状态转换之后最终调用

ContainerAssignedTransition.doTransition(TaskAttemptEventType.TA_ASSIGNED) 构建启动容器的进程(YarnChild)上下文环境

```

private static class ContainerAssignedTransition implements
    SingleArcTransition<TaskAttemptImpl, TaskAttemptEvent> {
    @SuppressWarnings({"unchecked"})
    @Override
    public void transition(final TaskAttemptImpl taskAttempt,

```

```

        TaskAttemptEvent event) {
    // eventType = TaskAttemptEventType.TA_ASSIGNED
    final TaskAttemptContainerAssignedEvent cEvent =
        (TaskAttemptContainerAssignedEvent) event;
    // 获取申请好的容器启动 Task
    Container container = cEvent.getContainer();
    taskAttempt.container = container;
    // this is a _real_ Task (classic Hadoop mapred flavor):
    taskAttempt.remoteTask = taskAttempt.createRemoteTask();
    taskAttempt.jvmID =
        new WrappedJvmID(taskAttempt.remoteTask.getTaskID().getJobID(),
            taskAttempt.remoteTask.isMapTask(),
            taskAttempt.container.getId().getContainerId());
    taskAttempt.taskAttemptListener.registerPendingTask(
        taskAttempt.remoteTask, taskAttempt.jvmID);

    taskAttempt.computeRackAndLocality();

    //launch the container
    //create the container object to be launched for a given Task attempt
    // 创建启动容器上下文的 ContainerLaunchContext
    // 也即容器启动 YarnChild 进程相关命令以及环境变量
    ContainerLaunchContext launchContext = createContainerLaunchContext(
        cEvent.getApplicationACLs(), taskAttempt.conf,
taskAttempt.jobToken,
        taskAttempt.remoteTask, taskAttempt.oldJobId, taskAttempt.jvmID,
        taskAttempt.taskAttemptListener, taskAttempt.credentials);
    // 调用
    ContainerLauncherRouter.handle(ContainerLauncher.EventType.CONTAINER_REMOTE_LAUNCH)
        taskAttempt.eventHandler
            .handle(new ContainerRemoteLaunchEvent(taskAttempt.attemptId,
                launchContext, container, taskAttempt.remoteTask));

    // send event to speculator that our container needs are satisfied
    // 发送事件给推测执行器 让其分析启动容器
    // 调用
    SpeculatorEventDispatcher.handle(Speculator.EventType.TASK_CONTAINER_NEED_UPDATE)
        taskAttempt.eventHandler.handle
            (new SpeculatorEvent(taskAttempt.getID().getTaskId(), -1));
    }
}

```

```

static ContainerLaunchContext createContainerLaunchContext(
    Map<ApplicationAccessType, String> applicationACLs,
    Configuration conf, Token<JobTokenIdentifier> jobToken, Task remoteTask,
    final org.apache.hadoop.mapred.JobID oldJobId,
    WrappedJvmID jvmID,
    TaskAttemptListener taskAttemptListener,
    Credentials credentials) {

    synchronized (commonContainersSpecLock) {

```

```

        if (commonContainersSpec == null) {
            // 构建通用的容器启动进程模板 此时还没有对应启动进程命令
            commonContainersSpec = createCommonContainerLaunchContext(
                applicationACLs, conf, jobToken, oldJobId, credentials);
        }
    }

    // Fill in the fields needed per-container that are missing in the common
    // spec.

    boolean userClassesTakesPrecedence =
        conf.getBoolean(MRJobConfig.MAPREDUCE_JOB_USER_CLASSPATH_FIRST,
false);

    // Setup environment by cloning from common env.
    Map<String, String> env = commonContainersSpec.getEnvironment();
    Map<String, String> myEnv = new HashMap<String, String>(env.size());
    myEnv.putAll(env);
    if (userClassesTakesPrecedence) {
        myEnv.put(Environment.CLASSPATH_PREPEND_DISTCACHE.name(), "true");
    }
    MapReduceChildJVM.setVMEnv(myEnv, remoteTask);

    // Set up the launch command
    // 构建启动容器进行命令 (也即启动容器啥进程 也即 YarnChild 进程)
    List<String> commands = MapReduceChildJVM.getVMCommand(
        taskAttemptListener.getAddress(), remoteTask, jvmID);

    // Duplicate the ByteBuffers for access by multiple containers.
    Map<String, ByteBuffer> myServiceData = new HashMap<String, ByteBuffer>();
    for (Entry<String, ByteBuffer> entry : commonContainersSpec
        .getServiceData().entrySet()) {
        myServiceData.put(entry.getKey(), entry.getValue().duplicate());
    }

    // Construct the actual Container
    // 构建容器上下文对象 也即准备好了容器启动 YarnChild 进程的所有准备工作
    ContainerLaunchContext container = ContainerLaunchContext.newInstance(
        commonContainersSpec.getLocalResources(), myEnv, commands,
        myServiceData, commonContainersSpec.getTokens().duplicate(),
        applicationACLs);

    return container;
}

```

```

public class MapReduceChildJVM {
    public static List<String> getVMCommand(
        InetAddress taskAttemptListenerAddr, Task task,
        JVMId jvmID) {

        TaskAttemptID attemptID = task.getTaskID();
    }
}

```

```

JobConf conf = task.conf;

Vector<String> vargs = new Vector<String>(8);

vargs.add(MRApps.crossPlatformifyMREnv(task.conf, Environment.JAVA_HOME)
    + "/bin/java");

// Add child (task) java-vm options.
//
// The following symbols if present in mapred.{map|reduce}.child.java.opts
// value are replaced:
// + @taskid@ is interpolated with value of TaskID.
// Other occurrences of @ will not be altered.
//
// Example with multiple arguments and substitutions, showing
// jvm GC logging, and start of a passwordless JVM JMX agent so can
// connect with jconsole and the likes to watch child memory, threads
// and get thread dumps.
//
// <property>
//   <name>mapred.map.child.java.opts</name>
//   <value>-Xmx 512M -verbose:gc -Xloggc:/tmp/@taskid@.gc \
//           -Dcom.sun.management.jmxremote.authenticate=false \
//           -Dcom.sun.management.jmxremote.ssl=false \
//   </value>
// </property>
//
// <property>
//   <name>mapred.reduce.child.java.opts</name>
//   <value>-Xmx 1024M -verbose:gc -Xloggc:/tmp/@taskid@.gc \
//           -Dcom.sun.management.jmxremote.authenticate=false \
//           -Dcom.sun.management.jmxremote.ssl=false \
//   </value>
// </property>
//
String javaOpts = getChildJavaOpts(conf, task.isMapTask());
javaOpts = javaOpts.replace("@taskid@", attemptID.toString());
String [] javaOptsSplit = javaOpts.split(" ");
for (int i = 0; i < javaOptsSplit.length; i++) {
    vargs.add(javaOptsSplit[i]);
}

Path childTmpDir = new Path(MRApps.crossPlatformifyMREnv(conf, Environment.PWD),
    YarnConfiguration.DEFAULT_CONTAINER_TEMP_DIR);
vargs.add("-Djava.io.tmpdir=" + childTmpDir);
MRApps.addLog4jSystemProperties(task, vargs, conf);

if (conf.getProfileEnabled()) {
    if (conf.getProfileTaskRange(task.isMapTask()
        ).isIncluded(task.getPartition())) {
        final String profileParams = conf.get(task.isMapTask()
            ? MRJobConfig.TASK_MAP_PROFILE_PARAMS

```

```

        : MRJobConfig.TASK_REDUCE_PROFILE_PARAMS, conf.getProfileParams());
    vargs.add(String.format(profileParams,
        getTaskLogFile(TaskLog.LogName.PROFILE)));
    }
}

// Add main class and its arguments
vargs.add(YarnChild.class.getName()); // main of Child
// pass TaskAttemptListener's address
vargs.add(taskAttemptListenerAddr.getAddress().getHostAddress());
vargs.add(Integer.toString(taskAttemptListenerAddr.getPort()));
vargs.add(attemptID.toString()); // pass task identifier

// Finally add the jvmID
vargs.add(String.valueOf(jvmID.getId()));
vargs.add("1>" + getTaskLogFile(TaskLog.LogName.STDOUT));
vargs.add("2>" + getTaskLogFile(TaskLog.LogName.STDERR));

// Final command
StringBuilder mergedCommand = new StringBuilder();
for (CharSequence str : vargs) {
    mergedCommand.append(str).append(" ");
}
Vector<String> vargsFinal = new Vector<String>(1);
vargsFinal.add(mergedCommand.toString());
return vargsFinal;
}
}

```

5.4.3.1.4 启动容器进程(YarnChild)相关环境准备好之后连接 NM 的 ContainerManagerImpl 服务 发送 RPC 请求启动容器 YarnChild 进程(调用 ContainerManagerImpl.startContainers() 类似启动 MRAppMaster 进程)

```

// 最终调用 ContainerLauncherImpl.EventProcessor.run()
/**
 * Setup and start the container on remote nodemanager.
 */
class EventProcessor implements Runnable {
    private ContainerLauncherEvent event;

    EventProcessor(ContainerLauncherEvent event) {
        // eventType = ContainerLauncher.EventType.CONTAINER_REMOTE_LAUNCH
        this.event = event;
    }

    @Override
    public void run() {
        // Processing the event
        // EventType: CONTAINER_REMOTE_LAUNCH
        // for container container_1684656010852_0006_01_000002
        // taskAttempt attempt_1684656010852_0006_m_000000_0
        LOG.info("Processing the event " + event.toString());
    }
}

```

```

// Load ContainerManager tokens before creating a connection.
// TODO: Do it only once per NodeManager.
ContainerId containerID = event.getContainerID();

Container c = getContainer(event);
switch (event.getType()) {

    case CONTAINER_REMOTE_LAUNCH:
        ContainerRemoteLaunchEvent launchEvent
            = (ContainerRemoteLaunchEvent) event;
        // 部署容器
        c.launch(launchEvent);
        break;

    case CONTAINER_REMOTE_CLEANUP:
        c.kill(event.getDumpContainerThreads());
        break;

    case CONTAINER_COMPLETED:
        c.done();
        break;

}
removeContainerIfDone(containerID);
}
}

```

```

// 调用 ContainerLauncherImpl.Container.launch() 发送 RPC 请求部署容器
@SuppressWarnings("unchecked")
public synchronized void launch(ContainerRemoteLaunchEvent event) {
    // Launching attempt_1684656010852_0006_m_000000_0
    LOG.info("Launching " + taskAttemptID);
    if (this.state == ContainerState.KILLED_BEFORE_LAUNCH) {
        state = ContainerState.DONE;
        sendContainerLaunchFailedMsg(taskAttemptID,
            "Container was killed before it was launched");
        return;
    }

    ContainerManagementProtocolProxyData proxy = null;
    try {
        // 获取 NM 的 RPC ContainerManagerImpl 服务客户端代理
        // 通讯协议接口为 ContainerManagementProtocol
        proxy = getCMPProxy(containerMgrAddress, containerID);

        // Construct the actual Container
        // 获取部署容器上下文对象 ContainerLaunchContext
        ContainerLaunchContext containerLaunchContext =
            event.getContainerLaunchContext();
    }
}

```

```

        // Now launch the actual container
        // 创建启动容器 RPC 请求体
        StartContainerRequest startRequest =
            StartContainerRequest.newInstance(containerLaunchContext,
                event.getContainerToken());
        List<StartContainerRequest> list = new
ArrayList<StartContainerRequest>();
        list.add(startRequest);
        StartContainersRequest requestList =
StartContainersRequest.newInstance(list);

        // 发送 RPC 请求启动容器
        // 调用 ContainerManagerImpl.startContainers()
        StartContainersResponse response =
            proxy.getContainerManagementProtocol()
                .startContainers(requestList);

        if (response.getFailedRequests() != null
            && response.getFailedRequests().containsKey(containerID)) {
            throw
response.getFailedRequests().get(containerID).deserialize();
        }
        ByteBuffer portInfo =
            response.getAllServicesMetaData().get(
                ShuffleHandler.MAPREDUCE_SHUFFLE_SERVICEID);
        int port = -1;
        if (portInfo != null) {
            port = ShuffleHandler.deserializeMetaData(portInfo);
        }
        LOG.info("Shuffle port returned by ContainerManager for "
            + taskAttemptID + " : " + port);

        if (port < 0) {
            this.state = ContainerState.FAILED;
            throw new IllegalStateException("Invalid shuffle port number "
                + port + " returned for " + taskAttemptID);
        }

        // after launching, send launched event to task attempt to move
        // it from ASSIGNED to RUNNING state
        context.getEventHandler().handle(
            new TaskAttemptContainerLaunchedEvent(taskAttemptID, port));
        this.state = ContainerState.RUNNING;
    } catch (Throwable t) {
        String message = "Container launch failed for " + containerID + " :
"
            + StringUtils.stringifyException(t);
        this.state = ContainerState.FAILED;
        sendContainerLaunchFailedMsg(taskAttemptID, message);
    } finally {
        if (proxy != null) {
            cmProxy.maybeCloseProxy(proxy);
        }
    }
}

```



```

    }
}
}

```

5.4.3.2 处理 ReduceTask 容器请求(类似启动 MapTask)

六 启动 MR 真正处理任务 YarnChild 进程

```

/**
 * The main() for MapReduce task processes.
 */
class YarnChild {
    public static void main(String[] args) throws Throwable {
        Thread.setDefaultUncaughtExceptionHandler(new
YarnUncaughtExceptionHandler());
        LOG.debug("Child starting");

        // 获取任务配置文件并解析 也即 job.xml
        final JobConf job = new JobConf(MRJobConfig.JOB_CONF_FILE);
        // Initing with our JobConf allows us to avoid loading confs twice
        // 初始化配置避免加载两次
        Limits.init(job);

        UserGroupInformation.setConfiguration(job);
        // MAPREDUCE-6565: need to set configuration for SecurityUtil.
        SecurityUtil.setConfiguration(job);

        /**
         * exec /bin/bash -c
         * "$JAVA_HOME/bin/java
         * -Djava.net.preferIPv4Stack=true
         * -Dhadoop.metrics.log.level=WARN
         * -Xmx820m
         * -Djava.io.tmpdir=$PWD/tmp
         * -Dlog4j.configuration=container-log4j.properties
         * -Dyarn.app.container.log.dir=/opt/app/hadoop-
3.1.3/logs/userlogs/application_1684656010852_0007/container_1684656010852_0007_01_0
00002
         * -Dyarn.app.container.log.filesize=0
         * -Dhadoop.root.logger=INFO,CLA
         * -Dhadoop.root.logfile=syslog
         * org.apache.hadoop.mapred.YarnChild
         * 192.168.6.102
         * 39956
         * attempt_1684656010852_0007_m_000000_0
         * 2
         * 1>/opt/app/hadoop-
3.1.3/logs/userlogs/application_1684656010852_0007/container_1684656010852_0007_01_0
00002/stdout

```

```

    * 2>/opt/app/hadoop-
3.1.3/logs/userlogs/application_1684656010852_0007/container_1684656010852_0007_01_0
00002/stderr
    */
    // AM 地址
    String host = args[0];
    int port = Integer.parseInt(args[1]);

    final InetSocketAddress address =
        NetUtils.createSocketAddrForHost(host, port);
    final TaskAttemptID firstTaskid = TaskAttemptID.forName(args[2]);
    Long jvmIdLong = Long.parseLong(args[3]);
    JVMId jvmId = new JVMId(firstTaskid.getJobID(),
        firstTaskid.getTaskType() == TaskType.MAP, jvmIdLong);

    CallerContext.setCurrent(
        new CallerContext.Builder("mr_" + firstTaskid.toString()).build());

    // initialize metrics
    DefaultMetricsSystem.initialize(
        StringUtils.camelize(firstTaskid.getTaskType().name()) + "Task");

    // Security framework already loaded the tokens into current ugi
    Credentials credentials =
        UserGroupInformation.getCurrentUser().getCredentials();

    // Executing with tokens:
    // [Kind: mapreduce.job, Service: job_1684656010852_0007,
    // Ident:
    (org.apache.hadoop.mapreduce.security.token.JobTokenIdentifier@22555ebf)]
    LOG.info("Executing with tokens: {}", credentials.getAllTokens());

    // Create TaskUmbilicalProtocol as actual task owner.
    UserGroupInformation taskOwner =

    UserGroupInformation.createRemoteUser(firstTaskid.getJobID().toString());
    Token<JobTokenIdentifier> jt = TokenCache.getJobToken(credentials);
    SecurityUtil.setTokenService(jt, address);
    taskOwner.addToken(jt);
    // 获取 AM(MRAppMaster) 的 Task 监听服务 TaskAttemptListenerImpl
    // 该 TaskAttemptListenerImpl 内部启动了一个 RPC Server 通讯协议为
    TaskUmbilicalProtocol
    // 这个获取 AM 的 RPC 服务 TaskAttemptListenerImpl 的 RPC 代理对象 (也即客户端)
    final TaskUmbilicalProtocol umbilical =
        taskOwner.doAs(new PrivilegedExceptionAction<TaskUmbilicalProtocol>
() {
    @Override
    public TaskUmbilicalProtocol run() throws Exception {
        return (TaskUmbilicalProtocol)
RPC.getProxy(TaskUmbilicalProtocol.class,
                TaskUmbilicalProtocol.versionID, address, job);
    }
}

```

```

    });

    // report non-pid to application master
   JvmContext context = new JvmContext(jvmId, "-1000");
    LOG.debug("PID: " + System.getenv().get("JVM_PID"));
    Task task = null;
    UserGroupInformation childUGI = null;
    ScheduledExecutorService logSyncer = null;

    try {
        int idleLoopCount = 0;
        JvmTask myTask = null;
        // poll for new task
        for (int idle = 0; null == myTask; ++idle) {
            long sleepTimeMillisecs = Math.min(idle * 500, 1500);
            // Sleeping for 0ms before retrying again. Got null now.
            LOG.info("sleeping for " + sleepTimeMillisecs
                    + "ms before retrying again. Got null now.");
            MILLISECONDS.sleep(sleepTimeMillisecs);
            // 发送 RPC 请求给 AM 的 TaskAttemptListenerImpl RPC 服务获取当前
YarnChild 进程执行任务
            myTask = umbilical.getTask(context);
        }
        if (myTask.shouldDie()) {
            return;
        }

        // 获取执行任务 可能是 MapTask 也可能是 ReduceTask
        task = myTask.getTask();
        YarnChild.taskid = task.getTaskID();

        // Create the job-conf and set credentials
        // 配置执行 Task 的一些配置信息以及证书相关的
        configureTask(job, task, credentials, jt);

        // log the system properties
        /**
         * /*****
         * [system properties]
         * os.name: Linux
         * os.version: 3.10.0-1062.el7.x86_64
         * java.home: /opt/app/jdk1.8.0_212/jre
         * java.runtime.version: 1.8.0_212-b10
         * java.vendor: Oracle Corporation
         * java.version: 1.8.0_212
         * java.vm.name: Java HotSpot(TM) 64-Bit Server VM

```

```
* java.class.path: /opt/app/hadoop-3.1.3/data/nm-local-  
dir/usercache/tanbs/appcache/application_1684656010852_0007/container_1684656010852_  
0007_01_000002:/opt/app/hadoop-3.1.3/etc/hadoop:/opt/app/hadoop-  
3.1.3/share/hadoop/common/hadoop-common-3.1.3-tests.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/hadoop-common-3.1.3.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/hadoop-kms-3.1.3.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/hadoop-nfs-3.1.3.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/jetty-io-9.3.24.v20180605.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/accessors-smart-1.2.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/jcip-annotations-1.0-1.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/animal-sniffer-annotations-1.17.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/netty-3.10.5.Final.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/asm-5.0.4.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/jersey-core-1.19.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/audience-annotations-0.5.0.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/nimbus-jose-jwt-4.41.1.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/avro-1.7.7.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/jetty-security-9.3.24.v20180605.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/checker-qual-2.5.2.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/jetty-server-9.3.24.v20180605.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/commons-beanutils-1.9.3.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/paranamer-2.3.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/commons-cli-1.2.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/jsr311-api-1.1.1.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/commons-codec-1.11.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/jersey-json-1.19.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/commons-collections-3.2.2.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/jul-to-slf4j-1.7.25.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/commons-compress-1.18.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/jersey-servlet-1.19.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/commons-configuration2-2.1.1.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/protobuf-java-2.5.0.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/commons-io-2.5.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/re2j-1.1.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/commons-lang-2.6.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/kerb-admin-1.0.1.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/commons-lang3-3.4.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/kerb-client-1.0.1.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/commons-logging-1.1.3.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/kerb-common-1.0.1.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/commons-math3-3.1.1.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/slf4j-api-1.7.25.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/commons-net-3.6.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/kerb-core-1.0.1.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/curator-client-2.13.0.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/kerb-crypto-1.0.1.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/curator-framework-2.13.0.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/kerb-identity-1.0.1.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/curator-recipes-2.13.0.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/jersey-server-1.19.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/error_prone_annotations-2.2.0.jar:/opt/app/hadoop-  
3.1.3/share/hadoop/common/lib/kerb-server-1.0.1.jar:/opt/app/hadoop-
```

3.1.3/share/hadoop/common/lib/failureaccess-1.0.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/gson-2.2.4.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/snappy-java-1.0.5.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/guava-27.0-jre.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/kerb-simplekdc-1.0.1.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/hadoop-annotations-3.1.3.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/kerb-util-1.0.1.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/hadoop-auth-3.1.3.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jsch-0.1.54.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/htrace-core4-4.1.0-incubating.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/stax2-api-3.1.4.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/httpclient-4.5.2.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/token-provider-1.0.1.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/httpcore-4.4.4.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/kerby-asn1-1.0.1.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/j2objc-annotations-1.1.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jettison-1.1.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jackson-annotations-2.7.8.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/kerby-config-1.0.1.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jackson-core-2.7.8.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/kerby-pkix-1.0.1.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jackson-core-asl-1.9.13.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/kerby-util-1.0.1.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jackson-databind-2.7.8.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/kerby-xdr-1.0.1.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jackson-jaxrs-1.9.13.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jetty-http-9.3.24.v20180605.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jackson-mapper-asl-1.9.13.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jackson-xc-1.9.13.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/log4j-1.2.17.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/javax.servlet-api-3.1.0.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/woodstox-core-5.0.3.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jaxb-api-2.2.11.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/metrics-core-3.2.4.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jaxb-impl-2.2.3-1.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jetty-servlet-9.3.24.v20180605.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/json-smart-2.3.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jetty-util-9.3.24.v20180605.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jsp-api-2.1.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jetty-webapp-9.3.24.v20180605.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jsr305-3.0.0.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/jetty-xml-9.3.24.v20180605.jar:/opt/app/hadoop-
3.1.3/share/hadoop/common/lib/listenablefuture-9999.0-empty-to-avoid-conflict-with-
guava.jar:/opt/app/hadoop-3.1.3/share/hadoop/common/lib/zookeeper-
3.4.13.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/hadoop-hdfs-3.1.3-
tests.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/hadoop-hdfs-
3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/hadoop-hdfs-client-3.1.3-
tests.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/hadoop-hdfs-client-
3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/hadoop-hdfs-httpfs-
3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/hadoop-hdfs-native-client-3.1.3-
tests.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/hadoop-hdfs-native-client-

3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/hadoop-hdfs-nfs-
3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/hadoop-hdfs-rbf-3.1.3-
tests.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/hadoop-hdfs-rbf-
3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jetty-http-
9.3.24.v20180605.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/accessors-smart-
1.2.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jaxb-impl-2.2.3-
1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/animal-sniffer-annotations-
1.17.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/log4j-
1.2.17.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/asm-
5.0.4.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jcip-annotations-1.0-
1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/audience-annotations-
0.5.0.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/netty-
3.10.5.Final.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/avro-
1.7.7.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jetty-io-
9.3.24.v20180605.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/checker-qual-
2.5.2.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jetty-security-
9.3.24.v20180605.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/commons-beanutils-
1.9.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/netty-all-
4.0.52.Final.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/commons-cli-
1.2.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/json-smart-
2.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/commons-codec-
1.11.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jersey-core-
1.19.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/commons-collections-
3.2.2.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jsr305-
3.0.0.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/commons-compress-
1.18.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jersey-server-
1.19.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/commons-configuration2-
2.1.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jsr311-api-
1.1.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/commons-daemon-
1.0.13.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/nimbus-jose-jwt-
4.41.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/commons-io-
2.5.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/okhttp-
2.7.5.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/commons-lang-
2.6.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerb-admin-
1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/commons-lang3-
3.4.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerb-client-
1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/commons-logging-
1.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerb-common-
1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/commons-math3-
3.1.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/okio-
1.6.0.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/commons-net-
3.6.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerb-core-
1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/curator-client-
2.13.0.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerb-crypto-
1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/curator-framework-
2.13.0.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerb-identity-
1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/curator-recipes-
2.13.0.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jersey-json-
1.19.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/error_prone_annotations-
2.2.0.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerb-server-
1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/failureaccess-
1.0.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/paranamer-

2.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/gson-2.2.4.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/protobuf-java-2.5.0.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/guava-27.0-jre.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerb-simplekdc-1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/hadoop-annotations-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerb-util-1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/hadoop-auth-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jetty-webapp-9.3.24.v20180605.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/htrace-core4-4.1.0-incubating.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/re2j-1.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/httpclient-4.5.2.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/snappy-java-1.0.5.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/httpcore-4.4.4.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerby-asn1-1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/j2objc-annotations-1.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jersey-servlet-1.19.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jackson-annotations-2.7.8.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerby-config-1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jackson-core-2.7.8.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerby-pkix-1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jackson-core-asl-1.9.13.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerby-util-1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jackson-databind-2.7.8.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/kerby-xdr-1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jackson-jaxrs-1.9.13.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jettison-1.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jackson-mapper-asl-1.9.13.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/leveldbjni-all-1.8.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jackson-xc-1.9.13.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/javax.servlet-api-3.1.0.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/stax2-api-3.1.4.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jaxb-api-2.2.11.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jetty-server-9.3.24.v20180605.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jetty-xml-9.3.24.v20180605.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jetty-servlet-9.3.24.v20180605.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jsch-0.1.54.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jetty-util-9.3.24.v20180605.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/json-simple-1.1.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/jetty-util-ajax-9.3.24.v20180605.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/listenablefuture-9999.0-empty-to-avoid-conflict-with-guava.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/token-provider-1.0.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/woodstox-core-5.0.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/hdfs/lib/zookeeper-3.4.13.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-api-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-applications-distributedshell-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-applications-unmanaged-am-launcher-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-client-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-common-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-registry-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-server-applicationhistoryservice-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-server-common-


```
3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-server-nodemanager-
3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-server-
resourcemanager-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-
server-router-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-server-
sharedcachemanager-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-
server-tests-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-server-
timeline-pluginstorage-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-
yarn-server-web-proxy-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-
services-api-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/hadoop-yarn-services-
core-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/HikariCP-java7-
2.4.12.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/aopalliance-
1.0.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/dnsjava-
2.1.7.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/ehcache-
3.3.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/fst-2.50.jar:/opt/app/hadoop-
3.1.3/share/hadoop/yarn/lib/geronimo-jcache_1.0_spec-1.0-alpha-
1.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/guice-4.0.jar:/opt/app/hadoop-
3.1.3/share/hadoop/yarn/lib/guice-servlet-4.0.jar:/opt/app/hadoop-
3.1.3/share/hadoop/yarn/lib/jackson-jaxrs-base-2.7.8.jar:/opt/app/hadoop-
3.1.3/share/hadoop/yarn/lib/jackson-jaxrs-json-provider-2.7.8.jar:/opt/app/hadoop-
3.1.3/share/hadoop/yarn/lib/jackson-module-jaxb-annotations-
2.7.8.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/java-util-
1.9.0.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/javax.inject-
1.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/jersey-client-
1.19.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/jersey-guice-
1.19.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/json-io-
2.5.1.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/metrics-core-
3.2.4.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/mssql-jdbc-
6.2.1.jre7.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/objenesis-
1.0.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/snakeyaml-
1.16.jar:/opt/app/hadoop-3.1.3/share/hadoop/yarn/lib/swagger-annotations-
1.5.4.jar:/opt/app/hadoop-3.1.3/share/hadoop/mapreduce/hadoop-mapreduce-client-app-
3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/mapreduce/hadoop-mapreduce-client-
common-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/mapreduce/hadoop-mapreduce-
client-core-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/mapreduce/hadoop-mapreduce-
client-hs-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/mapreduce/hadoop-mapreduce-
client-hs-plugins-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/mapreduce/hadoop-
mapreduce-client-jobclient-3.1.3-tests.jar:/opt/app/hadoop-
3.1.3/share/hadoop/mapreduce/hadoop-mapreduce-client-jobclient-
3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/mapreduce/hadoop-mapreduce-client-
nativetask-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/mapreduce/hadoop-mapreduce-
client-shuffle-3.1.3.jar:/opt/app/hadoop-3.1.3/share/hadoop/mapreduce/hadoop-
mapreduce-client-uploader-3.1.3.jar:/opt/app/hadoop-
3.1.3/share/hadoop/mapreduce/hadoop-mapreduce-examples-3.1.3.jar:/opt/app/hadoop-
3.1.3/share/hadoop/mapreduce/lib/hamcrest-core-1.3.jar:/opt/app/hadoop-
3.1.3/share/hadoop/mapreduce/lib/junit-
4.11.jar:/opt/app/hadoop-3.1.3/share/hadoop/mapreduce/lib/job.jar:/opt/app/hadoop-
3.1.3/data/nm-local-
dir/usercache/tanbs/appcache/application_1684656010852_0007/container_1684656010852_
0007_01_000002/job.jar
* java.io.tmpdir: /opt/app/hadoop-3.1.3/data/nm-local-
dir/usercache/tanbs/appcache/application_1684656010852_0007/container_1684656010852_
0007_01_000002/tmp
```



```

        * user.dir: /opt/app/hadoop-3.1.3/data/nm-local-
dir/usercache/tanbs/appcache/application_1684656010852_0007/container_1684656010852_
0007_01_000002
        * user.name: tanbs
        * *****
    */
    String systemPropsToLog = MRApps.getSystemPropertiesToLog(job);
    if (systemPropsToLog != null) {
        LOG.info(systemPropsToLog);
    }

    // Initiate Java VM metrics
    JvmMetrics.initSingleton(jvmId.toString(), job.getSessionId());
    childDUGI = UserGroupInformation.createRemoteUser(System
        .getenv(ApplicationConstants.Environment.USER.toString()));
    // Add tokens to new user so that it may execute its task correctly.
    childDUGI.addCredentials(credentials);

    // set job classloader if configured before invoking the task
    // 默认没有配置类加载器
    MRApps.setJobClassLoader(job);

    logSyncer = TaskLog.createLogSyncer();

    // Create a final reference to the task for the doAs block
    // 这里就是执行 MR 任务的最终 MapTask 或者 ReduceTask
    final Task taskFinal = task;
    childDUGI.doAs(new PrivilegedExceptionAction<Object>() {
        @Override
        public Object run() throws Exception {
            // use job-specified working directory
            setEncryptedSpillKeyIfRequired(taskFinal);

            FileSystem.get(job).setWorkingDirectory(job.getWorkingDirectory());
            // 调用 MapTask.run() 或者 ReduceTask.run()
            taskFinal.run(job, umbilical); // run the task
            return null;
        }
    });
} catch (FSError e) {
    LOG.error("FSError from child", e);
    if (!ShutdownHookManager.get().isShutdownInProgress()) {
        umbilical.fsError(taskid, e.getMessage());
    }
} catch (Exception exception) {
    LOG.warn("Exception running child : "
        + StringUtils.stringifyException(exception));
    try {
        if (task != null) {
            // do cleanup for the task
            if (childDUGI == null) { // no need to job into doAs block
                task.taskCleanup(umbilical);
            }
        }
    }
}

```

```

        } else {
            final Task taskFinal = task;
            childUGI.doAs(new PrivilegedExceptionAction<Object>() {
                @Override
                public Object run() throws Exception {
                    taskFinal.taskCleanup(umbilical);
                    return null;
                }
            });
        }
    }
} catch (Exception e) {
    LOG.info("Exception cleaning up: " +
StringUtils.stringifyException(e));
}
// Report back any failures, for diagnostic purposes
if (taskid != null) {
    if (!ShutdownHookManager.get().isShutdownInProgress()) {
        umbilical.fatalError(taskid,
            StringUtils.stringifyException(exception), false);
    }
}
} catch (Throwable throwable) {
    LOG.error("Error running child : "
        + StringUtils.stringifyException(throwable));
    if (taskid != null) {
        if (!ShutdownHookManager.get().isShutdownInProgress()) {
            Throwable tCause = throwable.getCause();
            String cause =
                tCause == null ? throwable.getMessage() : StringUtils
                    .stringifyException(tCause);
            umbilical.fatalError(taskid, cause, false);
        }
    }
} finally {
    RPC.stopProxy(umbilical);
    DefaultMetricsSystem.shutdown();
    TaskLog.syncLogsShutdown(logSyncer);
}
}
}

```

6.1 获取 AM(MRAppMaster) 的 Task 监听服务 TaskAttemptListenerImpl RPC 服务客户端代理

```
// 获取 AM(MRAppMaster) 的 Task 监听服务 TaskAttemptListenerImpl
// 该 TaskAttemptListenerImpl 内部启动了一个 RPC Server 通讯协议为
TaskUmbilicalProtocol
// 这个获取 AM 的 RPC 服务 TaskAttemptListenerImpl 的 RPC 代理对象 (也即客户端)
final TaskUmbilicalProtocol umbilical =
    taskOwner.doAs(new PrivilegedExceptionAction<TaskUmbilicalProtocol>
() {
    @Override
    public TaskUmbilicalProtocol run() throws Exception {
        return (TaskUmbilicalProtocol)
RPC.getProxy(TaskUmbilicalProtocol.class,
                TaskUmbilicalProtocol.versionID, address, job);
    }
});
```

6.2 发送 RPC 请求 TaskAttemptListenerImpl RPC 服务获取执行 Task

```
try {
    int idleLoopCount = 0;
    JvmTask myTask = null;
    // poll for new task
    for (int idle = 0; null == myTask; ++idle) {
        long sleepTimeMillisecs = Math.min(idle * 500, 1500);
        // sleeping for 0ms before retrying again. Got null now.
        LOG.info("sleeping for " + sleepTimeMillisecs
            + "ms before retrying again. Got null now.");
        MILLISECONDS.sleep(sleepTimeMillisecs);
        // 发送 RPC 请求给 AM 的 TaskAttemptListenerImpl RPC 服务获取当前
YarnChild 进程执行任务
        myTask = umbilical.getTask(context);
    }
}
```

6.3 执行 MapTask或者ReduceTask 业务逻辑

```
// Create a final reference to the task for the doAs block
// 这里就是执行 MR 任务的最终 MapTask 或者 ReduceTask
final Task taskFinal = task;
childUGI.doAs(new PrivilegedExceptionAction<Object>() {
    @Override
    public Object run() throws Exception {
        // use job-specified working directory
        setEncryptedSpillKeyIfRequired(taskFinal);

        FileSystem.get(job).setWorkingDirectory(job.getWorkingDirectory());
        // 调用 MapTask.run() 或者 ReduceTask.run()
        taskFinal.run(job, umbilical); // run the task
        return null;
    }
}
```

```
});
```

6.3.1 执行 Mapper 调用 MapTask.run()

```
/**
 * A Map task.
 */
@InterfaceAudience.LimitedPrivate({"MapReduce"})
@InterfaceStability.Unstable
public class MapTask extends Task {
    @Override
    public void run(final JobConf job, final TaskUmbilicalProtocol umbilical)
        throws IOException, ClassNotFoundException, InterruptedException {
        this.umbilical = umbilical;

        // MapTask
        if (isMapTask()) {
            // If there are no reducers then there won't be any sort. Hence the map
            // phase will govern the entire attempt's progress.
            if (conf.getNumReduceTasks() == 0) {
                mapPhase = getProgress().addPhase("map", 1.0f);
            } else {
                // If there are reducers then the entire attempt's progress will be
                // split between the map phase (67%) and the sort phase (33%).
                mapPhase = getProgress().addPhase("map", 0.667f);
                sortPhase = getProgress().addPhase("sort", 0.333f);
            }
        }

        TaskReporter reporter = startReporter(umbilical);

        // 默认使用 NEW API
        boolean useNewApi = job.getUseNewMapper();
        // 初始化 MapTask
        initialize(job, getJobID(), reporter, useNewApi);

        // check if it is a cleanupJobTask
        if (jobCleanup) {
            runJobCleanupTask(umbilical, reporter);
            return;
        }
        if (jobSetup) {
            runJobSetupTask(umbilical, reporter);
            return;
        }
        if (taskCleanup) {
            runTaskCleanupTask(umbilical, reporter);
            return;
        }

        if (useNewApi) {
            // 这里就是最底层执行 Mapper 业务逻辑入口

```

```

        runNewMapper(job, splitMetaInfo, umbilical, reporter);
    } else {
        runOldMapper(job, splitMetaInfo, umbilical, reporter);
    }
    done(umbilical, reporter);
}
}

```

```

@SuppressWarnings("unchecked")
private <INKEY, INVALUE, OUTKEY, OUTVALUE>
void runNewMapper(final JobConf job,
                  final TaskSplitIndex splitIndex,
                  final TaskUmbilicalProtocol umbilical,
                  TaskReporter reporter
) throws IOException, ClassNotFoundException,
    InterruptedException {
    // make a task context so we can get the classes
    // 创建 Task 上下文 TaskAttemptContextImpl
    org.apache.hadoop.mapreduce.TaskAttemptContext taskContext =
        new org.apache.hadoop.mapreduce.task.TaskAttemptContextImpl(job,
            getTaskID(),
            reporter);

    // make a mapper
    // 获取用户自定义 Mapper 类
    org.apache.hadoop.mapreduce.Mapper<INKEY, INVALUE, OUTKEY, OUTVALUE> mapper
    =
        (org.apache.hadoop.mapreduce.Mapper<INKEY, INVALUE, OUTKEY,
OUTVALUE>)
            ReflectionUtils.newInstance(taskContext.getMapperClass(),
job);

    // make the input format
    // 获取用户自定义输入类 默认 TextInputFormat
    org.apache.hadoop.mapreduce.InputFormat<INKEY, INVALUE> inputFormat =
        (org.apache.hadoop.mapreduce.InputFormat<INKEY, INVALUE>)
            ReflectionUtils.newInstance(taskContext.getInputFormatClass(), job);

    // rebuild the input split
    // 获取任务切片信息并解析当前 MapTask 处理哪些数据
    org.apache.hadoop.mapreduce.InputSplit split = null;
    split = getSplitDetails(new Path(splitIndex.getSplitLocation()),
        splitIndex.getStartOffset());
    // hdfs://hadoop102:8020/wordcount/input/LICENSE.txt:0+147145
    LOG.info("Processing split: " + split);

    // 根据输入类型以及切片信息构建创建读取数据对象 NewTrackingRecordReader
    org.apache.hadoop.mapreduce.RecordReader<INKEY, INVALUE> input =
        new NewTrackingRecordReader<INKEY, INVALUE>
            (split, inputFormat, reporter, taskContext);

    job.setBoolean(JobContext.SKIP_RECORDS, isSkipping());
    org.apache.hadoop.mapreduce.RecordWriter output = null;

```

```

// get an output object
// 判断 ReduceTask 个数的情况
if (job.getNumReduceTasks() == 0) {
    output =
        new NewDirectOutputCollector(taskContext, job, umbilical,
reporter);
} else {
    output = new NewOutputCollector(taskContext, job, umbilical, reporter);
}

// 创建 MapTask 上下文对象 MapContextImpl
org.apache.hadoop.mapreduce.MapContext<INKEY, INVALUE, OUTKEY, OUTVALUE>
    mapContext =
        new MapContextImpl<INKEY, INVALUE, OUTKEY, OUTVALUE>(job,
getTaskID(),
        input, output,
        committer,
        reporter, split);

// 包装 MapTask 上下文对象 MapContextImpl
org.apache.hadoop.mapreduce.Mapper<INKEY, INVALUE, OUTKEY, OUTVALUE>.Context
    mapperContext =
        new WrappedMapper<INKEY, INVALUE, OUTKEY, OUTVALUE>().getMapContext(
        mapContext);

try {
    // 初始化输入数据环境 调用 NewTrackingRecordReader.initialize()
    input.initialize(split, mapperContext);
    // 调用 Mapper.run() 执行读取数据处理业务逻辑
    mapper.run(mapperContext);
    mapPhase.complete();
    setPhase(TaskStatus.Phase.SORT);
    statusUpdate(umbilical);
    input.close();
    input = null;
    output.close(mapperContext);
    output = null;
} finally {
    closeQuietly(input);
    closeQuietly(output, mapperContext);
}
}

```

6.3.1 执行 Reducer调用 ReduceTask.run()

```

/**
 * A Reduce task.
 */
@interfaceAudience.Private
@interfacestability.Unstable

```

```

public class ReduceTask extends Task {
    @Override
    @SuppressWarnings("unchecked")
    public void run(JobConf job, final TaskUmbilicalProtocol umbilical)
        throws IOException, InterruptedException, ClassNotFoundException {
        job.setBoolean(JobContext.SKIP_RECORDS, isSkipping());

        if (isMapOrReduce()) {
            copyPhase = getProgress().addPhase("copy");
            sortPhase = getProgress().addPhase("sort");
            reducePhase = getProgress().addPhase("reduce");
        }
        // start thread that will handle communication with parent
        TaskReporter reporter = startReporter(umbilical);

        // 默认 true
        boolean useNewApi = job.getUseNewReducer();
        initialize(job, getJobID(), reporter, useNewApi);

        // check if it is a cleanupJobTask
        if (jobCleanup) {
            runJobCleanupTask(umbilical, reporter);
            return;
        }
        if (jobSetup) {
            runJobSetupTask(umbilical, reporter);
            return;
        }
        if (taskCleanup) {
            runTaskCleanupTask(umbilical, reporter);
            return;
        }

        // Initialize the codec
        // 初始化编码器 默认 DefaultCodec
        codec = initCodec();
        RawKeyValueIterator rIter = null;
        ShuffleConsumerPlugin shuffleConsumerPlugin = null;

        // 获取 Combiner 类
        Class combinerClass = conf.getCombinerClass();
        CombineOutputCollector combineCollector =
            (null != combinerClass) ?
                new CombineOutputCollector(reduceCombineOutputCounter,
                    reporter, conf) : null;

        Class<? extends ShuffleConsumerPlugin> clazz =
            job.getClass(MRConfig.SHUFFLE_CONSUMER_PLUGIN, Shuffle.class,
                ShuffleConsumerPlugin.class);

        shuffleConsumerPlugin = ReflectionUtils.newInstance(clazz, job);
    }
}

```

```

        // Using ShuffleConsumerPlugin:
        org.apache.hadoop.mapreduce.task.reduce.Shuffle@6622fc65
        LOG.info("Using ShuffleConsumerPlugin: " + shuffleConsumerPlugin);

        ShuffleConsumerPlugin.Context shuffleContext =
            new ShuffleConsumerPlugin.Context(getTaskID(), job,
            FileSystem.getLocal(job), umbilical,
                super.lDirAlloc, reporter, codec,
                combinerClass, combineCollector,
                spilledRecordsCounter, reduceCombineInputCounter,
                shuffledMapsCounter,
                reduceShuffleBytes, failedShuffleCounter,
                mergedMapOutputsCounter,
                taskStatus, copyPhase, sortPhase, this,
                mapOutputFile, localMapFiles);
        shuffleConsumerPlugin.init(shuffleContext);

        rIter = shuffleConsumerPlugin.run();

        // free up the data structures
        mapOutputFilesOnDisk.clear();

        sortPhase.complete(); // sort is complete
        setPhase(TaskStatus.Phase.REDUCE);
        statusUpdate(umbilical);
        Class keyClass = job.getMapOutputKeyClass();
        Class valueClass = job.getMapOutputValueClass();
        RawComparator comparator = job.getOutputValueGroupingComparator();

        if (useNewApi) {
            // Reducer 执行入口
            runNewReducer(job, umbilical, reporter, rIter, comparator,
                keyClass, valueClass);
        } else {
            runOldReducer(job, umbilical, reporter, rIter, comparator,
                keyClass, valueClass);
        }

        shuffleConsumerPlugin.close();
        done(umbilical, reporter);
    }
}

```

```

@SuppressWarnings("unchecked")
private <INKEY, INVALUE, OUTKEY, OUTVALUE>
void runNewReducer(JobConf job,
    final TaskUmbilicalProtocol umbilical,
    final TaskReporter reporter,
    RawKeyValueIterator rIter,
    RawComparator<INKEY> comparator,
    Class<INKEY> keyClass,

```



```

        Class<INVALUE> valueClass
    ) throws IOException, InterruptedException,
        ClassNotFoundException {
        // wrap value iterator to report progress.
        final RawKeyValueIterator rawIter = rIter;
        rIter = new RawKeyValueIterator() {
            public void close() throws IOException {
                rawIter.close();
            }

            public DataInputBuffer getKey() throws IOException {
                return rawIter.getKey();
            }

            public Progress getProgress() {
                return rawIter.getProgress();
            }

            public DataInputBuffer getValue() throws IOException {
                return rawIter.getValue();
            }

            public boolean next() throws IOException {
                boolean ret = rawIter.next();
                reporter.setProgress(rawIter.getProgress().getProgress());
                return ret;
            }
        };
        // make a task context so we can get the classes
        org.apache.hadoop.mapreduce.TaskAttemptContext taskContext =
            new org.apache.hadoop.mapreduce.task.TaskAttemptContextImpl(job,
                getTaskID(), reporter);
        // make a reducer
        org.apache.hadoop.mapreduce.Reducer<INKEY, INVALUE, OUTKEY, OUTVALUE>
reducer =
            (org.apache.hadoop.mapreduce.Reducer<INKEY, INVALUE, OUTKEY,
OUTVALUE>)
                ReflectionUtils.newInstance(taskContext.getReducerClass(),
job);
        org.apache.hadoop.mapreduce.RecordWriter<OUTKEY, OUTVALUE> trackedRW =
            new NewTrackingRecordWriter<OUTKEY, OUTVALUE>(this, taskContext);
        job.setBoolean("mapred.skip.on", isSkipping());
        job.setBoolean(JobContext.SKIP_RECORDS, isSkipping());
        org.apache.hadoop.mapreduce.Reducer.Context
            reducerContext = createReduceContext(reducer, job, getTaskID(),
                rIter, reduceInputKeyCounter,
                reduceInputValueCounter,
                trackedRW,
                committer,
                reporter, comparator, keyClass,
                valueClass);
        try {

```

```
        // 执行
        reducer.run(reducerContext);
    } finally {
        trackedRW.close(reducerContext);
    }
}
```

```
/**
 * Advanced application writers can use the
 * {@link #run(org.apache.hadoop.mapreduce.Reducer.Context)} method to
 * control how the reduce task works.
 */
public void run(Context context) throws IOException, InterruptedException {
    setup(context);
    try {
        while (context.nextKey()) {
            reduce(context.getCurrentKey(), context.getValues(), context);
            // If a back up store is used, reset it
            Iterator<VALUEIN> iter = context.getValues().iterator();
            if(iter instanceof ReduceContext.ValueIterator) {
                ((ReduceContext.ValueIterator<VALUEIN>)iter).resetBackupStore();
            }
        }
    } finally {
        cleanup(context);
    }
}
```